

Kafka Producers: Complete Developer Guide

A comprehensive refresher on Apache Kafka Producers, designed for both beginners and experienced developers. This README covers producer basics, partitioning strategies, reliability features, and real-world applications with detailed Java examples.

Table of Contents

-  [Producer Basics](#)
 - [KafkaProducer API](#)
 - [Asynchronous vs Synchronous Send](#)
 - [ProducerConfig Deep Dive](#)
-  [Partitioning & Ordering](#)
 - [Default Partitioner](#)
 - [Custom Partitioner](#)
 - [Ordering Guarantees](#)
-  [Reliability](#)
 - [Idempotent Producer](#)
 - [Transactional Producer](#)
 - [Delivery Semantics](#)
-  [Comprehensive Java Examples](#)
-  [Comparisons & Trade-offs](#)
-  [Common Pitfalls & Best Practices](#)
-  [Real-World Use Cases](#)
-  [Version Highlights](#)
-  [Additional Resources](#)

Producer Basics

KafkaProducer API

Simple Explanation

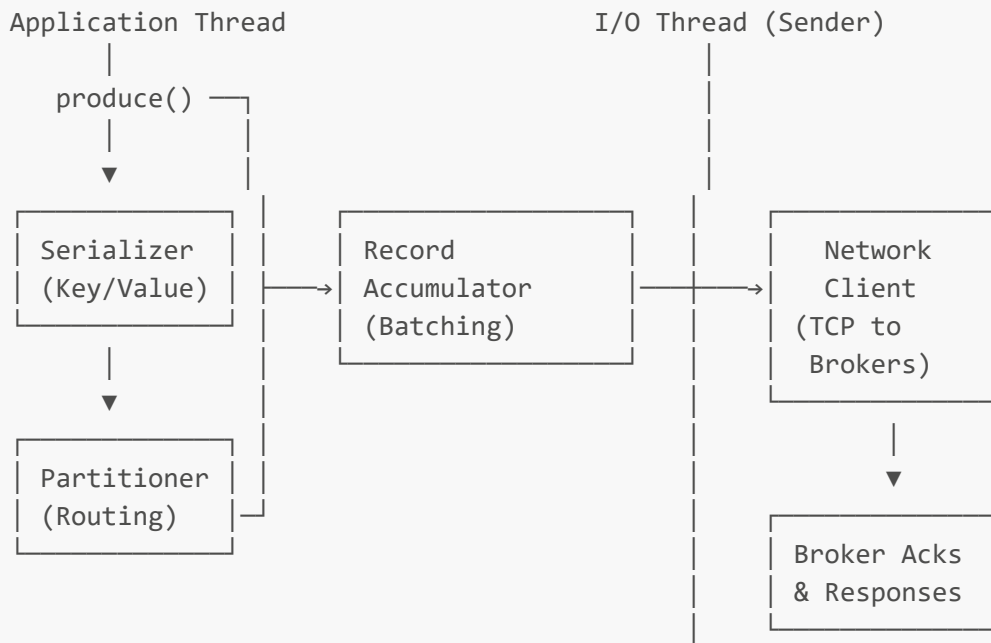
The `KafkaProducer` is the client API that applications use to publish records to Kafka topics. It's thread-safe and designed for high throughput with batching and asynchronous operations.

Problem It Solves

- **High Throughput:** Batches multiple messages for efficient network utilization
- **Fault Tolerance:** Handles retries and failures automatically
- **Scalability:** Distributes messages across topic partitions
- **Performance:** Asynchronous operations don't block application threads

Internal Architecture

KafkaProducer Internal Components:



Core Producer Configuration

```

Properties props = new Properties();

// Connection settings
props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(ProducerConfig.CLIENT_ID_CONFIG, "my-producer");

// Serialization
props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.StringSerializer");
props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.StringSerializer");

// Performance & Batching
props.put(ProducerConfig.BATCH_SIZE_CONFIG, 16384);           // 16KB batches
props.put(ProducerConfig.LINGER_MS_CONFIG, 5);               // Wait 5ms for batch
props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "gzip");    // Compress batches

// Reliability
props.put(ProducerConfig.ACKS_CONFIG, "all");                // Wait for all
// replicas
props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE); // Retry forever
props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);    // Prevent duplicates
  
```

Asynchronous vs Synchronous Send

Asynchronous Send (Recommended)

```

import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.serialization.StringSerializer;
import java.util.Properties;
import java.util.concurrent.Future;

public class AsyncProducerExample {
    private static final String TOPIC = "user-events";

    public static void main(String[] args) {
        Properties props = createProducerConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {

            for (int i = 0; i < 1000; i++) {
                String key = "user-" + i;
                String value = "login-event-" + System.currentTimeMillis();

                ProducerRecord<String, String> record =
                    new ProducerRecord<>(TOPIC, key, value);

                // Asynchronous send with callback
                producer.send(record, new Callback() {
                    @Override
                    public void onCompletion(RecordMetadata metadata, Exception
exception) {
                        if (exception != null) {
                            System.err.println("Failed to send record: " +
exception.getMessage());
                        }
                        // Implement retry logic or dead letter queue
                    } else {
                        System.out.printf("Sent record to partition %d, offset
%d, timestamp %d%n",
                                metadata.partition(), metadata.offset(),
                                metadata.timestamp());
                    }
                });

                // Don't block - continue sending more records
            }

            // Ensure all records are sent before closing
            producer.flush();

        } catch (Exception e) {
            System.err.println("Producer failed: " + e.getMessage());
        }
    }

    private static Properties createProducerConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    }
}

```

```

        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

        // High-throughput settings
        props.put(ProducerConfig.BATCH_SIZE_CONFIG, 16384);
        props.put(ProducerConfig.LINGER_MS_CONFIG, 5);
        props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "gzip");
        props.put(ProducerConfig.ACKS_CONFIG, "1");

        return props;
    }
}

```

Advantages:

- ☒ **High Throughput:** Non-blocking sends allow batching
- ☒ **Better Resource Utilization:** Application thread doesn't wait
- ☒ **Scalability:** Can send thousands of messages per second

Trade-offs:

- ⚠️ **Complexity:** Need to handle callbacks for error handling
- ⚠️ **Memory Usage:** Messages buffered in memory until sent

Synchronous Send (Use Sparingly)

```

public class SyncProducerExample {
    public static void main(String[] args) {
        Properties props = createProducerConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {

            for (int i = 0; i < 10; i++) {
                ProducerRecord<String, String> record =
                    new ProducerRecord<>("critical-events", "key-" + i, "critical-
data-" + i);

                try {
                    // Synchronous send - blocks until complete
                    RecordMetadata metadata = producer.send(record).get();
                    System.out.printf("Sent record synchronously to partition %d,
offset %d\n",
                                metadata.partition(), metadata.offset());
                } catch (Exception e) {
                    System.err.println("Failed to send record synchronously: " +
e.getMessage());
                    // Handle error immediately

```

```

        break;
    }
}
}
}
}

```

When to Use:

- 🚫 **Critical Messages:** When you need immediate confirmation
- 🚫 **Error Handling:** When you must handle failures immediately
- 🚫 **Ordering:** When strict order is required (with `max.in.flight.requests=1`)

Trade-offs:

- ✖ **Low Throughput:** Blocks on every send
- ✖ **Poor Performance:** No batching benefits
- ✖ **Latency:** Application waits for network round trips

ProducerConfig Deep Dive

Core Configuration Parameters

```

public class ProducerConfigExample {
    public static Properties createProductionConfig() {
        Properties props = new Properties();

        // === CONNECTION SETTINGS ===
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
            "broker1:9092,broker2:9092,broker3:9092");
        props.put(ProducerConfig.CLIENT_ID_CONFIG, "payment-service-producer");

        // === SERIALIZATION ===
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.connect.json.JsonSerializer"); // For JSON payloads

        // === RELIABILITY SETTINGS ===
        props.put(ProducerConfig.ACKS_CONFIG, "all"); // Wait for all
in-sync replicas
        props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE); // Retry
indefinitely
        props.put(ProducerConfig.RETRY_BACKOFF_MS_CONFIG, 100); // Wait 100ms
between retries
        props.put(ProducerConfig.REQUEST_TIMEOUT_MS_CONFIG, 30000); // 30s request
timeout
        props.put(ProducerConfig.DELIVERY_TIMEOUT_MS_CONFIG, 120000); // 2min
total timeout

        // === IDEMPOTENCE & ORDERING ===

```

```
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true); // Prevent
duplicates
        props.put(ProducerConfig.MAX_IN_FLIGHT_REQUESTS_PER_CONNECTION, 5); //
Default with idempotence

        // === BATCHING & PERFORMANCE ===
        props.put(ProducerConfig.BATCH_SIZE_CONFIG, 32768); // 32KB
batches
        props.put(ProducerConfig.LINGER_MS_CONFIG, 10); // Wait 10ms
to fill batch
        props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "gzip"); // Compress
data
        props.put(ProducerConfig.BUFFER_MEMORY_CONFIG, 67108864); // 64MB total
buffer

        // === NETWORK SETTINGS ===
        props.put(ProducerConfig.SEND_BUFFER_CONFIG, 131072); // 128KB TCP
send buffer
        props.put(ProducerConfig.RECEIVE_BUFFER_CONFIG, 65536); // 64KB TCP
receive buffer

        return props;
    }
}
```

Configuration Trade-offs

| Parameter | Higher Value | Lower Value |
|---------------|--------------------------------------|-----------------------------------|
| batch.size | ↑ Throughput, ↑ Latency, ↑ Memory | ↓ Latency, ↓ Throughput, ↓ Memory |
| linger.ms | ↑ Batching efficiency, ↑ Latency | ↓ Latency, ↓ Batching |
| acks | ↑ Durability, ↓ Throughput | ↓ Durability, ↑ Throughput |
| retries | ↑ Reliability, ↑ Latency | ↓ Reliability, ↓ Latency |
| buffer.memory | ↑ Buffering capacity, ↑ Memory usage | ↓ Memory, may block sends |

Environment-Specific Configurations

```
public class EnvironmentConfigs {

    // Development - Low latency, simple debugging
    public static Properties developmentConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BootstrapServersConfig, "localhost:9092");
        props.put(ProducerConfig.AcksConfig, "1"); // Leader
only
        props.put(ProducerConfig.RetriesConfig, 3); // Limited
retries
        props.put(ProducerConfig.LingerMsConfig, 0); // No
    }
}
```

```

batching delay
    props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, false);    //
Simplicity
    return props;
}

// Production - High reliability and performance
public static Properties productionConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BootstrapServersConfig,
        "prod-broker1:9092,prod-broker2:9092,prod-broker3:9092");
    props.put(ProducerConfig.AcksConfig, "all");                    // Maximum
durability
    props.put(ProducerConfig.RetriesConfig, Integer.MAX_VALUE);    // Retry
forever
    props.put(ProducerConfig.LingerMsConfig, 10);                  // Optimize
batching
    props.put(ProducerConfig.BatchSizeConfig, 32768);              // Larger
batches
    props.put(ProducerConfig.CompressionTypeConfig, "gzip");       // Reduce
bandwidth
    props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);     // Exactly-
once
    return props;
}

// High-throughput - Maximum performance
public static Properties highThroughputConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BootstrapServersConfig, "localhost:9092");
    props.put(ProducerConfig.AcksConfig, "0");                     // Fire and
forget
    props.put(ProducerConfig.LingerMsConfig, 100);                 //
Aggressive batching
    props.put(ProducerConfig.BatchSizeConfig, 65536);              // 64KB
batches
    props.put(ProducerConfig.CompressionTypeConfig, "lz4");        // Fast
compression
    props.put(ProducerConfig.BufferMemoryConfig, 134217728);       // 128MB
buffer
    return props;
}
}

```

Partitioning & Ordering

Default Partitioner

How Default Partitioner Works

Partitioning Logic:

1. If record has a key:
partition = hash(key) % num_partitions
2. If record has no key:
partition = round-robin (sticky partitioner in Kafka 2.4+)

Examples:

Topic: "orders" (3 partitions)

Key-based:

"user123" → hash("user123") % 3 = 1 → Partition 1

"user456" → hash("user456") % 3 = 0 → Partition 0

"user789" → hash("user789") % 3 = 2 → Partition 2

No key:

null → Round-robin → Partition 0, 1, 2, 0, 1, 2...

Default Partitioner Java Example

```
import org.apache.kafka.clients.producer.*;
import java.util.Properties;

public class DefaultPartitionerExample {

    public static void main(String[] args) {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {
            // Key-based partitioning - all messages for same user go to same
            partition
            sendWithKey(producer, "user123", "User 123 placed order #1");
            sendWithKey(producer, "user123", "User 123 placed order #2"); // Same
            partition
            sendWithKey(producer, "user456", "User 456 placed order #1"); //
            Different partition

            // Round-robin partitioning - no key
            sendWithoutKey(producer, "Anonymous order #1");
            sendWithoutKey(producer, "Anonymous order #2");
            sendWithoutKey(producer, "Anonymous order #3");
        }
    }
}
```



```

        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    private static void sendWithKey(KafkaProducer<String, String> producer,
                                    String key, String value) {
        ProducerRecord<String, String> record =
            new ProducerRecord<>("orders", key, value);

        producer.send(record, (metadata, exception) -> {
            if (exception == null) {
                System.out.printf("Key: %s → Partition: %d, Offset: %d\n",
                                   key, metadata.partition(), metadata.offset());
            }
        });
    }

    private static void sendWithoutKey(KafkaProducer<String, String> producer,
                                       String value) {
        ProducerRecord<String, String> record =
            new ProducerRecord<>("orders", null, value);

        producer.send(record, (metadata, exception) -> {
            if (exception == null) {
                System.out.printf("No key → Partition: %d, Offset: %d\n",
                                   metadata.partition(), metadata.offset());
            }
        });
    }
}

```

Sticky Partitioner (Kafka 2.4+)

Problem with Old Round-Robin:

- Poor batching efficiency when no key provided
- Each message could go to different partition, reducing batch effectiveness

Sticky Partitioner Solution:

```





// Old behavior (pre-2.4):
// msg1 → partition 0, msg2 → partition 1, msg3 → partition 2, msg4 → partition
// 0...

// Sticky partitioner (2.4+):
// msg1-10 → partition 0, msg11-20 → partition 1, msg21-30 → partition 2...
// Sticks to partition until batch is sent, then switches

```

Custom Partitioner

When to Use Custom Partitioner

-  **Business Logic:** Route messages based on business rules
-  **Load Balancing:** Custom distribution strategies
-  **Data Locality:** Keep related data together
-  **Performance:** Optimize for specific access patterns

Custom Partitioner Implementation

```
import org.apache.kafka.clients.producer.Partitioner;
import org.apache.kafka.common.Cluster;
import org.apache.kafka.common.PartitionInfo;
import java.util.List;
import java.util.Map;

// Route VIP customers to specific partition for priority processing
public class VipCustomerPartitioner implements Partitioner {

    private static final int VIP_PARTITION = 0; // Dedicate partition 0 for VIP
    customers

    @Override
    public int partition(String topic, Object key, byte[] keyBytes,
                        Object value, byte[] valueBytes, Cluster cluster) {

        List<PartitionInfo> partitions = cluster.partitionsForTopic(topic);
        int numPartitions = partitions.size();

        if (key == null) {
            // No key - use round-robin for remaining partitions
            return (int) (Math.random() * (numPartitions - 1)) + 1;
        }

        String keyStr = key.toString();

        // VIP customers go to partition 0
        if (keyStr.startsWith("vip-")) {
            return VIP_PARTITION;
        }

        // Premium customers get better partitions (1-2)
        if (keyStr.startsWith("premium-")) {
            return (keyStr.hashCode() & Integer.MAX_VALUE) % 2 + 1;
        }

        // Regular customers use remaining partitions (3+)
        if (numPartitions > 3) {
            return (keyStr.hashCode() & Integer.MAX_VALUE) % (numPartitions - 3) +
3;
        }

        // Fallback to default behavior
    }
}
```

```

        return (keyStr.hashCode() & Integer.MAX_VALUE) % numPartitions;
    }

    @Override
    public void close() {
        // Cleanup resources if needed
    }

    @Override
    public void configure(Map<String, ?> configs) {
        // Configuration if needed
    }
}

// Geographic partitioner - route by region
public class GeographicPartitioner implements Partitioner {

    @Override
    public int partition(String topic, Object key, byte[] keyBytes,
                        Object value, byte[] valueBytes, Cluster cluster) {

        if (key == null) {
            return 0; // Default partition
        }

        String keyStr = key.toString();
        int numPartitions = cluster.partitionsForTopic(topic).size();

        // Extract region from key (format: "region-user-id")
        String[] parts = keyStr.split("-");
        if (parts.length >= 2) {
            String region = parts[0].toLowerCase();

            switch (region) {
                case "us": return 0;
                case "eu": return 1;
                case "asia": return 2;
                case "latam": return 3;
                default:
                    return (keyStr.hashCode() & Integer.MAX_VALUE) %
numPartitions;
            }
        }

        return (keyStr.hashCode() & Integer.MAX_VALUE) % numPartitions;
    }

    @Override
    public void close() {}

    @Override
    public void configure(Map<String, ?> configs) {}
}

```

Using Custom Partitioner

```
public class CustomPartitionerExample {

    public static void main(String[] args) {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // Configure custom partitioner
        props.put(ProducerConfig.PARTITIONER_CLASS_CONFIG,
            "com.example.VipCustomerPartitioner");

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {

            // VIP customer - goes to partition 0
            sendMessage(producer, "vip-customer-001", "VIP order: $10000");

            // Premium customer - goes to partition 1-2
            sendMessage(producer, "premium-customer-123", "Premium order: $1000");

            // Regular customer - goes to partition 3+
            sendMessage(producer, "regular-customer-456", "Regular order: $100");

            // Geographic routing
            sendMessage(producer, "us-customer-789", "US order");
            sendMessage(producer, "eu-customer-321", "European order");
            sendMessage(producer, "asia-customer-654", "Asian order");

        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    private static void sendMessage(KafkaProducer<String, String> producer,
                                    String key, String value) {
        ProducerRecord<String, String> record =
            new ProducerRecord<>("customer-orders", key, value);

        producer.send(record, (metadata, exception) -> {
            if (exception == null) {
                System.out.printf("Key: %s → Partition: %d (Custom Logic)%n",
                    key, metadata.partition());
            } else {
                System.err.println("Failed to send: " + exception.getMessage());
            }
        });
    }
}
```


```
}
}
```

Ordering Guarantees

Single Partition Ordering

Single Partition = Total Order

Partition 0:



Consumer reads: M1, M2, M3, M4, M5 (always same order)

```
// Ensuring single partition for ordering
public class OrderedMessagesExample {

    public static void main(String[] args) {
        Properties props = createProducerConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {

            String userId = "user123"; // Same key = same partition

            // These messages will be ordered within the partition
            sendOrderedMessage(producer, userId, "User created account");
            sendOrderedMessage(producer, userId, "User verified email");
            sendOrderedMessage(producer, userId, "User updated profile");
            sendOrderedMessage(producer, userId, "User made first purchase");

            producer.flush(); // Ensure all sent

        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    private static void sendOrderedMessage(KafkaProducer<String, String> producer,
                                           String userId, String event) {
        ProducerRecord<String, String> record =
            new ProducerRecord<>("user-events", userId, event);

        // For strict ordering, consider using synchronous send
        try {
```

```
        RecordMetadata metadata = producer.send(record).get();
        System.out.printf("Ordered event for %s: %s → Partition %d, Offset %d%n",
            userId, event, metadata.partition(), metadata.offset());
    } catch (Exception e) {
        System.err.println("Failed to send ordered message: " +
e.getMessage());
    }
}

private static Properties createProducerConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringSerializer");
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringSerializer");

    // Settings for strict ordering
    props.put(ProducerConfig.MAX_IN_FLIGHT_REQUESTS_PER_CONNECTION, 1);
    props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE);
    props.put(ProducerConfig.ACKS_CONFIG, "all");

    return props;
}
```

Multiple Partitions - No Global Order

Multiple Partitions = No Global Order

Partition 0: [M1, M3, M5] ← User A events
Partition 1: [M2, M4, M6] ← User B events

Consumer Group:
Consumer 1 reads P0: M1, M3, M5
Consumer 2 reads P1: M2, M4, M6

Global consumption order could be: M1, M2, M3, M4, M5, M6
OR: M2, M1, M4, M3, M6, M5
OR: Any interleaving!

Ordering Configuration Trade-offs

| Configuration | Ordering Guarantee | Performance Impact |
|--------------------------|--|--|
| max.in.flight.requests=1 | <input checked="" type="checkbox"/> Strict per partition | <input checked="" type="checkbox"/> Lower throughput |

| Configuration | Ordering Guarantee | Performance Impact |
|--|--|----------------------|
| <code>max.in.flight.requests=5</code> + <code>enable.idempotence=true</code> | <input checked="" type="checkbox"/> Order preserved with retries | ★ Good throughput |
| <code>max.in.flight.requests=5</code> + <code>enable.idempotence=false</code> | ✗ May reorder on retry | ★★ Higher throughput |

Reliability

Idempotent Producer

Problem It Solves

Producer Retry Duplication:

Without Idempotence:

1. Producer sends message M1
2. Broker receives M1, but ACK is lost
3. Producer retries, sends M1 again
4. Result: M1 appears twice in log

With Idempotence:

1. Producer sends M1 with sequence number
2. Broker receives M1, but ACK is lost
3. Producer retries M1 with same sequence number
4. Broker recognizes duplicate, ignores retry
5. Result: M1 appears once in log

Internal Mechanism

Idempotent Producer Internals:

Producer ID (PID): Unique producer identifier
Sequence Number: Per-partition counter starting from 0

Message Format:

| | | | |
|-------|-------|----------|---------|
| PID | Epoch | Sequence | Message |
| 12345 | 0 | 0 | "Hello" |

Broker State:

Partition 0: Last seen (PID=12345, Epoch=0, Seq=2)
Partition 1: Last seen (PID=12345, Epoch=0, Seq=5)

Idempotent Producer Java Example

```
import org.apache.kafka.clients.producer.*;
import java.util.Properties;
import java.util.concurrent.ExecutionException;

public class IdempotentProducerExample {

    public static void main(String[] args) {
        Properties props = createIdempotentProducerConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {

            // Simulate sending critical financial transactions
            for (int i = 0; i < 10; i++) {
                String transactionId = "txn-" + i;
                String transactionData = String.format(
                    "Transfer $%.2f from account A to account B",
                    1000.0 + (i * 100));

                sendIdempotentMessage(producer, transactionId, transactionData);
            }

            // Force any buffered messages to be sent
            producer.flush();

        } catch (Exception e) {
            System.err.println("Idempotent producer failed: " + e.getMessage());
            e.printStackTrace();
        }

        private static void sendIdempotentMessage(KafkaProducer<String, String>
producer,
                                                String key, String value) {
            ProducerRecord<String, String> record =
                new ProducerRecord<>("financial-transactions", key, value);

            producer.send(record, new Callback() {
                @Override
                public void onCompletion(RecordMetadata metadata, Exception exception)
            {
                if (exception != null) {
                    if (exception instanceof
org.apache.kafka.common.errors.OutOfOrderSequenceException) {
                        System.err.println("Out of order sequence - this shouldn't
happen with idempotence!");
                    } else if (exception instanceof
org.apache.kafka.common.errors.DuplicateSequenceNumberException) {
                        System.out.println("Duplicate detected and handled by
broker");
                    }
                }
            }
        }
    }
}
```



```

        } else {
            System.err.println("Send failed: " +
exception.getMessage());
        }
    } else {
        System.out.printf("Idempotent send: %s → P%d:O%d%n",
            key, metadata.partition(), metadata.offset());
    }
}
});
}

private static Properties createIdempotentProducerConfig() {
    Properties props = new Properties();

    // Basic configuration
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringSerializer");
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringSerializer");

    // Idempotence configuration (Kafka 3.0+ defaults)
    props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);

    // Required settings for idempotence
    props.put(ProducerConfig.ACKS_CONFIG, "all"); // Must wait
for all replicas
    props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE); // Must
enable retries
    props.put(ProducerConfig.MAX_IN_FLIGHT_REQUESTS_PER_CONNECTION, 5); // Max
5 for idempotence

    // Additional reliability
    props.put(ProducerConfig.RETRY_BACKOFF_MS_CONFIG, 100);
    props.put(ProducerConfig.REQUEST_TIMEOUT_MS_CONFIG, 30000);

    return props;
}
}

```

Idempotence Limitations & Considerations

```

public class IdempotenceLimitations {

    public void demonstrateLimitations() {
        // 1. Producer restart loses PID - can't prevent duplicates across
restarts
        System.out.println("Limitation 1: Producer restart breaks idempotence");

        // 2. Only works within producer session
    }
}

```

```
        System.out.println("Limitation 2: Only prevents retries within same
producer instance");

        // 3. Only prevents duplicates from producer retries, not application
retries
        System.out.println("Limitation 3: Application-level retries can still
create duplicates");
    }

    // For cross-session idempotence, use transactional producer
    public Properties transactionalConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // Transactional settings
        props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "payment-processor-1");
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true); //
Automatically true
        props.put(ProducerConfig.ACKS_CONFIG, "all");

        return props;
    }
}
```

Transactional Producer

Problem It Solves

- **Exactly-Once Processing:** Guaranteed exactly-once semantics across producer restarts
- **Atomic Multi-Topic Writes:** All messages in transaction committed together
- **Read Committed:** Consumers only see committed messages

Transactional Producer Architecture

Transactional Producer Flow:

1. Producer gets Transaction Coordinator assignment

2. Producer begins transaction

3. Producer sends messages to multiple topics/partitions

4. Producer commits transaction

5. Transaction Coordinator marks all messages as committed

6. Consumers with isolation.level=read_committed see messages

Transaction Log:

| | | | |
|--------|--------|------------|-----------|
| TXN_ID | Status | Partitions | Timestamp |
|--------|--------|------------|-----------|

| | | | |
|-----------|-----------|-------------------------|------------|
| payment-1 | ONGOING | orders:0 inventory:1 | 1693875600 |
| payment-1 | COMMITTED | orders:0 inventory:1 | 1693875605 |

Comprehensive Transactional Producer Example

```
import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.errors.ProducerFencedException;
import org.apache.kafka.common.errors.OutOfOrderSequenceException;
import org.apache.kafka.common.errors.AuthorizationException;
import java.util.Properties;

public class TransactionalProducerExample {

    private static final String ORDERS_TOPIC = "orders";
    private static final String INVENTORY_TOPIC = "inventory";
    private static final String PAYMENTS_TOPIC = "payments";

    public static void main(String[] args) {
        Properties props = createTransactionalProducerConfig();

        KafkaProducer<String, String> producer = new KafkaProducer<>(props);

        // Initialize transactions
        producer.initTransactions();

        try {
            // Process multiple orders atomically
            for (int i = 0; i < 5; i++) {
                processOrderTransactionally(producer, "order-" + i, i * 100.0);
            }
        } catch (Exception e) {
            System.err.println("Transactional processing failed: " +
e.getMessage());
        } finally {
            producer.close();
        }
    }

    private static void processOrderTransactionally(KafkaProducer<String, String>
producer,
                                                    String orderId, double amount)
    {
        try {
            // Begin transaction
            producer.beginTransaction();

```

```

        // Send order creation event
        ProducerRecord<String, String> orderRecord = new ProducerRecord<>(
            ORDERS_TOPIC, orderId,
            String.format("Order created: %s, amount: $%.2f", orderId,
amount));
        producer.send(orderRecord);

        // Reserve inventory
        ProducerRecord<String, String> inventoryRecord = new ProducerRecord<>(
            INVENTORY_TOPIC, orderId,
            String.format("Inventory reserved for order: %s", orderId));
        producer.send(inventoryRecord);

        // Process payment
        ProducerRecord<String, String> paymentRecord = new ProducerRecord<>(
            PAYMENTS_TOPIC, orderId,
            String.format("Payment processed: %s, amount: $%.2f", orderId,
amount));
        producer.send(paymentRecord);

        // Simulate business logic that might fail
        if (orderId.equals("order-3")) {
            throw new RuntimeException("Payment processor temporarily
unavailable");
        }

        // Commit transaction - all messages become visible atomically
        producer.commitTransaction();
        System.out.println("Transaction committed successfully for " +
orderId);

    } catch (ProducerFencedException e) {
        // Another producer with same transactional.id is active
        System.err.println("Producer fenced: " + e.getMessage());
        producer.close();
    } catch (OutOfOrderSequenceException | AuthorizationException e) {
        // Unrecoverable errors
        System.err.println("Unrecoverable error: " + e.getMessage());
        producer.close();
    } catch (Exception e) {
        // Abort transaction on any error
        System.err.println("Transaction failed for " + orderId + ": " +
e.getMessage());
        try {
            producer.abortTransaction();
            System.out.println("Transaction aborted for " + orderId);
        } catch (Exception abortEx) {
            System.err.println("Failed to abort transaction: " +
abortEx.getMessage());
        }
    }
}

private static Properties createTransactionalProducerConfig() {

```

```

        Properties props = new Properties();

        // Basic configuration
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // Transactional configuration
        props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "order-processor-1");

        // These are automatically set with transactional.id
        // props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
        // props.put(ProducerConfig.ACKS_CONFIG, "all");
        // props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE);

        // Performance settings
        props.put(ProducerConfig.BATCH_SIZE_CONFIG, 16384);
        props.put(ProducerConfig.LINGER_MS_CONFIG, 5);

        return props;
    }
}

```

Transactional Consumer (Read Committed)

```

import org.apache.kafka.clients.consumer.*;
import java.time.Duration;
import java.util.Arrays;
import java.util.Properties;

public class TransactionalConsumerExample {

    public static void main(String[] args) {
        Properties props = createTransactionalConsumerConfig();

        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
        {
            consumer.subscribe(Arrays.asList("orders", "inventory", "payments"));

            while (true) {
                ConsumerRecords<String, String> records =
consumer.poll(Duration.ofMillis(1000));

                for (ConsumerRecord<String, String> record : records) {
                    // Only see committed messages due to
isolation.level=read_committed
                    System.out.printf("Consumed committed record: topic=%s,
key=%s, value=%s, " +
                        "partition=%d, offset=%d%n",

```

```

        record.topic(), record.key(), record.value(),
        record.partition(), record.offset());
    }

    // Commit offsets
    consumer.commitSync();
}
}

private static Properties createTransactionalConsumerConfig() {
    Properties props = new Properties();
    props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ConsumerConfig.GROUP_ID_CONFIG, "transactional-consumer-group");
    props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringDeserializer");
    props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
        "org.apache.kafka.common.serialization.StringDeserializer");

    // Only read committed messages (default is read_uncommitted)
    props.put(ConsumerConfig.ISOLATION_LEVEL_CONFIG, "read_committed");

    props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
    props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);

    return props;
}
}

```

Delivery Semantics

At-Most-Once (Fire and Forget)

```

public class AtMostOnceProducer {

    public static Properties createConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // At-most-once settings - fire and forget
        props.put(ProducerConfig.ACKS_CONFIG, "0"); // Don't wait for
ack
        props.put(ProducerConfig.RETRIES_CONFIG, 0); // No retries
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, false);

        return props;
    }
}

```

```

    public static void main(String[] args) {
        Properties props = createConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {
            for (int i = 0; i < 100; i++) {
                ProducerRecord<String, String> record =
                    new ProducerRecord<>("metrics", "sensor-" + i, "temperature:"
+ (20 + i));

                // Fire and forget - don't care about result
                producer.send(record);
            }
        }

        System.out.println("Messages sent (at-most-once) - some may be lost, none
duplicated");
    }
}

```

At-Least-Once (Standard Retry)

```

public class AtLeastOnceProducer {

    public static Properties createConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // At-least-once settings
        props.put(ProducerConfig.ACKS_CONFIG, "all"); // Wait for all
replicas
        props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE); // Retry on
failure
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, false); // Allow
duplicates

        return props;
    }

    public static void main(String[] args) {
        Properties props = createConfig();

        try (KafkaProducer<String, String> producer = new KafkaProducer<>(props))
        {
            for (int i = 0; i < 100; i++) {
                ProducerRecord<String, String> record =

```

```

        new ProducerRecord<>("events", "event-" + i, "data-" + i);

        producer.send(record, (metadata, exception) -> {
            if (exception != null) {
                System.err.println("Failed after retries: " +
exception.getMessage());
            } else {
                System.out.printf("Sent: partition=%d, offset=%d%n",
                    metadata.partition(), metadata.offset());
            }
        });
    }

    System.out.println("Messages sent (at-least-once) - none lost, duplicates
possible");
}
}

```

Exactly-Once (Idempotent/Transactional)

```

public class ExactlyOnceProducer {

    // Exactly-once with idempotence (single producer session)
    public static Properties createIdempotentConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // Exactly-once within producer session
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
        // Other settings automatically configured: acks=all, retries=MAX,
max.in.flight=5

        return props;
    }

    // Exactly-once across restarts (transactional)
    public static Properties createTransactionalConfig() {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.StringSerializer");

        // Exactly-once across producer restarts
        props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "exactly-once-

```



```

producer");
    // Idempotence automatically enabled

    return props;
}

public static void main(String[] args) {
    // Example with transactional exactly-once
    Properties props = createTransactionalConfig();

    KafkaProducer<String, String> producer = new KafkaProducer<>(props);
    producer.initTransactions();

    try {
        producer.beginTransaction();

        for (int i = 0; i < 10; i++) {
            ProducerRecord<String, String> record =
                new ProducerRecord<>("financial-events", "txn-" + i,
                    "amount:" + (1000.0 * i));
            producer.send(record);
        }

        producer.commitTransaction();
        System.out.println("Transaction committed - exactly-once guarantee");

    } catch (Exception e) {
        producer.abortTransaction();
        System.err.println("Transaction aborted: " + e.getMessage());
    } finally {
        producer.close();
    }
}
}

```

Comprehensive Java Examples

Complete Producer Application

```

import org.apache.kafka.clients.producer.*;
import org.apache.kafka.common.serialization.StringSerializer;
import com.fasterxml.jackson.databind.ObjectMapper;
import com.fasterxml.jackson.core.JsonProcessingException;

import java.util.Properties;
import java.util.concurrent.CountDownLatch;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;
import java.util.concurrent.atomic.AtomicInteger;

```

```

/**
 * Production-ready Kafka Producer with error handling, metrics, and monitoring
 */
public class ProductionKafkaProducerExample {

    private final KafkaProducer<String, String> producer;
    private final ObjectMapper objectMapper;
    private final AtomicInteger successCount = new AtomicInteger(0);
    private final AtomicInteger errorCount = new AtomicInteger(0);

    public ProductionKafkaProducerExample(Properties config) {
        this.producer = new KafkaProducer<>(config);
        this.objectMapper = new ObjectMapper();
    }

    public static void main(String[] args) {
        Properties config = createProductionConfig();
        ProductionKafkaProducerExample producerApp =
            new ProductionKafkaProducerExample(config);

        // Simulate high-throughput application
        producerApp.runHighThroughputTest();

        // Cleanup
        producerApp.close();
    }

    private static Properties createProductionConfig() {
        Properties props = new Properties();

        // Cluster connection
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
            "broker1:9092,broker2:9092,broker3:9092");
        props.put(ProducerConfig.CLIENT_ID_CONFIG, "order-service-producer-v1.0");

        // Serialization
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

        // Reliability - exactly-once
        props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
        props.put(ProducerConfig.ACKS_CONFIG, "all");
        props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE);
        props.put(ProducerConfig.RETRY_BACKOFF_MS_CONFIG, 100);

        // Performance optimization
        props.put(ProducerConfig.BATCH_SIZE_CONFIG, 32768); // 32KB batches
        props.put(ProducerConfig.LINGER_MS_CONFIG, 20); // 20ms batching
window
        props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "gzip"); // Compression
        props.put(ProducerConfig.BUFFER_MEMORY_CONFIG, 67108864); // 64MB buffer
    }
}

```

```

        // Timeouts
        props.put(ProducerConfig.REQUEST_TIMEOUT_MS_CONFIG, 30000);
        props.put(ProducerConfig.DELIVERY_TIMEOUT_MS_CONFIG, 120000);

        return props;
    }

    public void runHighThroughputTest() {
        int numThreads = 4;
        int messagesPerThread = 1000;
        CountDownLatch latch = new CountDownLatch(numThreads);
        ExecutorService executor = Executors.newFixedThreadPool(numThreads);

        long startTime = System.currentTimeMillis();

        for (int t = 0; t < numThreads; t++) {
            final int threadId = t;
            executor.submit(() -> {
                try {
                    for (int i = 0; i < messagesPerThread; i++) {
                        sendOrderEvent(threadId, i);
                    }
                } finally {
                    latch.countDown();
                }
            });
        }

        try {
            latch.await(); // Wait for all threads to complete
            producer.flush(); // Ensure all messages are sent

            long duration = System.currentTimeMillis() - startTime;

            System.out.printf("Performance Results:%n");
            System.out.printf("Total messages: %d%n", numThreads *
messagesPerThread);
            System.out.printf("Successful: %d%n", successCount.get());
            System.out.printf("Errors: %d%n", errorCount.get());
            System.out.printf("Duration: %d ms%n", duration);
            System.out.printf("Throughput: %.2f messages/sec%n",
                (numThreads * messagesPerThread * 1000.0) / duration);

        } catch (InterruptedException e) {
            Thread.currentThread().interrupt();
        } finally {
            executor.shutdown();
        }
    }

    private void sendOrderEvent(int threadId, int messageId) {
        try {
            // Create order event object
            OrderEvent orderEvent = new OrderEvent(

```

```

        "order-" + threadId + "-" + messageId,
        "user-" + (messageId % 100),
        100.0 + (messageId % 500),
        System.currentTimeMillis()
    );

    String key = orderEvent.getUserId(); // Partition by user
    String value = objectMapper.writeValueAsString(orderEvent);

    ProducerRecord<String, String> record =
        new ProducerRecord<>("order-events", key, value);

    // Add headers for tracing/monitoring
    record.headers().add("source-service", "order-service".getBytes());
    record.headers().add("thread-id",
String.valueOf(threadId).getBytes());

    // Send asynchronously with callback
    producer.send(record, new Callback() {
        @Override
        public void onCompletion(RecordMetadata metadata, Exception
exception) {
            if (exception == null) {
                successCount.incrementAndGet();
                if (messageId % 100 == 0) {
                    System.out.printf("Thread %d sent message %d to
partition %d%n",
                                threadId, messageId, metadata.partition());
                }
            } else {
                errorCount.incrementAndGet();
                System.err.printf("Thread %d failed to send message %d:
%s%n",
                                threadId, messageId, exception.getMessage());

                // Handle different exception types
                handleProducerException(exception, orderEvent);
            }
        }
    });

} catch (JsonProcessingException e) {
    System.err.println("JSON serialization error: " + e.getMessage());
    errorCount.incrementAndGet();
}

}

private void handleProducerException(Exception exception, OrderEvent
orderEvent) {
    if (exception instanceof
org.apache.kafka.common.errors.RetriableException) {
        // Will be retried automatically by producer
        System.out.println("Retriable error - producer will retry: " +
exception.getMessage());
    }
}

```

```

        } else if (exception instanceof
org.apache.kafka.common.errors.RecordTooLargeException) {
            // Message too large - need to handle specially
            System.err.println("Message too large, sending to DLQ: " +
orderEvent.getOrderId());
            // Send to dead letter queue or break into smaller messages
        } else {
            // Other non-retriable errors
            System.err.println("Non-retriable error: " + exception.getMessage());
            // Could send to DLQ or alert monitoring system
        }
    }

    public void close() {
        if (producer != null) {
            producer.close();
        }
    }

    // Data class for order events
    static class OrderEvent {
        private String orderId;
        private String userId;
        private double amount;
        private long timestamp;

        public OrderEvent(String orderId, String userId, double amount, long
timestamp) {
            this.orderId = orderId;
            this.userId = userId;
            this.amount = amount;
            this.timestamp = timestamp;
        }

        // Getters for JSON serialization
        public String getOrderId() { return orderId; }
        public String getUserId() { return userId; }
        public double getAmount() { return amount; }
        public long getTimestamp() { return timestamp; }
    }
}

```

Producer with Custom Serializer

```

import org.apache.kafka.common.serialization.Serializer;
import com.fasterxml.jackson.databind.ObjectMapper;
import java.util.Map;

// Custom JSON serializer for complex objects
public class JsonSerializer<T> implements Serializer<T> {
    private final ObjectMapper objectMapper = new ObjectMapper();
}

```

```

@Override
public void configure(Map<String, ?> configs, boolean isKey) {
    // Configuration if needed
}

@Override
public byte[] serialize(String topic, T data) {
    try {
        return objectMapper.writeValueAsBytes(data);
    } catch (Exception e) {
        throw new RuntimeException("Error serializing JSON", e);
    }
}

@Override
public void close() {
    // Cleanup if needed
}
}

// Usage with custom serializer
public class CustomSerializerExample {

    public static void main(String[] args) {
        Properties props = new Properties();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
JsonSerializer.class);

        try (KafkaProducer<String, OrderEvent> producer = new KafkaProducer<>
(props)) {
            OrderEvent order = new OrderEvent("order-123", "user-456", 999.99,
System.currentTimeMillis());

            ProducerRecord<String, OrderEvent> record =
                new ProducerRecord<>("orders", order.getUserId(), order);

            producer.send(record, (metadata, exception) -> {
                if (exception == null) {
                    System.out.println("Custom serialized object sent
successfully");
                } else {
                    System.err.println("Failed to send: " +
exception.getMessage());
                }
            });
        }
    }
}

```

Comparisons & Trade-offs

Delivery Semantics Comparison

| Delivery Semantic | Configuration | Guarantees | Performance | Use Cases |
|----------------------|---|--|------------------------|----------------------------------|
| At-Most-Once | <code>acks=0, retries=0</code> | Messages may be lost, never duplicated | ★★★ Highest throughput | Metrics, logs, non-critical data |
| At-Least-Once | <code>acks=all, retries>0</code> , no idempotence | No data loss, duplicates possible | ★★ Good performance | Event streaming, analytics |
| Exactly-Once | <code>enable.idempotence=true</code> or transactional | No loss, no duplicates | ★ Higher latency | Financial systems, critical data |

Producer Configuration Trade-offs

| Parameter | Higher Value Impact | Lower Value Impact | Recommendation |
|-------------------------------|-----------------------------------|----------------------------|---|
| <code>batch.size</code> | ↑ Throughput, ↑ Latency, ↑ Memory | ↓ Latency, ↓ Throughput | 16KB-32KB for most use cases |
| <code>linger.ms</code> | ↑ Batching, ↑ Latency | ↓ Latency, ↓ Batching | 5-20ms for balanced performance |
| <code>acks</code> | ↑ Durability, ↓ Throughput | ↓ Durability, ↑ Throughput | <code>all</code> for production systems |
| <code>compression.type</code> | ↓ Bandwidth, ↑ CPU | ↑ Bandwidth, ↓ CPU | <code>gzip</code> for bandwidth, <code>lz4</code> for speed |

Partitioning Strategy Comparison

| Strategy | Advantages | Disadvantages | When to Use |
|--------------------|---------------------------------|------------------------------------|-------------------------------------|
| Key-based | Data locality, ordering per key | Uneven distribution if keys skewed | User events, entity updates |
| Round-robin | Even distribution | No data locality | Logs, metrics without keys |
| Custom | Business-specific optimization | Implementation complexity | Geographic routing, priority queues |

Common Pitfalls & Best Practices

1. Configuration Mistakes

✗ Blocking the Application Thread

```
// DON'T - Synchronous sends hurt performance
for (int i = 0; i < 10000; i++) {
    try {
        producer.send(record).get(); // Blocks on every send!
    } catch (Exception e) {
        // Handle error
    }
}
```

```
// DO - Use async sends with callbacks
for (int i = 0; i < 10000; i++) {
    producer.send(record, (metadata, exception) -> {
        if (exception != null) {
            // Handle error asynchronously
            handleError(exception);
        }
    });
}
producer.flush(); // Wait for all sends to complete
```

✗ Ignoring Producer Exceptions

```
// DON'T - Fire and forget without error handling
producer.send(record);
```

```
// DO - Always handle exceptions
producer.send(record, new Callback() {
    @Override
    public void onCompletion(RecordMetadata metadata, Exception exception) {
        if (exception != null) {
            if (exception instanceof RetriableException) {
                // Will be retried automatically
                logger.warn("Retriable error: {}", exception.getMessage());
            } else {
                // Non-retriable - handle immediately
                logger.error("Failed to send message", exception);
                sendToDeadLetterQueue(record);
            }
        }
    }
});
```


2. Memory and Resource Issues

✗ Unbounded Producer Memory Usage

```
// DON'T - Can cause OutOfMemoryError
Properties props = new Properties();
props.put(ProducerConfig.BUFFER_MEMORY_CONFIG, Long.MAX_VALUE); // Dangerous!
props.put(ProducerConfig.BATCH_SIZE_CONFIG, 1048576); // 1MB batches
props.put(ProducerConfig.LINGER_MS_CONFIG, 60000); // Wait 60 seconds
```

```
// DO - Set reasonable limits
Properties props = new Properties();
props.put(ProducerConfig.BUFFER_MEMORY_CONFIG, 67108864); // 64MB limit
props.put(ProducerConfig.BATCH_SIZE_CONFIG, 16384); // 16KB batches
props.put(ProducerConfig.LINGER_MS_CONFIG, 100); // 100ms max wait
props.put(ProducerConfig.MAX_BLOCK_MS_CONFIG, 5000); // Block max 5s when buffer
full
```

✗ Producer Resource Leaks

```
// DON'T - Forget to close producer
public void sendMessage(String message) {
    KafkaProducer<String, String> producer = new KafkaProducer<>(props);
    producer.send(new ProducerRecord<>("topic", message));
    // Producer not closed - resource leak!
}
```

```
// DO - Use try-with-resources or proper cleanup
private final KafkaProducer<String, String> producer; // Shared instance

public void sendMessage(String message) {
    producer.send(new ProducerRecord<>("topic", message));
}

@PreDestroy
public void cleanup() {
    if (producer != null) {
        producer.close(Duration.ofSeconds(10)); // Graceful close
    }
}
```

3. Ordering and Idempotence Issues

✗ Losing Message Order on Retries

```
// DON'T - Can reorder messages on retry
Properties props = new Properties();
props.put(ProducerConfig.MAX_IN_FLIGHT_REQUESTS_PER_CONNECTION, 10);
props.put(ProducerConfig.RETRIES_CONFIG, 3);
props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, false);
```

```
// DO - Enable idempotence for ordering guarantees
Properties props = new Properties();
props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
// Automatically sets: acks=all, retries=MAX_VALUE, max.in.flight.requests=5
```

4. Performance Anti-patterns

✗ Creating New Producer Per Message

```
// DON'T - Very expensive
public void sendMessage(String message) {
    Properties props = createProducerConfig();
    try (KafkaProducer<String, String> producer = new KafkaProducer<>(props)) {
        producer.send(new ProducerRecord<>("topic", message));
    } // Producer closed after every message!
}
```

```
// DO - Reuse producer instances
@Component
public class MessageService {
    private final KafkaProducer<String, String> producer;

    public MessageService() {
        this.producer = new KafkaProducer<>(createProducerConfig());
    }

    public void sendMessage(String message) {
        producer.send(new ProducerRecord<>("topic", message));
    }

    @PreDestroy
    public void close() {
        producer.close();
    }
}
```

Best Practices Summary

☑ Producer Configuration Best Practices

1. **Enable idempotence** by default (`enable.idempotence=true`)
2. **Use appropriate acks** setting (`all` for production)
3. **Set reasonable timeouts** (`request.timeout.ms=30000`)
4. **Configure proper batching** (`batch.size=16384`, `linger.ms=5-20`)
5. **Use compression** (`compression.type=gzip` or `lz4`)

☑ Application Best Practices

1. **Reuse producer instances** - they're thread-safe
2. **Handle exceptions properly** - distinguish retrieable vs non-retrieable
3. **Use async sends** - don't block application threads
4. **Monitor metrics** - track success rate, latency, errors
5. **Implement circuit breakers** - handle prolonged failures gracefully

☑ Operational Best Practices

1. **Monitor producer metrics** - `record-send-rate`, `record-error-rate`
2. **Set up alerting** - on error rates, buffer exhaustion
3. **Plan for backpressure** - handle when Kafka is slow/down
4. **Test failure scenarios** - broker failures, network partitions
5. **Capacity planning** - monitor throughput trends



Real-World Use Cases

1. E-commerce Order Processing

```
@Service
public class OrderService {
    private final KafkaProducer<String, OrderEvent> producer;

    @Transactional
    public void processOrder(Order order) {
        // Save to database
        orderRepository.save(order);

        // Publish event for downstream services
        OrderEvent event = new OrderEvent(order.getId(), order.getUserId(),
            order.getAmount(), OrderStatus.CREATED);

        ProducerRecord<String, OrderEvent> record =
            new ProducerRecord<>("order-events", order.getUserId(), event);

        producer.send(record, (metadata, exception) -> {
            if (exception != null) {
                // Compensating action or alert
            }
        })
    }
}
```

```

        logger.error("Failed to publish order event", exception);
        orderEventFailureHandler.handle(order, exception);
    }
});
}
}

```

Why Kafka for Orders:

- **Decoupling:** Order service doesn't depend on inventory, payment services
- **Reliability:** Events are persisted and can be replayed
- **Scalability:** Multiple consumers can handle different aspects

2. Real-time Analytics and Monitoring

```

@Component
public class MetricsProducer {
    private final KafkaProducer<String, MetricEvent> producer;

    @Scheduled(fixedDelay = 1000) // Every second
    public void publishSystemMetrics() {
        // CPU usage
        double cpuUsage = systemMonitor.getCpuUsage();
        publishMetric("system.cpu.usage", cpuUsage, "percentage");

        // Memory usage
        long memoryUsed = systemMonitor.getMemoryUsed();
        publishMetric("system.memory.used", memoryUsed, "bytes");

        // Request rate
        double requestRate = applicationMonitor.getRequestRate();
        publishMetric("app.requests.rate", requestRate, "requests/sec");
    }

    private void publishMetric(String metricName, double value, String unit) {
        MetricEvent event = new MetricEvent(metricName, value, unit,
            System.currentTimeMillis(), "hostname");

        // Partition by metric name for consistent routing
        ProducerRecord<String, MetricEvent> record =
            new ProducerRecord<>("system-metrics", metricName, event);

        producer.send(record); // Fire and forget for metrics
    }
}

```

Configuration for High-Volume Metrics:

```
Properties metricsProducerConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");

    // Optimize for throughput over reliability
    props.put(ProducerConfig.ACKS_CONFIG, "1"); // Leader ack only
    props.put(ProducerConfig.BATCH_SIZE_CONFIG, 65536); // 64KB batches
    props.put(ProducerConfig.LINGER_MS_CONFIG, 100); // Wait 100ms for batch
    props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "lz4"); // Fast compression

    return props;
}
```

3. Financial Transaction Processing

```
@Service
public class PaymentProcessor {
    private final KafkaProducer<String, PaymentEvent> producer;

    @Transactional
    public void processPayment(PaymentRequest request) {
        try {
            // Initialize transaction
            producer.initTransactions();
            producer.beginTransaction();

            // Validate payment
            PaymentValidation validation = paymentValidator.validate(request);
            if (!validation.isValid()) {
                throw new PaymentValidationException(validation.getErrors());
            }

            // Process payment with external service
            PaymentResult result = externalPaymentService.processPayment(request);

            // Record events atomically
            PaymentEvent event = new PaymentEvent(request.getPaymentId(),
                result.getStatus(), request.getAmount(),
                System.currentTimeMillis());

            producer.send(new ProducerRecord<>("payment-events",
                request.getPaymentId(), event));

            // Update account balance event
            AccountEvent accountEvent = new AccountEvent(request.getAccountId(),
                AccountEventType.DEBIT, request.getAmount());

            producer.send(new ProducerRecord<>("account-events",
                request.getAccountId(), accountEvent));
        }
    }
}
```

```

        // Commit transaction
        producer.commitTransaction();

    } catch (Exception e) {
        producer.abortTransaction();
        throw new PaymentProcessingException("Payment failed", e);
    }
}
}

```

Transactional Configuration:

```

Properties transactionalConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "payment-processor-1");
    // Exactly-once semantics automatically configured
    return props;
}

```

4. IoT Sensor Data Streaming

```

@Service
public class IoTDataProducer {
    private final KafkaProducer<String, SensorReading> producer;

    public void publishSensorReading(String sensorId, SensorReading reading) {
        // Custom partitioner routes by geographic region
        ProducerRecord<String, SensorReading> record =
            new ProducerRecord<>("sensor-data", sensorId, reading);

        // Add metadata headers
        record.headers().add("sensor-type", reading.getSensorType().getBytes());
        record.headers().add("location", reading.getLocation().getBytes());
        record.headers().add("timestamp",
            String.valueOf(reading.getTimestamp()).getBytes());

        producer.send(record, (metadata, exception) -> {
            if (exception != null) {
                // Log error but don't fail - sensor data is high volume
                logger.warn("Failed to send sensor reading for {}: {}",
                    sensorId, exception.getMessage());

                // Could implement local buffering or retry queue
                sensorFailureBuffer.add(record);
            }
        });
    }
}

```

5. Log Aggregation from Microservices

```
@Component
public class StructuredLogger {
    private final KafkaProducer<String, LogEvent> producer;

    public void logEvent(String service, LogLevel level, String message,
                        Map<String, Object> context) {
        LogEvent event = LogEvent.builder()
            .service(service)
            .level(level)
            .message(message)
            .context(context)
            .timestamp(System.currentTimeMillis())
            .hostname(getHostname())
            .build();

        // Partition by service for consistent routing
        ProducerRecord<String, LogEvent> record =
            new ProducerRecord<>("application-logs", service, event);

        producer.send(record); // Async - don't slow down application
    }
}

// Usage in application code
@RestController
public class UserController {

    @Autowired
    private StructuredLogger logger;

    @PostMapping("/users")
    public ResponseEntity<User> createUser(@RequestBody CreateUserRequest request)
    {
        Map<String, Object> context = Map.of(
            "userId", request.getUserId(),
            "email", request.getEmail(),
            "source", "web-app"
        );





        logger.logEvent("user-service", LogLevel.INFO,
            "Creating new user", context);

        // ... business logic ...

        return ResponseEntity.ok(user);
    }
}
```

Version Highlights

Kafka 4.0 (September 2025) - Current Latest

-  **Default KRaft Mode:** ZooKeeper completely removed
-  **New Consumer Protocol (KIP-848):** Faster rebalancing, better partition assignment
-  **Enhanced Producer Batching:** Improved sticky partitioner performance
-  **Java 17+ Required:** For brokers (Java 11+ for clients)

Kafka 3.x Series Producer Features

- **3.6** (Oct 2023): Producer metrics improvements, better error reporting
- **3.5** (Jun 2023): Enhanced idempotent producer stability
- **3.4** (Feb 2023): Improved transactional producer performance
- **3.3** (Oct 2022): Producer client optimizations for KRaft
- **3.2** (May 2022): Better producer error handling and retry logic
- **3.1** (Jan 2022): Producer connection pooling improvements
- **3.0** (Sep 2021): **Idempotence enabled by default**, improved batching

Key Producer Evolution

| Version | Producer Feature | Impact |
|---------|---------------------------------|-----------------------------------|
| 4.0 | Enhanced batching algorithms | Better throughput |
| 3.0 | Idempotence by default | Exactly-once semantics out-of-box |
| 2.8 | Producer improvements for KRaft | Better metadata handling |
| 2.4 | Sticky partitioner | Better batching for null keys |
| 2.1 | Zstandard compression | Better compression ratios |
| 1.1 | Headers support | Metadata in messages |
| 0.11 | Idempotent producer | Exactly-once semantics |
| 0.10 | Message timestamps | Time-based operations |
| 0.9 | Producer rewrite | Much better performance |

Current Recommendations (2025)

```
// Modern Kafka 4.0 producer configuration
public static Properties modernProducerConfig() {
    Properties props = new Properties();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG, StringSerializer.class);
    props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);

    // Kafka 4.0 best practices
    props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true); // Default since
```



```
3.0
    props.put(ProducerConfig.ACKS_CONFIG, "all");           // Strong
durability
    props.put(ProducerConfig.COMPRESSION_TYPE_CONFIG, "zstd"); // Best
compression (2.1+)

    // Optimized batching (improvements in 4.0)
    props.put(ProducerConfig.BATCH_SIZE_CONFIG, 32768);     // 32KB
    props.put(ProducerConfig.LINGER_MS_CONFIG, 10);         // 10ms

    return props;
}
```

Additional Resources

Official Documentation

- [Kafka Producer API Documentation](#)
- [Producer Configuration Reference](#)
- [Kafka Improvement Proposals - Producers](#)

Learning Resources

- [Confluent Producer Tutorial](#)
- [Apache Kafka: Producer Deep Dive](#)
- [Kafka Producer Best Practices](#)

Tools & Monitoring

- [JMX Metrics for Producers](#)
- [Kafka Producer Performance Testing](#)


Troubleshooting

- [Producer Error Handling](#)
- [Common Producer Issues](#)

Last Updated: September 2025

Kafka Version: 4.0.0

Java Compatibility: 11+ (clients), 17+ (recommended)

 **Pro Tip:** Start with idempotence enabled and `acks=all` for reliability, then optimize for performance based on your specific use case requirements.