Kafka Consumers: Complete Developer Guide

A comprehensive refresher on Apache Kafka Consumers, designed for both beginners and experienced developers. This README covers consumer basics, consumer groups, offset management, and real-world applications with detailed Java examples.

Table of Contents

- Consumer API
 - KafkaConsumer Basics
 - Poll Loop Implementation
 - Auto vs Manual Commit
- **11** Consumer Groups
 - Rebalancing Strategies
 - Static Membership
 - Group Coordination
- III Offset Management
 - _consumer_offsets Topic
 - Reset Policies
 - Lag Monitoring
- Comprehensive Java Examples
- Kommon Pitfalls & Best Practices
- Real-World Use Cases
- Wersion Highlights
- Additional Resources

Consumer API

KafkaConsumer Basics

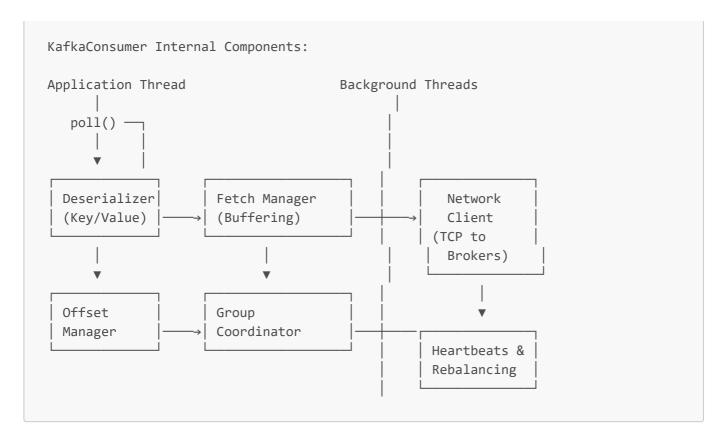
Simple Explanation

The KafkaConsumer is the client API that applications use to read records from Kafka topics. It provides high-level abstractions for subscribing to topics, managing offsets, and handling rebalancing in consumer groups.

Problem It Solves

- Scalable Data Consumption: Multiple consumers can process data in parallel
- Fault Tolerance: Automatic failover and recovery
- Order Preservation: Maintains order within partitions
- Offset Management: Tracks what messages have been processed

Internal Architecture



Basic Consumer Configuration

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.serialization.StringDeserializer;
import java.util.Properties;
public class BasicConsumerConfig {
    public static Properties createBasicConsumerConfig() {
        Properties props = new Properties();
        // Connection settings
        props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP ID CONFIG, "my-consumer-group");
        props.put(ConsumerConfig.CLIENT_ID_CONFIG, "my-consumer-client");
        // Deserialization
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
            StringDeserializer.class.getName());
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
            StringDeserializer.class.getName());
        // Offset behavior
        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
        props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, true);
        props.put(ConsumerConfig.AUTO COMMIT INTERVAL MS CONFIG, 1000);
        // Performance settings
        props.put(ConsumerConfig.FETCH MIN BYTES CONFIG, 1024);
                                                                          // Min
1KB per fetch
```

Poll Loop Implementation

Basic Poll Loop Pattern

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.TopicPartition;
import java.time.Duration;
import java.util.Arrays;
import java.util.Properties;
public class BasicConsumerExample {
   private static final String TOPIC = "user-events";
   private volatile boolean running = true;
   public static void main(String[] args) {
        BasicConsumerExample consumer = new BasicConsumerExample();
        consumer.consume();
   }
   public void consume() {
        Properties props = BasicConsumerConfig.createBasicConsumerConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            // Subscribe to topic(s)
            consumer.subscribe(Arrays.asList(TOPIC));
            // Main consumption loop
            while (running) {
                try {
                    // Poll for records with timeout
                    ConsumerRecords<String, String> records =
                        consumer.poll(Duration.ofMillis(1000));
                    // Process each record
                    for (ConsumerRecord<String, String> record : records) {
                        processRecord(record);
                    }
                    // Commit offsets (if manual commit is enabled)
```

```
// consumer.commitSync();
                } catch (Exception e) {
                    System.err.println("Error during consumption: " +
e.getMessage());
                    // Handle error - might want to continue or break based on
error type
                }
            }
        } catch (Exception e) {
            System.err.println("Consumer failed: " + e.getMessage());
        }
    }
    private void processRecord(ConsumerRecord<String, String> record) {
        System.out.printf("Consumed record: topic=%s, partition=%d, offset=%d, " +
            "key=%s, value=%s, timestamp=%d%n",
            record.topic(), record.partition(), record.offset(),
            record.key(), record.value(), record.timestamp());
        // Your business logic here
        try {
            // Simulate processing time
            Thread.sleep(10);
            // Process the record
            handleUserEvent(record.key(), record.value());
        } catch (InterruptedException e) {
            Thread.currentThread().interrupt();
            running = false;
        } catch (Exception e) {
            System.err.println("Failed to process record: " + e.getMessage());
            // Decide whether to skip, retry, or send to DLQ
        }
    }
    private void handleUserEvent(String userId, String eventData) {
        // Business logic implementation
        System.out.println("Processing event for user: " + userId);
    }
    public void shutdown() {
        running = false;
    }
}
```

Advanced Poll Loop with Error Handling

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.errors.*;
import java.time.Duration;
import java.util.Arrays;
import java.util.concurrent.atomic.AtomicBoolean;
public class RobustConsumerExample {
    private final AtomicBoolean running = new AtomicBoolean(true);
    private final String topic;
    private final Properties config;
    public RobustConsumerExample(String topic, Properties config) {
        this.topic = topic;
        this.config = config;
    }
    public void consume() {
        while (running.get()) {
            try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>
(config)) {
                consumer.subscribe(Arrays.asList(topic), new RebalanceListener());
                while (running.get()) {
                    try {
                        ConsumerRecords<String, String> records =
                            consumer.poll(Duration.ofMillis(1000));
                        if (records.isEmpty()) {
                            continue;
                        }
                        processRecords(records);
                        commitOffsets(consumer, records);
                    } catch (WakeupException e) {
                        // Shutdown signal
                        System.out.println("Consumer wakeup called");
                        break;
                    } catch (CommitFailedException e) {
                        // Commit failed - usually during rebalance
                        System.err.println("Commit failed: " + e.getMessage());
                        // Continue - offsets will be reset by rebalance
                    } catch (AuthorizationException e) {
                        // Not authorized - fatal error
                        System.err.println("Not authorized: " + e.getMessage());
                        break;
                    } catch (InvalidOffsetException e) {
                        // Invalid offset - reset
```

```
System.err.println("Invalid offset: " + e.getMessage());
                        consumer.seekToBeginning(e.partitions());
                    } catch (RetriableException e) {
                        // Retriable error - log and continue
                        System.err.println("Retriable error: " + e.getMessage());
                    } catch (Exception e) {
                        // Unexpected error
                        System.err.println("Unexpected error: " + e.getMessage());
                        e.printStackTrace();
                        break;
                    }
                }
            } catch (Exception e) {
                System.err.println("Consumer initialization failed: " +
e.getMessage());
                if (running.get()) {
                    // Wait before retry
                    try {
                        Thread.sleep(5000);
                    } catch (InterruptedException ie) {
                        Thread.currentThread().interrupt();
                        break;
                    }
                }
            }
        }
        System.out.println("Consumer shut down");
   }
   private void processRecords(ConsumerRecords<String, String> records) {
        for (ConsumerRecord<String, String> record : records) {
            try {
                // Process individual record
                processRecord(record);
            } catch (Exception e) {
                System.err.printf("Failed to process record at offset %d: %s%n",
                    record.offset(), e.getMessage());
                // Decide on error handling strategy
                handleProcessingError(record, e);
            }
        }
   }
   private void processRecord(ConsumerRecord<String, String> record) {
        // Your business logic here
        System.out.printf("Processing: %s:%s%n", record.key(), record.value());
```

```
private void handleProcessingError(ConsumerRecord<String, String> record,
Exception e) {
        // Error handling strategies:
        if (e instanceof IllegalArgumentException) {
            // Invalid data - skip record
            System.out.println("Skipping invalid record");
        } else if (e instanceof java.net.ConnectException) {
            // External service unavailable - could retry or send to DLQ
            System.out.println("External service unavailable - sending to DLQ");
            sendToDeadLetterQueue(record);
        } else {
            // Unknown error - could fail fast or continue
            System.err.println("Unknown processing error - continuing");
        }
    }
    private void sendToDeadLetterQueue(ConsumerRecord<String, String> record) {
        // Implementation to send failed records to DLQ
        System.out.println("Sending to DLQ: " + record.key());
    }
    private void commitOffsets(KafkaConsumer<String, String> consumer,
                              ConsumerRecords<String, String> records) {
        try {
            consumer.commitSync();
        } catch (CommitFailedException e) {
            System.err.println("Commit failed: " + e.getMessage());
        }
    }
    public void shutdown() {
        running.set(false);
    }
    // Rebalance listener for cleanup
    private class RebalanceListener implements ConsumerRebalanceListener {
        @Override
        public void
onPartitionsRevoked(java.util.Collection<org.apache.kafka.common.TopicPartition>
partitions) {
            System.out.println("Partitions revoked: " + partitions);
            // Commit current offsets before losing partitions
        }
        @Override
        public void
onPartitionsAssigned(java.util.Collection<org.apache.kafka.common.TopicPartition>
partitions) {
            System.out.println("Partitions assigned: " + partitions);
            // Initialize any partition-specific state
```

```
}
}
}
```

Auto vs Manual Commit

Auto Commit (Default)

```
public class AutoCommitConsumerExample {
    public static Properties createAutoCommitConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "auto-commit-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        // Auto commit settings (default)
        props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, true);
        props.put(ConsumerConfig.AUTO_COMMIT_INTERVAL_MS_CONFIG, 5000); // Every 5
seconds
        return props;
    }
    public static void main(String[] args) {
        Properties props = createAutoCommitConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            consumer.subscribe(Arrays.asList("orders"));
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, String> record : records) {
                    // Process record
                    processOrder(record.value());
                    // Offsets committed automatically every 5 seconds
                    // No manual commit needed
                }
                // Note: If processing fails after auto-commit,
                // messages might be lost on restart
            }
        }
```

```
private static void processOrder(String orderData) {
    System.out.println("Processing order: " + orderData);
    // Business logic here
}
```

Auto Commit Trade-offs:

- Simplicity: No manual offset management
- Performance: No commit overhead in processing loop
- **X** At-most-once: Messages can be lost if processing fails after commit
- X Less Control: Cannot commit based on processing success

Manual Commit (Recommended for Critical Data)

```
public class ManualCommitConsumerExample {
    public static Properties createManualCommitConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP ID CONFIG, "manual-commit-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        // Manual commit settings
        props.put(ConsumerConfig.ENABLE AUTO COMMIT CONFIG, false); // Disable
auto commit
        // Reduce fetch size for better control
        props.put(ConsumerConfig.MAX POLL RECORDS CONFIG, 100);
        return props;
    }
    public static void main(String[] args) {
        Properties props = createManualCommitConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            consumer.subscribe(Arrays.asList("financial-transactions"));
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                if (records.isEmpty()) {
                    continue;
                }
```

```
// Process all records in batch
                boolean allProcessedSuccessfully = true;
                for (ConsumerRecord<String, String> record : records) {
                        processFinancialTransaction(record);
                    } catch (Exception e) {
                        System.err.println("Processing failed for record: " +
record.offset());
                        allProcessedSuccessfully = false;
                        break; // Stop processing this batch
                    }
                }
                // Only commit if all records processed successfully
                if (allProcessedSuccessfully) {
                    try {
                        // Synchronous commit - blocks until complete
                        consumer.commitSync();
                        System.out.println("Batch committed successfully");
                    } catch (CommitFailedException e) {
                        System.err.println("Commit failed: " + e.getMessage());
                        // Handle commit failure - maybe reprocess
                    }
                } else {
                    System.out.println("Skipping commit due to processing
failures");
                    // Could implement retry logic here
                }
            }
        }
    }
    private static void processFinancialTransaction(ConsumerRecord<String, String>
record) {
        System.out.printf("Processing transaction: %s at offset %d%n",
            record.value(), record.offset());
        // Simulate processing that could fail
        if (record.value().contains("invalid")) {
            throw new IllegalArgumentException("Invalid transaction data");
        }
        // Actual business logic here
        // Save to database, call external service, etc.
   }
}
```

Asynchronous Manual Commit

```
public class AsyncCommitConsumerExample {
    public static void main(String[] args) {
        Properties props = ManualCommitConsumerExample.createManualCommitConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            consumer.subscribe(Arrays.asList("user-events"));
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, String> record : records) {
                    processUserEvent(record);
                }
                // Asynchronous commit - doesn't block
                consumer.commitAsync(new OffsetCommitCallback() {
                    @Override
                    public void
onComplete(java.util.Map<org.apache.kafka.common.TopicPartition,</pre>
org.apache.kafka.clients.consumer.OffsetAndMetadata> offsets,
                                          Exception exception) {
                        if (exception != null) {
                            System.err.println("Commit failed: " +
exception.getMessage());
                            // Could implement retry logic
                        } else {
                            System.out.println("Async commit successful");
                    }
                });
            }
        }
    }
    private static void processUserEvent(ConsumerRecord<String, String> record) {
        System.out.println("Processing user event: " + record.value());
    }
}
```

Fine-grained Offset Control

```
import java.util.HashMap;
import java.util.Map;

public class FinegrainedOffsetConsumerExample {
```

```
public static void main(String[] args) {
        Properties props = ManualCommitConsumerExample.createManualCommitConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            consumer.subscribe(Arrays.asList("orders"));
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                // Track offsets per partition
                Map<TopicPartition, OffsetAndMetadata> currentOffsets = new
HashMap<>();
                for (ConsumerRecord<String, String> record : records) {
                    try {
                        processOrder(record);
                        // Track the offset for this partition
                        currentOffsets.put(
                            new TopicPartition(record.topic(),
record.partition()),
                            new OffsetAndMetadata(record.offset() + 1, "Processed
successfully")
                        );
                    } catch (Exception e) {
                        System.err.printf("Failed to process record at offset %d:
%s%n",
                            record.offset(), e.getMessage());
                        // Don't update offset for failed record
                        break;
                    }
                }
                // Commit specific offsets
                if (!currentOffsets.isEmpty()) {
                    try {
                        consumer.commitSync(currentOffsets);
                        System.out.println("Committed offsets: " +
currentOffsets);
                    } catch (CommitFailedException e) {
                        System.err.println("Offset commit failed: " +
e.getMessage());
                    }
                }
            }
        }
    }
    private static void processOrder(ConsumerRecord<String, String> record) {
```

£ Consumer Groups

Rebalancing Strategies

Range Assignor (Default)

```
public class RangeAssignorExample {
    public static Properties createRangeAssignorConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "range-assignor-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        // Range assignor (default)
        props.put(ConsumerConfig.PARTITION_ASSIGNMENT_STRATEGY_CONFIG,
            "org.apache.kafka.clients.consumer.RangeAssignor");
        return props;
    }
    * Range Assignor Distribution:
     * Topic: orders (6 partitions: 0,1,2,3,4,5)
     * 3 consumers: C1, C2, C3
     * Assignment:
     * C1: partitions 0,1 (6/3 = 2)
     * C2: partitions 2,3 (6/3 = 2)
     * C3: partitions 4,5 (6/3 = 2)
     * If 4 consumers:
     * C1: partitions 0,1 (first 6%4=2 consumers get extra)
     * C2: partitions 2,3
     * C3: partition 4
     * C4: partition 5
     */
}
```

```
public class RoundRobinAssignorExample {
    public static Properties createRoundRobinAssignorConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "roundrobin-assignor-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        // Round robin assignor
        props.put(ConsumerConfig.PARTITION ASSIGNMENT STRATEGY CONFIG,
            "org.apache.kafka.clients.consumer.RoundRobinAssignor");
        return props;
    }
    /*
     * Round Robin Assignor Distribution:
     * Topic: orders (6 partitions: 0,1,2,3,4,5)
     * 3 consumers: C1, C2, C3
     * Assignment (round-robin):
     * C1: partitions 0,3
     * C2: partitions 1,4
     * C3: partitions 2,5
     * Better distribution across multiple topics
}
```

Sticky Assignor (Recommended)

```
props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, 30000);
props.put(ConsumerConfig.HEARTBEAT_INTERVAL_MS_CONFIG, 3000);

return props;
}

/*
   * Sticky Assignor Benefits:
   * 1. Minimizes partition reassignment during rebalance
   * 2. Maintains balanced distribution
   * 3. Reduces consumer startup time after rebalance
   *
   * Example:
   * Initial: C1[0,1], C2[2,3], C3[4,5]
   * C2 leaves: C1[0,1,2], C3[4,5,3] (C1 and C3 keep their partitions)
   * vs Round Robin: C1[0,2,4], C3[1,3,5] (all partitions reassigned)
   */
}
```

Cooperative Sticky Assignor (Kafka 2.4+)

```
public class CooperativeStickyAssignorExample {
    public static Properties createCooperativeStickyConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP ID CONFIG, "cooperative-sticky-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE DESERIALIZER CLASS CONFIG,
StringDeserializer.class);
        // Cooperative sticky assignor - incremental rebalancing
        props.put(ConsumerConfig.PARTITION ASSIGNMENT STRATEGY CONFIG,
            "org.apache.kafka.clients.consumer.CooperativeStickyAssignor");
        // Optimize for cooperative rebalancing
        props.put(ConsumerConfig.SESSION TIMEOUT MS CONFIG, 30000);
        props.put(ConsumerConfig.HEARTBEAT INTERVAL MS CONFIG, 3000);
        props.put(ConsumerConfig.MAX POLL INTERVAL MS CONFIG, 300000);
        return props;
    }
    public static void main(String[] args) {
        Properties props = createCooperativeStickyConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            // Rebalance listener to observe cooperative rebalancing
```

```
consumer.subscribe(Arrays.asList("orders"), new
ConsumerRebalanceListener() {
                @Override
                public void
onPartitionsRevoked(java.util.Collection<TopicPartition> partitions) {
                    System.out.println("Partitions revoked (cooperative): " +
partitions);
                    // With cooperative rebalancing, this may be called multiple
times
                    // Only revoked partitions are stopped, others continue
processing
                }
                @Override
                public void
onPartitionsAssigned(java.util.Collection<TopicPartition> partitions) {
                    System.out.println("Partitions assigned (cooperative): " +
partitions);
                    // New partitions can start processing while others continue
                }
                @Override
                public void onPartitionsLost(java.util.Collection<TopicPartition>
partitions) {
                    System.out.println("Partitions lost: " + partitions);
                    // Handle partition loss scenario
                }
            });
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, String> record : records) {
                    System.out.printf("Processing record from partition %d: %s%n",
                        record.partition(), record.value());
                }
                consumer.commitAsync();
            }
        }
    }
     * Cooperative Rebalancing Benefits:
     * 1. No "stop-the-world" rebalancing
     * 2. Unaffected consumers continue processing
     * 3. Faster rebalancing overall
     * 4. Better availability during membership changes
     * Process:
     * 1. First rebalance: Identify partitions to be revoked
     * 2. Only affected consumers stop processing revoked partitions
     * 3. Second rebalance: Reassign revoked partitions
```

```
* 4. Processing resumes with new assignment

*/
}
```

Static Membership

Static Consumer Configuration

```
public class StaticMembershipExample {
    public static Properties createStaticMemberConfig(String memberId) {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "static-membership-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        // Static membership configuration
        props.put(ConsumerConfig.GROUP_INSTANCE_ID_CONFIG, memberId); // Unique
static ID
        // Longer session timeout for static members
        props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, 60000); // 60
seconds
        props.put(ConsumerConfig.HEARTBEAT_INTERVAL_MS_CONFIG, 6000); // 6 seconds
        // Cooperative rebalancing works well with static membership
        props.put(ConsumerConfig.PARTITION ASSIGNMENT STRATEGY CONFIG,
            "org.apache.kafka.clients.consumer.CooperativeStickyAssignor");
        return props;
    }
    public static void main(String[] args) {
        if (args.length != 1) {
            System.err.println("Usage: java StaticMembershipExample <member-id>");
            System.exit(1);
        }
        String memberId = args[0]; // e.g., "consumer-1", "consumer-2"
        Properties props = createStaticMemberConfig(memberId);
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            consumer.subscribe(Arrays.asList("orders"), new
ConsumerRebalanceListener() {
                @Override
                public void
```

```
onPartitionsRevoked(java.util.Collection<TopicPartition> partitions) {
                    System.out.printf("[%s] Partitions revoked: %s%n", memberId,
partitions);
                }
                @Override
                public void
onPartitionsAssigned(java.util.Collection<TopicPartition> partitions) {
                    System.out.printf("[%s] Partitions assigned: %s%n", memberId,
partitions);
                }
            });
            System.out.printf("Started static member: %s%n", memberId);
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, String> record : records) {
                    System.out.printf("[%s] Processing from partition %d: %s%n",
                        memberId, record.partition(), record.value());
                }
                if (!records.isEmpty()) {
                    consumer.commitSync();
                }
            }
        } catch (Exception e) {
            System.err.printf("[%s] Consumer failed: %s%n", memberId,
e.getMessage());
        }
    }
     * Static Membership Benefits:
     * 1. Avoid unnecessary rebalances:
          - Temporary disconnects don't trigger rebalance
          - Restart with same group.instance.id keeps same partitions
     * 2. Faster recovery:
         - No partition reassignment on temporary failures
          - Immediate resumption when consumer reconnects
     * 3. Better for stateful processing:
         - Maintains partition-to-consumer affinity
          - Useful for local state/caches tied to partitions
     * Usage scenarios:
     * - Kafka Streams applications
     * - Consumers with local state tied to partitions
     * - Applications sensitive to rebalance overhead
```

```
*/
}
```

Group Coordination

Custom Rebalance Listener with State Management

```
import java.util.Map;
import java.util.concurrent.ConcurrentHashMap;
public class StatefulConsumerExample {
    private final Map<Integer, String> partitionState = new ConcurrentHashMap<>();
    private final String consumerId;
    public StatefulConsumerExample(String consumerId) {
        this.consumerId = consumerId;
    }
    public void consume() {
        Properties props = createConsumerConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            consumer.subscribe(Arrays.asList("user-events"), new
StateAwareRebalanceListener());
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, String> record : records) {
                    processRecordWithState(record);
                consumer.commitSync();
            }
        }
    }
    private Properties createConsumerConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "stateful-consumer-group");
        props.put(ConsumerConfig.CLIENT ID CONFIG, consumerId);
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);
```

```
// Static membership for better state management
        props.put(ConsumerConfig.GROUP_INSTANCE_ID_CONFIG, consumerId);
        props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, 45000);
        return props;
   }
   private void processRecordWithState(ConsumerRecord<String, String> record) {
        int partition = record.partition();
        String value = record.value();
       // Use partition-local state
        String currentState = partitionState.get(partition);
        String newState = updateState(currentState, value);
        partitionState.put(partition, newState);
        System.out.printf("[%s] Partition %d: %s -> %s%n",
            consumerId, partition, currentState, newState);
   }
   private String updateState(String currentState, String newValue) {
        // State update logic based on new message
        return currentState == null ? newValue : currentState + "," + newValue;
   }
   private class StateAwareRebalanceListener implements ConsumerRebalanceListener
{
       @Override
        public void onPartitionsRevoked(java.util.Collection<TopicPartition>
partitions) {
            System.out.printf("[%s] Partitions revoked: %s%n", consumerId,
partitions);
            // Save partition state before losing partitions
            for (TopicPartition partition : partitions) {
                String state = partitionState.remove(partition.partition());
                if (state != null) {
                    savePartitionState(partition, state);
                }
            }
            // Commit offsets before rebalance
            try {
                // Note: consumer reference not available here in this example
                // In real implementation, you'd pass consumer reference
                System.out.println("Committing offsets before rebalance");
            } catch (Exception e) {
                System.err.println("Failed to commit offsets: " + e.getMessage());
            }
        }
        @Override
```

```
public void onPartitionsAssigned(java.util.Collection<TopicPartition>
            System.out.printf("[%s] Partitions assigned: %s%n", consumerId,
partitions);
            // Restore state for newly assigned partitions
            for (TopicPartition partition : partitions) {
                String state = loadPartitionState(partition);
                if (state != null) {
                    partitionState.put(partition.partition(), state);
                }
            }
        }
       @Override
        public void onPartitionsLost(java.util.Collection<TopicPartition>
partitions) {
            System.out.printf("[%s] Partitions lost: %s%n", consumerId,
partitions);
            // Clean up state for lost partitions
            for (TopicPartition partition : partitions) {
                partitionState.remove(partition.partition());
            }
        }
   }
   private void savePartitionState(TopicPartition partition, String state) {
        // Save state to external store (database, file, etc.)
        System.out.printf("Saving state for %s: %s%n", partition, state);
   }
   private String loadPartitionState(TopicPartition partition) {
        // Load state from external store
        System.out.printf("Loading state for %s%n", partition);
        return null; // Return actual state from storage
   }
   public static void main(String[] args) {
        String consumerId = args.length > 0 ? args[0] : "consumer-1";
        new StatefulConsumerExample(consumerId).consume();
}
```

Offset Management

__consumer_offsets Topic

Understanding the Internal Topic

```
public class ConsumerOffsetsExplorer {
    public static void main(String[] args) {
        // Consumer offsets are stored in __consumer_offsets topic
        System.out.println("Consumer Offsets Topic Structure:");
        System.out.println("Topic: __consumer_offsets");
        System.out.println("Partitions: 50 (default, configurable with
offsets.topic.num.partitions)");
        System.out.println("Replication Factor: 3 (default, configurable with
offsets.topic.replication.factor)");
        System.out.println("Cleanup Policy: compact (to keep latest offsets
only)");
        exploreOffsetStorage();
    }
    private static void exploreOffsetStorage() {
        * Offset Storage Format:
         * Key: [group.id, topic, partition]
         * Value: [offset, metadata, commit_timestamp, expire_timestamp]
         * Example:
         * Key: ["my-consumer-group", "orders", 0]
         * Value: [1234, "processed successfully", 1693875600000, 1693962000000]
         * Partition Assignment:
         * hash(group.id) % offsets.topic.num.partitions
         */
        String groupId = "my-consumer-group";
        int offsetsTopicPartitions = 50; // default
        int partitionForGroup = Math.abs(groupId.hashCode()) %
offsetsTopicPartitions;
        System.out.printf("Group '%s' offsets stored in consumer offsets
partition %d%n",
            groupId, partitionForGroup);
    }
    // Method to programmatically access offset information
    public static void monitorConsumerOffsets(String groupId) {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "offset-monitor-group");
        props.put(ConsumerConfig.KEY DESERIALIZER CLASS CONFIG,
            "org.apache.kafka.common.serialization.ByteArrayDeserializer");
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
            "org.apache.kafka.common.serialization.ByteArrayDeserializer");
        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
        try (KafkaConsumer<byte[], byte[]> consumer = new KafkaConsumer<>(props))
```

```
// Subscribe to consumer offsets topic
            consumer.subscribe(Arrays.asList("__consumer_offsets"));
            System.out.println("Monitoring consumer offsets...");
            while (true) {
                ConsumerRecords<byte[], byte[]> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<byte[], byte[]> record : records) {
                    // Decode offset commit messages
                    // Note: This is complex as it requires understanding Kafka's
internal format
                    // In practice, use AdminClient or JMX metrics instead
                    System.out.printf("Offset commit in partition %d at offset
%d%n",
                        record.partition(), record.offset());
                }
            }
        }
   }
}
```

Programmatic Offset Management

```
import org.apache.kafka.clients.admin.*;
import org.apache.kafka.common.TopicPartition;
import java.util.Arrays;
import java.util.Map;
import java.util.concurrent.ExecutionException;
public class OffsetManagementExample {
    public static void main(String[] args) throws ExecutionException,
InterruptedException {
        Properties adminProps = new Properties();
        adminProps.put(AdminClientConfig.BOOTSTRAP_SERVERS_CONFIG,
"localhost:9092");
        try (AdminClient adminClient = AdminClient.create(adminProps)) {
            String groupId = "my-consumer-group";
            // List consumer groups
            listConsumerGroups(adminClient);
            // Get consumer group details
            describeConsumerGroup(adminClient, groupId);
```

```
// List consumer group offsets
           listConsumerGroupOffsets(adminClient, groupId);
           // Reset consumer group offsets
           // resetConsumerGroupOffsets(adminClient, groupId);
       }
   }
    private static void listConsumerGroups(AdminClient adminClient)
           throws ExecutionException, InterruptedException {
       ListConsumerGroupsResult result = adminClient.listConsumerGroups();
       System.out.println("Consumer Groups:");
        result.all().get().forEach(group -> {
           System.out.printf(" Group ID: %s, State: %s, Protocol: %s%n",
                group.groupId(), group.state().orElse("UNKNOWN"),
                group.protocolType().orElse("UNKNOWN"));
       });
   }
   private static void describeConsumerGroup(AdminClient adminClient, String
groupId)
           throws ExecutionException, InterruptedException {
       DescribeConsumerGroupsResult result =
           adminClient.describeConsumerGroups(Arrays.asList(groupId));
       ConsumerGroupDescription description = result.all().get().get(groupId);
       System.out.printf("Consumer Group: %s%n", groupId);
       System.out.printf(" State: %s%n", description.state());
        System.out.printf(" Protocol Type: %s%n", description.protocolType());
        System.out.printf(" Protocol: %s%n", description.protocolData());
       System.out.printf(" Coordinator: %s%n", description.coordinator());
       System.out.println(" Members:");
       description.members().forEach(member -> {
           System.out.printf("
                                 Member ID: %s%n", member.memberId());
           System.out.printf("
                                 Client ID: %s%n", member.clientId());
           System.out.printf(" Host: %s%n", member.host());
           System.out.printf(" Assigned Partitions: %s%n",
member.assignment().topicPartitions());
       });
   }
   private static void listConsumerGroupOffsets(AdminClient adminClient, String
groupId)
           throws ExecutionException, InterruptedException {
       ListConsumerGroupOffsetsResult result =
           adminClient.listConsumerGroupOffsets(groupId);
       Map<TopicPartition, OffsetAndMetadata> offsets =
```

```
result.partitionsToOffsetAndMetadata().get();
        System.out.printf("Offsets for Consumer Group: %s%n", groupId);
        offsets.forEach((partition, offset) -> {
            System.out.printf(" %s: offset=%d, metadata=%s%n",
                partition, offset.offset(), offset.metadata());
       });
   }
   private static void resetConsumerGroupOffsets(AdminClient adminClient, String
groupId)
           throws ExecutionException, InterruptedException {
        // Reset to earliest offsets for all partitions of a topic
       TopicPartition partition = new TopicPartition("orders", 0);
       Map<TopicPartition, OffsetAndMetadata> resetOffsets = Map.of(
            partition, new OffsetAndMetadata(OL, "Reset to beginning")
        );
       AlterConsumerGroupOffsetsResult result =
            adminClient.alterConsumerGroupOffsets(groupId, resetOffsets);
        result.all().get();
       System.out.printf("Reset offsets for group %s%n", groupId);
   }
}
```

Reset Policies

Auto Offset Reset Configuration

```
public class OffsetResetPoliciesExample {

    // Reset to earliest available offset
    public static Properties createEarliestResetConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "earliest-reset-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
        StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
        StringDeserializer.class);

        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

        return props;
    }

    // Reset to latest offset (skip existing messages)
```

```
public static Properties createLatestResetConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "latest-reset-group");
        props.put(ConsumerConfig.KEY DESERIALIZER CLASS CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "latest");
        return props;
    }
    // Throw exception on missing offset
    public static Properties createNoneResetConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP ID CONFIG, "none-reset-group");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "none");
        return props;
    }
    public static void demonstrateResetPolicies() {
        // Scenario 1: New consumer group with earliest
        System.out.println("=== EARLIEST Reset Policy ===");
        demonstrateEarliestReset();
        // Scenario 2: New consumer group with latest
        System.out.println("\n=== LATEST Reset Policy ===");
        demonstrateLatestReset();
        // Scenario 3: Error handling with none
        System.out.println("\n=== NONE Reset Policy ===");
        demonstrateNoneReset();
    }
    private static void demonstrateEarliestReset() {
        Properties props = createEarliestResetConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            consumer.subscribe(Arrays.asList("orders"));
            // This will start reading from the earliest available offset
            // for new consumer groups or when committed offset is invalid
```

```
ConsumerRecords<String, String> records =
                consumer.poll(Duration.ofMillis(5000));
            System.out.printf("Records fetched with earliest reset: %d%n",
records.count());
        } catch (Exception e) {
            System.err.println("Error with earliest reset: " + e.getMessage());
   }
   private static void demonstrateLatestReset() {
        Properties props = createLatestResetConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            consumer.subscribe(Arrays.asList("orders"));
            // This will start reading from the latest offset
            // (skipping all existing messages)
            ConsumerRecords<String, String> records =
                consumer.poll(Duration.ofMillis(5000));
            System.out.printf("Records fetched with latest reset: %d%n",
records.count());
        } catch (Exception e) {
            System.err.println("Error with latest reset: " + e.getMessage());
   }
    private static void demonstrateNoneReset() {
        Properties props = createNoneResetConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
            consumer.subscribe(Arrays.asList("orders"));
            // This will throw an exception if no committed offset exists
            ConsumerRecords<String, String> records =
                consumer.poll(Duration.ofMillis(5000));
            System.out.printf("Records fetched with none reset: %d%n",
records.count());
        } catch (org.apache.kafka.clients.consumer.NoOffsetForPartitionException
e) {
            System.err.println("Expected exception with none reset: " +
e.getMessage());
        } catch (Exception e) {
            System.err.println("Unexpected error: " + e.getMessage());
```

```
public static void main(String[] args) {
    demonstrateResetPolicies();
}
```

Manual Offset Reset

```
public class ManualOffsetResetExample {
    public static void main(String[] args) {
        Properties props = createManualOffsetConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
            // Manual partition assignment (not using consumer groups)
            TopicPartition partition = new TopicPartition("orders", 0);
            consumer.assign(Arrays.asList(partition));
            // Different manual reset strategies
            demonstrateManualResets(consumer, partition);
        } catch (Exception e) {
            System.err.println("Manual offset reset failed: " + e.getMessage());
   }
   private static Properties createManualOffsetConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.ENABLE AUTO COMMIT CONFIG, false);
       // No group.id for manual assignment
       return props;
   }
   private static void demonstrateManualResets(KafkaConsumer<String, String>
consumer,
                                              TopicPartition partition) {
        // 1. Seek to beginning
        System.out.println("=== Seek to Beginning ===");
        consumer.seekToBeginning(Arrays.asList(partition));
        pollAndPrint(consumer, "Beginning");
       // 2. Seek to end
```

```
System.out.println("\n=== Seek to End ===");
        consumer.seekToEnd(Arrays.asList(partition));
        pollAndPrint(consumer, "End");
        // 3. Seek to specific offset
        System.out.println("\n=== Seek to Specific Offset ===");
        consumer.seek(partition, 10L);
        pollAndPrint(consumer, "Offset 10");
        // 4. Seek by timestamp
        System.out.println("\n=== Seek by Timestamp ===");
        seekByTimestamp(consumer, partition);
        // 5. Get offset information
        System.out.println("\n=== Offset Information ===");
        printOffsetInformation(consumer, partition);
    }
    private static void pollAndPrint(KafkaConsumer<String, String> consumer,
String context) {
        ConsumerRecords<String, String> records =
            consumer.poll(Duration.ofMillis(2000));
        System.out.printf("%s - Records: %d%n", context, records.count());
        records.forEach(record -> {
            System.out.printf(" Offset: %d, Key: %s, Value: %s%n",
                record.offset(), record.key(), record.value());
        });
    }
    private static void seekByTimestamp(KafkaConsumer<String, String> consumer,
                                      TopicPartition partition) {
        // Seek to 1 hour ago
        long oneHourAgo = System.currentTimeMillis() - (60 * 60 * 1000);
        Map<TopicPartition, Long> timestampsToSearch = Map.of(partition,
oneHourAgo);
        Map<TopicPartition, OffsetAndTimestamp> offsetsForTimes =
            consumer.offsetsForTimes(timestampsToSearch);
        OffsetAndTimestamp offsetAndTimestamp = offsetsForTimes.get(partition);
        if (offsetAndTimestamp != null) {
            consumer.seek(partition, offsetAndTimestamp.offset());
            System.out.printf("Seeking to timestamp %d, offset %d%n",
                oneHourAgo, offsetAndTimestamp.offset());
            pollAndPrint(consumer, "Timestamp seek");
        } else {
            System.out.println("No offset found for timestamp");
```

```
private static void printOffsetInformation(KafkaConsumer<String, String>
consumer,
                                             TopicPartition partition) {
       // Beginning offsets
        Map<TopicPartition, Long> beginningOffsets =
            consumer.beginningOffsets(Arrays.asList(partition));
        // End offsets
       Map<TopicPartition, Long> endOffsets =
            consumer.endOffsets(Arrays.asList(partition));
        // Current position
       long currentPosition = consumer.position(partition);
        System.out.printf("Partition: %s%n", partition);
        System.out.printf(" Beginning offset: %d%n",
beginningOffsets.get(partition));
        System.out.printf(" End offset: %d%n", endOffsets.get(partition));
        System.out.printf(" Current position: %d%n", currentPosition);
        System.out.printf(" Available records: %d%n",
            endOffsets.get(partition) - currentPosition);
   }
}
```

Lag Monitoring

JMX-based Lag Monitoring

```
import javax.management.MBeanServer;
import javax.management.ObjectName;
import java.lang.management.ManagementFactory;
import java.util.Set;
public class ConsumerLagMonitoring {
   public static void main(String[] args) {
        // Start a consumer to generate metrics
       Thread consumerThread = new Thread(() -> runConsumer());
        consumerThread.start();
       // Monitor consumer lag via JMX
       Thread monitorThread = new Thread(() -> monitorLagViaMx());
       monitorThread.start();
   }
    private static void runConsumer() {
        Properties props = createMonitoredConsumerConfig();
        try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
```

```
consumer.subscribe(Arrays.asList("orders"));
            while (true) {
                ConsumerRecords<String, String> records =
                    consumer.poll(Duration.ofMillis(1000));
                // Simulate processing time to create lag
                for (ConsumerRecord<String, String> record : records) {
                    try {
                        Thread.sleep(100); // Simulate slow processing
                    } catch (InterruptedException e) {
                        Thread.currentThread().interrupt();
                        return;
                    }
                }
                consumer.commitSync();
            }
        }
    }
    private static Properties createMonitoredConsumerConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "monitored-consumer-group");
        props.put(ConsumerConfig.CLIENT_ID_CONFIG, "lag-monitored-consumer");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.ENABLE AUTO COMMIT CONFIG, false);
        // Enable JMX metrics
        props.put("metric.reporters",
"org.apache.kafka.common.metrics.JmxReporter");
        return props;
    }
    private static void monitorLagViaMx() {
        MBeanServer server = ManagementFactory.getPlatformMBeanServer();
        try {
            while (true) {
                Thread.sleep(10000); // Check every 10 seconds
                // Query consumer lag metrics
                Set<ObjectName> consumerMetrics = server.queryNames(
                    new ObjectName("kafka.consumer:type=consumer-fetch-manager-
metrics,client-id=*"),
                    null
                );
```

```
System.out.println("\n=== Consumer Lag Metrics ===");
                for (ObjectName objName : consumerMetrics) {
                    try {
                        // Records lag max
                        Object recordsLagMax = server.getAttribute(objName,
"records-lag-max");
                        // Records consumed rate
                        Object recordsConsumedRate = server.getAttribute(objName,
"records-consumed-rate");
                        // Fetch latency average
                        Object fetchLatencyAvg = server.getAttribute(objName,
"fetch-latency-avg");
                        String clientId = objName.getKeyProperty("client-id");
                        System.out.printf("Client: %s%n", clientId);
                        System.out.printf(" Records Lag Max: %s%n",
recordsLagMax);
                        System.out.printf(" Records Consumed Rate: %.2f/sec%n",
recordsConsumedRate);
                        System.out.printf(" Fetch Latency Avg: %.2f ms%n",
fetchLatencyAvg);
                    } catch (Exception e) {
                        System.err.println("Error reading metric: " +
e.getMessage());
                    }
                }
                // Query coordinator metrics
                monitorCoordinatorMetrics(server);
            }
        } catch (InterruptedException e) {
            Thread.currentThread().interrupt();
        }
    }
    private static void monitorCoordinatorMetrics(MBeanServer server) {
        try {
            Set<ObjectName> coordinatorMetrics = server.queryNames(
                new ObjectName("kafka.consumer:type=consumer-coordinator-
metrics,client-id=*"),
                null
            );
            System.out.println("\n=== Coordinator Metrics ===");
            for (ObjectName objName : coordinatorMetrics) {
                try {
                    // Heartbeat rate
```

```
Object heartbeatRate = server.getAttribute(objName,
"heartbeat-rate");
                    // Join rate
                    Object joinRate = server.getAttribute(objName, "join-rate");
                    // Sync rate
                    Object syncRate = server.getAttribute(objName, "sync-rate");
                    String clientId = objName.getKeyProperty("client-id");
                    System.out.printf("Client: %s%n", clientId);
                    System.out.printf(" Heartbeat Rate: %.2f/sec%n",
heartbeatRate);
                    System.out.printf(" Join Rate: %.2f/sec%n", joinRate);
                    System.out.printf(" Sync Rate: %.2f/sec%n", syncRate);
                } catch (Exception e) {
                    // Metric might not be available
            }
        } catch (Exception e) {
            System.err.println("Error reading coordinator metrics: " +
e.getMessage());
        }
    }
}
```

Application-Level Lag Monitoring

```
import java.util.Map;
import java.util.concurrent.Executors;
import java.util.concurrent.ScheduledExecutorService;
import java.util.concurrent.TimeUnit;
public class ApplicationLevelLagMonitor {
    private final AdminClient adminClient;
   private final ScheduledExecutorService scheduler;
   private final String groupId;
   public ApplicationLevelLagMonitor(String bootstrapServers, String groupId) {
        Properties adminProps = new Properties();
        adminProps.put(AdminClientConfig.BOOTSTRAP_SERVERS_CONFIG,
bootstrapServers);
        this.adminClient = AdminClient.create(adminProps);
        this.scheduler = Executors.newScheduledThreadPool(1);
       this.groupId = groupId;
   }
```

```
public void startMonitoring() {
        scheduler.scheduleAtFixedRate(this::checkLag, 0, 30, TimeUnit.SECONDS);
    private void checkLag() {
        try {
            // Get consumer group offsets
            ListConsumerGroupOffsetsResult offsetsResult =
                adminClient.listConsumerGroupOffsets(groupId);
            Map<TopicPartition, OffsetAndMetadata> consumerOffsets =
                offsetsResult.partitionsToOffsetAndMetadata().get();
            if (consumerOffsets.isEmpty()) {
                System.out.println("No consumer offsets found for group: " +
groupId);
                return;
            }
            // Get latest offsets for topics
            Set<TopicPartition> partitions = consumerOffsets.keySet();
            Map<TopicPartition, ListOffsetsResult.ListOffsetsResultInfo>
latestOffsets =
                adminClient.listOffsets(
                    partitions.stream().collect(java.util.stream.Collectors.toMap(
                        tp -> tp,
                        tp -> OffsetSpec.latest()
                    ))
                ).all().get();
            System.out.printf("\n=== Consumer Lag Report (%s) ===%n",
                java.time.LocalDateTime.now());
            System.out.printf("Group: %s%n", groupId);
            long totalLag = 0;
            for (TopicPartition partition : partitions) {
                OffsetAndMetadata consumerOffset = consumerOffsets.get(partition);
                ListOffsetsResult.ListOffsetsResultInfo latestOffset =
latestOffsets.get(partition);
                if (consumerOffset != null && latestOffset != null) {
                    long lag = latestOffset.offset() - consumerOffset.offset();
                    totalLag += lag;
                    System.out.printf(" %s: consumer=%d, latest=%d, lag=%d%n",
                        partition, consumerOffset.offset(), latestOffset.offset(),
lag);
                    // Alert on high lag
                    if (lag > 1000) {
                        alertHighLag(partition, lag);
```

```
}
            System.out.printf("Total Lag: %d messages%n", totalLag);
            // Alert on total high lag
            if (totalLag > 5000) {
                alertHighTotalLag(totalLag);
            }
        } catch (Exception e) {
            System.err.println("Error checking consumer lag: " + e.getMessage());
        }
    }
    private void alertHighLag(TopicPartition partition, long lag) {
        System.err.printf("ALERT: High lag detected for %s: %d messages%n",
            partition, lag);
        // In real implementation:
        // - Send notification to monitoring system
        // - Trigger alert to ops team
        // - Log to alerting service
    }
    private void alertHighTotalLag(long totalLag) {
        System.err.printf("ALERT: High total lag detected for group %s: %d
messages%n",
            groupId, totalLag);
        // In real implementation:
       // - Check consumer health
        // - Consider scaling up consumers
        // - Alert operations team
    }
    public void shutdown() {
        scheduler.shutdown();
        adminClient.close();
    }
    public static void main(String[] args) {
        ApplicationLevelLagMonitor monitor =
            new ApplicationLevelLagMonitor("localhost:9092", "my-consumer-group");
        monitor.startMonitoring();
        // Shutdown hook
        Runtime.getRuntime().addShutdownHook(new Thread(monitor::shutdown));
        // Keep main thread alive
        try {
            Thread.sleep(Long.MAX_VALUE);
        } catch (InterruptedException e) {
```

```
Thread.currentThread().interrupt();
}
}
}
```

Comprehensive Java Examples

Production-Ready Consumer Application

```
import org.apache.kafka.clients.consumer.*;
import org.apache.kafka.common.errors.*;
import org.apache.kafka.common.TopicPartition;
import com.fasterxml.jackson.databind.ObjectMapper;
import java.time.Duration;
import java.util.*;
import java.util.concurrent.atomic.AtomicBoolean;
import java.util.concurrent.atomic.AtomicLong;
 * Production-ready Kafka Consumer with comprehensive error handling,
 * monitoring, and resilience features.
 */
public class ProductionKafkaConsumerExample {
    private final Properties config;
    private final String topicName;
    private final AtomicBoolean running = new AtomicBoolean(true);
    private final AtomicLong processedMessages = new AtomicLong(∅);
    private final AtomicLong failedMessages = new AtomicLong(♥);
    private final ObjectMapper objectMapper = new ObjectMapper();
    public ProductionKafkaConsumerExample(String topicName, Properties config) {
        this.topicName = topicName;
        this.config = config;
    public static void main(String[] args) {
        Properties config = createProductionConsumerConfig();
        ProductionKafkaConsumerExample consumerApp =
            new ProductionKafkaConsumerExample("order-events", config);
        // Setup shutdown hook
        Runtime.getRuntime().addShutdownHook(new Thread(consumerApp::shutdown));
        // Start consuming
        consumerApp.consume();
    }
    private static Properties createProductionConsumerConfig() {
```

```
Properties props = new Properties();
       // Connection settings
       props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
           "broker1:9092,broker2:9092,broker3:9092");
       props.put(ConsumerConfig.GROUP_ID_CONFIG, "order-processing-service");
       props.put(ConsumerConfig.CLIENT_ID_CONFIG, "order-processor-" +
UUID.randomUUID());
       // Serialization
       props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
           "org.apache.kafka.common.serialization.StringDeserializer");
       props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
           "org.apache.kafka.common.serialization.StringDeserializer");
       // Reliability settings
       props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false); // Manual
commit
       props.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
       // Performance settings
       wait
       props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, 100);
                                                                // 100
records max
       props.put(ConsumerConfig.MAX_POLL_INTERVAL_MS_CONFIG, 300000); // 5 min
processing
       // Group coordination
       props.put(ConsumerConfig.SESSION TIMEOUT MS CONFIG, 30000); // 30s
session
       props.put(ConsumerConfig.HEARTBEAT INTERVAL MS CONFIG, 3000); // 3s
heartbeat
       // Rebalancing
       props.put(ConsumerConfig.PARTITION_ASSIGNMENT_STRATEGY_CONFIG,
           "org.apache.kafka.clients.consumer.CooperativeStickyAssignor");
       // Static membership for stability
       props.put(ConsumerConfig.GROUP INSTANCE ID CONFIG,
           "order-processor-" + System.getenv().getOrDefault("HOSTNAME",
"localhost"));
       return props;
   }
   public void consume() {
       while (running.get()) {
           try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>
(config)) {
               consumer.subscribe(Arrays.asList(topicName), new
RebalanceHandler());
```

```
System.out.println("Consumer started, waiting for messages...");
                while (running.get()) {
                    try {
                        ConsumerRecords<String, String> records =
                            consumer.poll(Duration.ofMillis(1000));
                        if (records.isEmpty()) {
                            continue;
                        }
                        processRecordsBatch(records);
                        commitOffsets(consumer);
                        // Print statistics
                        printStatistics();
                    } catch (WakeupException e) {
                        System.out.println("Consumer wakeup called");
                        break;
                    } catch (CommitFailedException e) {
                        System.err.println("Commit failed during rebalance: " +
e.getMessage());
                        // Continue processing - rebalance will reset offsets
                    } catch (InvalidOffsetException e) {
                        System.err.println("Invalid offset, seeking to beginning:
" + e.getMessage());
                        consumer.seekToBeginning(e.partitions());
                    } catch (AuthorizationException e) {
                        System.err.println("Authorization failed: " +
e.getMessage());
                        break; // Fatal error
                    } catch (InterruptException e) {
                        System.err.println("Consumer interrupted: " +
e.getMessage());
                        Thread.currentThread().interrupt();
                        break;
                    } catch (Exception e) {
                        System.err.println("Unexpected consumer error: " +
e.getMessage());
                        e.printStackTrace();
                        // Decide whether to continue or stop
                        if (isFatalError(e)) {
                            break;
                        // Wait before retry
```

```
try {
                            Thread.sleep(5000);
                        } catch (InterruptedException ie) {
                            Thread.currentThread().interrupt();
                            break;
                        }
                    }
                }
            } catch (Exception e) {
                System.err.println("Consumer initialization failed: " +
e.getMessage());
                if (running.get()) {
                    // Wait before retry
                    try {
                        Thread.sleep(10000);
                    } catch (InterruptedException ie) {
                        Thread.currentThread().interrupt();
                        break;
                    }
                }
            }
        }
        System.out.println("Consumer stopped");
    }
    private void processRecordsBatch(ConsumerRecords<String, String> records) {
        System.out.printf("Processing batch of %d records%n", records.count());
        for (ConsumerRecord<String, String> record : records) {
            try {
                processMessage(record);
                processedMessages.incrementAndGet();
            } catch (Exception e) {
                failedMessages.incrementAndGet();
                handleProcessingError(record, e);
            }
        }
    }
    private void processMessage(ConsumerRecord<String, String> record) throws
Exception {
        // Deserialize message
        OrderEvent orderEvent = objectMapper.readValue(record.value(),
OrderEvent.class);
        // Business logic processing
        System.out.printf("Processing order: %s for user: %s, amount: $%.2f%n",
            orderEvent.getOrderId(), orderEvent.getUserId(),
orderEvent.getAmount());
```

```
// Simulate processing time
        Thread.sleep(50);
        // Validate order
        if (orderEvent.getAmount() <= ∅) {
            throw new IllegalArgumentException("Invalid order amount");
        // Process order (save to database, call external services, etc.)
        processOrder(orderEvent);
        // Update metrics
        updateProcessingMetrics(orderEvent);
    }
    private void processOrder(OrderEvent orderEvent) {
        // Simulate order processing
        // In real implementation:
        // - Save to database
        // - Call payment service
        // - Update inventory
        // - Send confirmation
        System.out.printf("Order %s processed successfully%n",
orderEvent.getOrderId());
    }
    private void updateProcessingMetrics(OrderEvent orderEvent) {
        // Update application metrics
        // In real implementation:
        // - Update Micrometer/Prometheus metrics
       // - Send to monitoring system
        // - Update business KPIs
    }
    private void handleProcessingError(ConsumerRecord<String, String> record,
Exception e) {
        System.err.printf("Failed to process record at offset %d: %s%n",
            record.offset(), e.getMessage());
        if (e instanceof IllegalArgumentException) {
            // Invalid data - skip and continue
            System.err.println("Skipping invalid record");
            logFailedRecord(record, e, "INVALID_DATA");
        } else if (e instanceof java.net.ConnectException) {
            // External service unavailable
            System.err.println("External service unavailable - sending to retry
queue");
            sendToRetryQueue(record, e);
        } else if (e instanceof java.sql.SQLException) {
            // Database error - might be transient
            System.err.println("Database error - sending to retry queue");
```

```
sendToRetryQueue(record, e);
        } else {
            // Unknown error - send to DLQ
            System.err.println("Unknown error - sending to dead letter queue");
            sendToDeadLetterQueue(record, e);
        }
    }
   private void sendToRetryQueue(ConsumerRecord<String, String> record, Exception
e) {
        // Implementation to send to retry topic
        System.out.printf("Sending to retry queue: key=%s, partition=%d,
offset=%d%n",
            record.key(), record.partition(), record.offset());
    }
    private void sendToDeadLetterQueue(ConsumerRecord<String, String> record,
Exception e) {
        // Implementation to send to DLQ topic
        System.out.printf("Sending to DLQ: key=%s, partition=%d, offset=%d%n",
            record.key(), record.partition(), record.offset());
    }
    private void logFailedRecord(ConsumerRecord<String, String> record, Exception
e, String reason) {
        System.err.printf("Failed record logged: reason=%s, key=%s, offset=%d,
error=%s%n",
            reason, record.key(), record.offset(), e.getMessage());
    }
    private void commitOffsets(KafkaConsumer<String, String> consumer) {
        try {
            consumer.commitSync();
        } catch (CommitFailedException e) {
            System.err.println("Offset commit failed: " + e.getMessage());
            // Could implement retry logic or continue
        }
    }
    private void printStatistics() {
        long processed = processedMessages.get();
        long failed = failedMessages.get();
        if (processed % 100 == 0 && processed > 0) {
            System.out.printf("Statistics: Processed=%d, Failed=%d, Success
Rate=%.2f%%n",
                processed, failed,
                (processed * 100.0) / (processed + failed));
        }
    }
    private boolean isFatalError(Exception e) {
        return e instanceof AuthorizationException ||
```

```
e instanceof SecurityException ||
           e instanceof OutOfMemoryError;
}
public void shutdown() {
    System.out.println("Shutting down consumer...");
    running.set(false);
}
// Rebalance handler for proper cleanup
private class RebalanceHandler implements ConsumerRebalanceListener {
    @Override
    public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
        System.out.println("Partitions revoked: " + partitions);
        // Cleanup before losing partitions
        // - Flush any pending work
       // - Save partition-specific state
        // - Commit current offsets handled by main loop
    }
   @Override
    public void onPartitionsAssigned(Collection<TopicPartition> partitions) {
        System.out.println("Partitions assigned: " + partitions);
        // Initialize for new partitions
        // - Load partition-specific state
        // - Reset processing metrics
        // - Log partition assignment for monitoring
    }
   @Override
    public void onPartitionsLost(Collection<TopicPartition> partitions) {
        System.err.println("Partitions lost: " + partitions);
        // Handle partition loss
        // - Clean up partition-specific resources
        // - Log error for monitoring
        // - Alert operations if needed
    }
}
// Data class for order events
public static class OrderEvent {
    private String orderId;
    private String userId;
    private double amount;
    private long timestamp;
    // Constructors, getters, setters
    public OrderEvent() {}
    public OrderEvent(String orderId, String userId, double amount, long
```

```
timestamp) {
            this.orderId = orderId;
            this.userId = userId;
            this.amount = amount;
            this.timestamp = timestamp;
        }
       // Getters
        public String getOrderId() { return orderId; }
        public String getUserId() { return userId; }
        public double getAmount() { return amount; }
        public long getTimestamp() { return timestamp; }
       // Setters
        public void setOrderId(String orderId) { this.orderId = orderId; }
        public void setUserId(String userId) { this.userId = userId; }
        public void setAmount(double amount) { this.amount = amount; }
        public void setTimestamp(long timestamp) { this.timestamp = timestamp; }
   }
}
```

A Comparisons & Trade-offs

Consumer Configuration Trade-offs

| Parameter | Higher Value | Lower Value | Recommendation |
|-----------------------|--|--|-----------------------------------|
| max.poll.records | † Throughput, † Memory, † Processing time | ↓ Memory, ↓ Latency | 100-1000 based on message size |
| fetch.min.bytes | ↑ Throughput, ↑ Latency | ↓ Latency, ↓ Throughput | 1KB-100KB based on traffic |
| session.timeout.ms | ↑ Tolerance to GC pauses, ↓ Failure detection | ↑ Failure detection, ↓ GC tolerance | 30s for production |
| heartbeat.interval.ms | ↓ Network overhead | ↑ Failure detection speed | 1/3 of session timeout |

Commit Strategy Comparison

| Strategy | Data Guarantees | Performance | Use Case |
|--------------|-------------------------------|----------------------|------------------------------------|
| Auto Commit | At-most-once (potential loss) | ★ ★ ★ High | Metrics, logs, non-critical |
| Sync Commit | At-least-once | ★ Lower | Financial, critical data |
| Async Commit | At-least-once | ★ ★ Good | High-volume, some duplicates OK |

| Strategy | Data Guarantees | Performance | Use Case |
|------------|-----------------|------------------|-----------------------------------|
| Manual per | At least ones | ♣ Laurant | Critical processing fine control |
| Record | At-least-once | ★ Lowest | Critical processing, fine control |

Rebalancing Strategy Comparison

| Strategy | Distribution | Partition Movement | Use Case |
|-----------------------|------------------------------------|--------------------------|----------------------------------|
| Range | Can be uneven with multiple topics | Minimal | Single topic, ordered processing |
| Round Robin | Even distribution | Higher movement | Multiple topics, balanced load |
| Sticky | Even + sticky | Minimal movement | Stateful consumers |
| Cooperative Sticky | Even + sticky + incremental | Minimal + no downtime | Production systems |

Common Pitfalls & Best Practices

1. Poll Loop Mistakes

✗ Blocking Operations in Poll Loop

```
// DON'T - Blocking operations in poll loop
while (true) {
    ConsumerRecords<String, String> records =
    consumer.poll(Duration.ofMillis(1000));

    for (ConsumerRecord<String, String> record : records) {
        // This blocks the consumer thread!
        callSlowExternalService(record.value()); // 5 seconds per call
    }

    consumer.commitSync();
}
```

```
// DO - Use separate threads for heavy processing
ExecutorService executor = Executors.newFixedThreadPool(10);
while (running.get()) {
   ConsumerRecords<String, String> records =
   consumer.poll(Duration.ofMillis(1000));
   List<Future<Void>> futures = new ArrayList<>();
```

```
for (ConsumerRecord<String, String> record : records) {
        Future<Void> future = executor.submit(() -> {
            processRecord(record);
            return null;
        });
       futures.add(future);
   }
   // Wait for all processing to complete
   for (Future<Void> future : futures) {
       try {
            future.get(30, TimeUnit.SECONDS); // Timeout to prevent hanging
       } catch (Exception e) {
            System.err.println("Processing failed: " + e.getMessage());
        }
   }
   consumer.commitSync();
}
```

X Infinite Poll Timeout

```
// DON'T - Infinite timeout blocks shutdown
while (running) {
    ConsumerRecords<String, String> records =
        consumer.poll(Duration.ofMillis(Long.MAX_VALUE)); // Never returns!

    // Processing...
}
```

```
// DO - Use reasonable timeout for responsive shutdown
while (running.get()) {
    ConsumerRecords<String, String> records =
        consumer.poll(Duration.ofMillis(1000)); // 1 second timeout

    if (records.isEmpty()) {
        continue; // Check running flag frequently
    }

    // Process records...
}
```

2. Offset Management Issues

X Committing Before Processing

```
// DON'T - Commit before processing (at-most-once, potential loss)
ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(1000));

// Committing first - dangerous!
consumer.commitSync();

for (ConsumerRecord<String, String> record : records) {
   processRecord(record); // If this fails, message is lost forever
}
```

```
// DO - Commit after successful processing
ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(1000));
boolean allProcessedSuccessfully = true;
for (ConsumerRecord<String, String> record : records) {
    try {
        processRecord(record);
    } catch (Exception e) {
        System.err.println("Processing failed: " + e.getMessage());
        allProcessedSuccessfully = false;
        break; // Stop processing batch
    }
}
// Only commit if all records processed successfully
if (allProcessedSuccessfully) {
    consumer.commitSync();
}
```

X Not Handling Commit Failures

```
// DON'T - Ignore commit failures
consumer.commitSync(); // What if this fails?
```

```
// DO - Handle commit failures appropriately
try {
    consumer.commitSync();
} catch (CommitFailedException e) {
    System.err.println("Commit failed: " + e.getMessage());

    // Strategy options:
    // 1. Continue (rebalance will reset offsets)
    // 2. Retry commit
    // 3. Alert monitoring system
```

```
alertCommitFailure(e);
}
```

3. Rebalancing Issues

X Long Processing Blocking Rebalance

```
// DON'T - Long processing in rebalance listener
class BadRebalanceListener implements ConsumerRebalanceListener {
    @Override
    public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
        // This blocks rebalancing!
        saveMassiveStateToDatabase(); // 30 seconds
        consumer.commitSync(); // Another potential delay
    }
}
```

```
// DO - Quick cleanup in rebalance listener
class GoodRebalanceListener implements ConsumerRebalanceListener {
   @Override
   public void onPartitionsRevoked(Collection<TopicPartition> partitions) {
        System.out.println("Partitions revoked: " + partitions);
       // Quick operations only
       try {
            consumer.commitSync(Duration.ofSeconds(5)); // Timeout commit
        } catch (Exception e) {
            System.err.println("Quick commit failed: " + e.getMessage());
        }
        // Async cleanup for heavy operations
        CompletableFuture.runAsync(() -> {
            saveMassiveStateToDatabase();
       });
   }
}
```

4. Memory Management Problems

X Unbounded Record Processing

```
// DON'T - Process unlimited records without backpressure
props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, Integer.MAX_VALUE);
props.put(ConsumerConfig.FETCH_MAX_BYTES_CONFIG, Integer.MAX_VALUE);
// This can cause OutOfMemoryError!
```

```
ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(1000));
List<String> allData = new ArrayList<>();
for (ConsumerRecord<String, String> record : records) {
   allData.add(record.value()); // Unbounded memory growth
}
```

```
// DO - Set reasonable limits and implement backpressure
props.put(ConsumerConfig.MAX POLL RECORDS CONFIG, 500);
props.put(ConsumerConfig.FETCH_MAX_BYTES_CONFIG, 52428800); // 50MB
ConsumerRecords<String, String> records = consumer.poll(Duration.ofMillis(1000));
// Process in batches to control memory
int batchSize = 100;
List<ConsumerRecord<String, String>> batch = new ArrayList<>();
for (ConsumerRecord<String, String> record : records) {
    batch.add(record);
    if (batch.size() >= batchSize) {
        processBatch(batch);
        batch.clear(); // Free memory
    }
}
// Process remaining records
if (!batch.isEmpty()) {
    processBatch(batch);
}
```

Best Practices Summary

☑ Consumer Configuration Best Practices

- 1. **Use manual commits** for critical data (enable.auto.commit=false)
- 2. Set appropriate timeouts (session.timeout.ms=30000, heartbeat.interval.ms=3000)
- 3. Limit poll size (max.poll.records=500 based on processing speed)
- 4. Use cooperative sticky assignor for better rebalancing
- 5. **Enable static membership** for stable applications

☑ Processing Best Practices

- 1. **Keep poll loop lightweight** offload heavy processing to threads
- 2. **Implement proper error handling** distinguish retriable vs fatal errors
- 3. **Commit after processing** ensure at-least-once semantics
- 4. Handle rebalances gracefully quick cleanup only
- 5. Monitor consumer lag set up alerting for high lag

2025-09-24

Operational Best Practices

- 1. Monitor key metrics lag, commit rate, rebalance frequency
- 2. **Set up proper logging** include partition, offset, timestamp info
- 3. Implement circuit breakers handle downstream failures
- 4. Use dead letter queues for failed messages
- 5. **Test failure scenarios** consumer restarts, network partitions

Real-World Use Cases

1. Order Processing System

```
@Service
public class OrderProcessingConsumer {
    private final OrderService orderService;
    private final InventoryService inventoryService;
    private final PaymentService paymentService;
    private final NotificationService notificationService;
    @EventListener
    public void consumeOrderEvents() {
        Properties props = createOrderConsumerConfig();
        try (KafkaConsumer<String, OrderEvent> consumer = new KafkaConsumer<>
(props)) {
            consumer.subscribe(Arrays.asList("order-events"), new
OrderRebalanceListener());
            while (running) {
                ConsumerRecords<String, OrderEvent> records =
                    consumer.poll(Duration.ofMillis(1000));
                // Process orders in transaction
                processOrdersInTransaction(records);
                consumer.commitSync();
            }
        }
    }
    @Transactional
    private void processOrdersInTransaction(ConsumerRecords<String, OrderEvent>
records) {
        for (ConsumerRecord<String, OrderEvent> record : records) {
            OrderEvent event = record.value();
            try {
                // Process order atomically
                Order order = orderService.processOrder(event);
                inventoryService.reserveItems(order);
                paymentService.processPayment(order);
```

```
notificationService.sendConfirmation(order);
                logger.info("Order processed: {}", event.getOrderId());
            } catch (InsufficientInventoryException e) {
                // Business logic error - send notification
                notificationService.sendInventoryAlert(event);
            } catch (PaymentFailedException e) {
                // Payment failed - update order status
                orderService.markPaymentFailed(event.getOrderId());
            } catch (Exception e) {
                // Unexpected error - send to DLQ
                sendToDeadLetterQueue(record, e);
                throw e; // Fail transaction
            }
       }
   }
}
```

2. Real-time Analytics Consumer

```
@Component
public class AnalyticsConsumer {
    private final MetricsCollector metricsCollector;
    private final TimeWindowAggregator aggregator;
    public void consumeUserEvents() {
        Properties props = createAnalyticsConsumerConfig();
        try (KafkaConsumer<String, UserEvent> consumer = new KafkaConsumer<>>
(props)) {
            consumer.subscribe(Arrays.asList("user-events"));
            // Use time-window processing
            Map<String, UserSession> activeSessions = new HashMap<>();
            while (running) {
                ConsumerRecords<String, UserEvent> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, UserEvent> record : records) {
                    UserEvent event = record.value();
                    // Update user session
                    String userId = event.getUserId();
                    UserSession session = activeSessions.computeIfAbsent(userId,
                        k -> new UserSession(userId));
```

```
session.addEvent(event);
                    // Collect real-time metrics
                    metricsCollector.recordEvent(event.getEventType());
                    // Update aggregations
                    aggregator.updateTimeWindow(event);
                    // Expire old sessions
                    expireOldSessions(activeSessions);
                }
                // Publish aggregated metrics
                publishAggregatedMetrics();
                consumer.commitAsync();
        }
   }
    private Properties createAnalyticsConsumerConfig() {
        Properties props = new Properties();
        props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        props.put(ConsumerConfig.GROUP_ID_CONFIG, "analytics-consumer-group");
        // Optimize for throughput
        props.put(ConsumerConfig.FETCH_MIN_BYTES_CONFIG, 100000); // 100KB
        props.put(ConsumerConfig.FETCH_MAX_WAIT_MS_CONFIG, 500); // 500ms
        props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, 1000); // 1000 records
        // Use async commits for better performance
        props.put(ConsumerConfig.ENABLE AUTO COMMIT CONFIG, false);
       return props;
   }
}
```

3. Event Sourcing Consumer

```
@Service
public class EventSourcingConsumer {

    private final EventStore eventStore;
    private final ProjectionUpdater projectionUpdater;

public void consumeEventStream() {
        Properties props = createEventSourcingConfig();

        try (KafkaConsumer<String, DomainEvent> consumer = new KafkaConsumer<> (props)) {

        // Manual partition assignment for event sourcing
```

```
TopicPartition partition = new TopicPartition("domain-events", 0);
            consumer.assign(Arrays.asList(partition));
            // Start from last processed offset
            long lastProcessedOffset = eventStore.getLastProcessedOffset();
            consumer.seek(partition, lastProcessedOffset);
            while (running) {
                ConsumerRecords<String, DomainEvent> records =
                    consumer.poll(Duration.ofMillis(1000));
                for (ConsumerRecord<String, DomainEvent> record : records) {
                    DomainEvent event = record.value();
                    // Ensure exactly-once processing with idempotency
                    if (eventStore.isEventProcessed(event.getEventId())) {
                        continue; // Skip already processed events
                    }
                    // Apply event atomically
                    applyEventAtomically(event, record.offset());
                }
                // Commit only after successful processing
                consumer.commitSync();
            }
        }
   }
   @Transactional
   private void applyEventAtomically(DomainEvent event, long offset) {
        // Store event
        eventStore.appendEvent(event, offset);
        // Update projections
        projectionUpdater.updateProjections(event);
       // Update last processed offset
        eventStore.updateLastProcessedOffset(offset);
   }
}
```

4. CDC (Change Data Capture) Consumer

```
@Service
public class CDCConsumer {

   private final DatabaseSyncService syncService;
   private final SearchIndexService searchService;
   private final CacheService cacheService;
```

```
public void consumeDBChanges() {
        Properties props = createCDCConsumerConfig();
        try (KafkaConsumer<String, DatabaseChangeEvent> consumer = new
KafkaConsumer<>(props)) {
            consumer.subscribe(Arrays.asList("db-changes"));
            while (running) {
                ConsumerRecords<String, DatabaseChangeEvent> records =
                    consumer.poll(Duration.ofMillis(1000));
                // Group changes by table for efficient processing
                Map<String, List<DatabaseChangeEvent>> changesByTable =
                    groupChangesByTable(records);
                // Process changes by table
                for (Map.Entry<String, List<DatabaseChangeEvent>> entry :
changesByTable.entrySet()) {
                    String tableName = entry.getKey();
                    List<DatabaseChangeEvent> changes = entry.getValue();
                    processTableChanges(tableName, changes);
                }
                consumer.commitSync();
            }
        }
   }
   private void processTableChanges(String tableName, List<DatabaseChangeEvent>
changes) {
       for (DatabaseChangeEvent change : changes) {
            try {
                switch (change.getOperation()) {
                    case INSERT:
                        handleInsert(tableName, change);
                        break;
                    case UPDATE:
                        handleUpdate(tableName, change);
                        break;
                    case DELETE:
                        handleDelete(tableName, change);
                        break;
                }
            } catch (Exception e) {
                logger.error("Failed to process change for table {}: {}",
                    tableName, e.getMessage());
                // Could implement retry logic or send to DLQ
            }
        }
   }
   private void handleUpdate(String tableName, DatabaseChangeEvent change) {
```

```
// Update secondary database
syncService.updateRecord(tableName, change.getAfter());

// Update search index
searchService.updateDocument(change.getRecordId(), change.getAfter());

// Invalidate cache
cacheService.invalidate(tableName, change.getRecordId());
}
```

Wersion Highlights

Kafka 4.0 (September 2025) - Current Latest

- * New Consumer Protocol (KIP-848): Faster rebalancing, better partition assignment
- * Enhanced Cooperative Rebalancing: Further reduced downtime during rebalances
- **Improved Lag Monitoring**: Better JMX metrics and monitoring capabilities
- A Java 17+ Required: For brokers (Java 11+ for clients)

Kafka 3.x Series Consumer Features

- 3.6 (Oct 2023): Consumer metrics improvements, better error reporting
- 3.5 (Jun 2023): Enhanced static membership stability
- 3.4 (Feb 2023): Improved consumer group rebalancing performance
- 3.3 (Oct 2022): Consumer optimizations for KRaft mode
- 3.2 (May 2022): Better consumer coordinator design
- 3.1 (Jan 2022): Improved incremental cooperative rebalancing
- 3.0 (Sep 2021): Consumer performance optimizations

Key Consumer Evolution

| Version | Consumer Feature | Impact |
|---------|----------------------------------|--------------------------|
| 4.0 | New consumer protocol (KIP-848) | 50% faster rebalancing |
| 2.4 | Cooperative sticky assignor | Incremental rebalancing |
| 2.3 | Static membership (KIP-345) | Reduced rebalances |
| 2.0 | Improved rebalancing protocol | Better failure detection |
| 1.1 | Headers support in consumer | Message metadata access |
| 0.11 | Exactly-once consumption support | Transactional reading |
| 0.10 | Timestamp support | Time-based seeking |
| 0.9 | Consumer rewrite | Much better performance |

Current Recommendations (2025)

```
// Modern Kafka 4.0 consumer configuration
public static Properties modernConsumerConfig() {
    Properties props = new Properties();
    props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    props.put(ConsumerConfig.GROUP_ID_CONFIG, "modern-consumer-group");
    props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
    props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
    // Kafka 4.0 best practices
    props.put(ConsumerConfig.PARTITION_ASSIGNMENT_STRATEGY_CONFIG,
        "org.apache.kafka.clients.consumer.CooperativeStickyAssignor");
    // Static membership for stability
    props.put(ConsumerConfig.GROUP INSTANCE ID CONFIG,
        "consumer-" + System.getenv("HOSTNAME"));
    // Optimized settings
    props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, 30000);
    props.put(ConsumerConfig.HEARTBEAT_INTERVAL_MS_CONFIG, 3000);
    props.put(ConsumerConfig.MAX_POLL_RECORDS_CONFIG, 500);
    props.put(ConsumerConfig.FETCH_MIN_BYTES_CONFIG, 50000);
    // Manual offset management
    props.put(ConsumerConfig.ENABLE AUTO COMMIT CONFIG, false);
    return props;
}
```

Additional Resources

Official Documentation

- Kafka Consumer API Documentation
- Consumer Configuration Reference
- Consumer Group Protocol

Learning Resources

- Confluent Consumer Tutorial
- Apache Kafka Consumer Deep Dive
- Consumer Best Practices

Nools & Monitoring

- Consumer Lag Monitoring
- JMX Metrics for Consumers
- kafka-consumer-groups.sh CLI

™ Troubleshooting

- Consumer Rebalancing Issues
- Common Consumer Problems
- Performance Tuning Guide

Last Updated: September 2025

Kafka Version: 4.0.0

Java Compatibility: 11+ (clients), 17+ (recommended)

Pro Tip: Start with cooperative sticky assignor and static membership for production systems. Monitor consumer lag closely and implement proper error handling with retry queues and dead letter topics.