

# Kafka Connect: Complete Developer Guide

---

A comprehensive refresher on Apache Kafka Connect, designed for both beginners and experienced developers. This README covers connectors, deployment modes, configuration, customization, and real-world applications with detailed Java examples.

## Table of Contents

- [Basics](#)
    - [Source vs Sink Connectors](#)
    - [Standalone vs Distributed Mode](#)
    - [Config Management via REST API](#)
  - [Common Connectors](#)
    - [JDBC Source/Sink](#)
    - [Debezium CDC](#)
    - [Elasticsearch/S3 Connectors](#)
  - [Customization](#)
    - [Writing Custom Connectors](#)
    - [Single Message Transforms \(SMTs\)](#)
  - [Comprehensive Java Examples](#)
  - [Comparisons & Trade-offs](#)
  - [Common Pitfalls & Best Practices](#)
  - [Real-World Use Cases](#)
  - [Version Highlights](#)
  - [Additional Resources](#)
- 

## Basics

### Simple Explanation

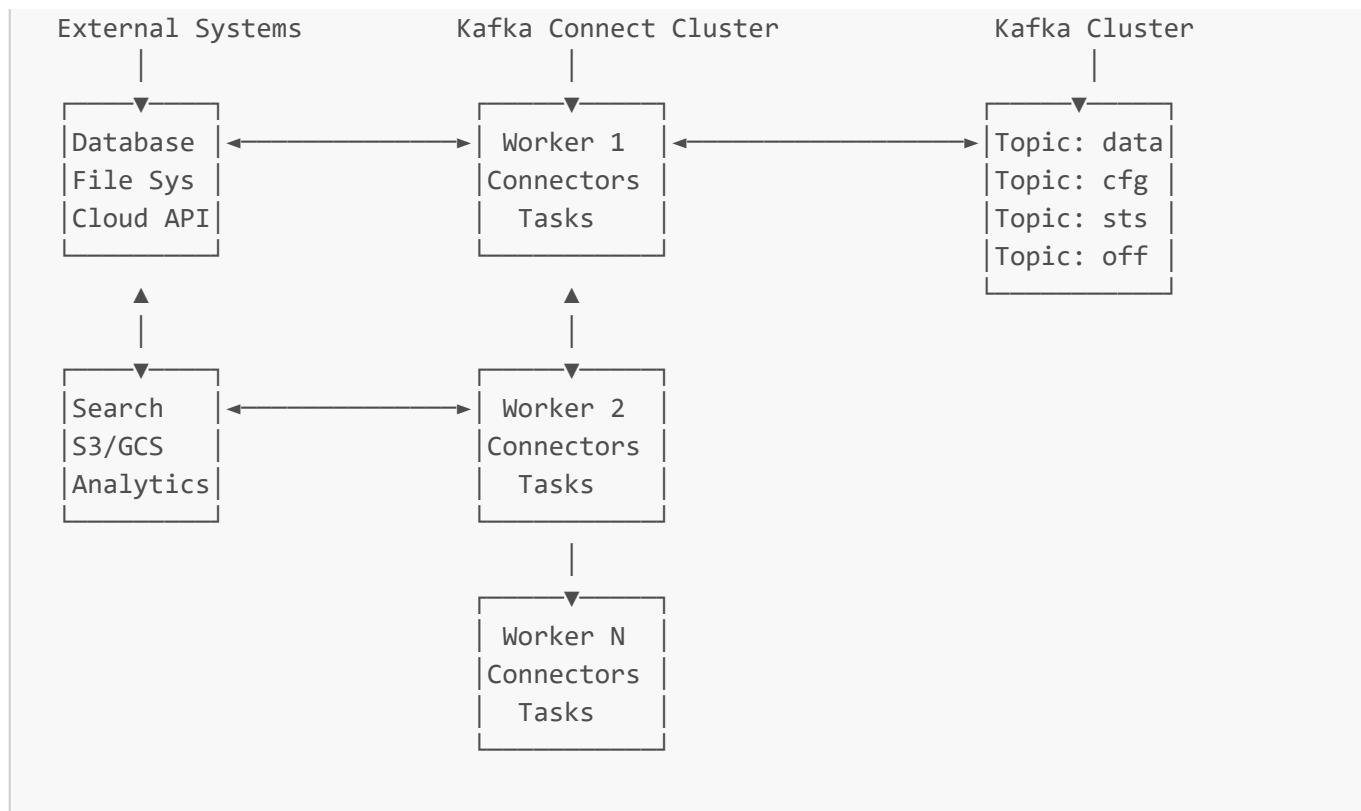
Kafka Connect is a framework for building and running reusable connectors that move data between Apache Kafka and other systems. It provides a REST API for managing connectors, automatic offset management, distributed operation, and fault tolerance.

### Problem It Solves

- **Data Integration:** Eliminates custom code for moving data in/out of Kafka
- **Scalability:** Distributes connector tasks across multiple workers
- **Fault Tolerance:** Automatic recovery and exactly-once delivery guarantees
- **Operational Simplicity:** Centralized management via REST API

### Internal Architecture

Kafka Connect Architecture:



## How It Works Under the Hood

1. **Workers:** JVM processes that execute connector code and manage tasks
2. **Connectors:** Define how to copy data (splitting into tasks)
3. **Tasks:** Actual units of work that copy data
4. **Offset Storage:** Tracks progress in external systems and Kafka
5. **Configuration Storage:** Stores connector configurations in Kafka topics

## Relevant Data Structures

- **connect-configs:** Compacted topic storing connector configurations
- **connect-offsets:** Tracks source connector progress (external system offsets)
- **connect-status:** Monitors connector and task status

## Source vs Sink Connectors

### Source Connectors - Import Data INTO Kafka

```
// Example: Database Source Connector Flow
External Database → Source Connector → Kafka Topic

// Flow:
1. Poll external system for new data
2. Convert to Kafka records
3. Send to Kafka topics
4. Track offsets in external system
```

## Sink Connectors - Export Data FROM Kafka

```
// Example: Elasticsearch Sink Connector Flow  
Kafka Topic → Sink Connector → Elasticsearch Index  
  
// Flow:  
1. Consume records from Kafka  
2. Transform to target format  
3. Write to external system  
4. Commit Kafka offsets
```

## Source Connector Example Configuration

```
import java.util.Properties;  
import java.util.Map;  
import java.util.HashMap;  
  
public class SourceConnectorConfig {  
  
    public static Map<String, String> createJDBCSourceConfig() {  
        Map<String, String> config = new HashMap<>();  
  
        // Connector class  
        config.put("connector.class",  
"io.confluent.connect.jdbc.JdbcSourceConnector");  
  
        // Connection properties  
        config.put("connection.url", "jdbc:postgresql://localhost:5432/mydb");  
        config.put("connection.user", "user");  
        config.put("connection.password", "password");  
  
        // Source configuration  
        config.put("table.whitelist", "users,orders,products");  
        config.put("mode", "incrementing");  
        config.put("incrementing.column.name", "id");  
        config.put("timestamp.column.name", "updated_at");  
  
        // Topic naming  
        config.put("topic.prefix", "postgres-");  
  
        // Polling configuration  
        config.put("poll.interval.ms", "5000");  
        config.put("batch.max.rows", "100");  
  
        return config;  
    }  
}
```

## Sink Connector Example Configuration

```

public class SinkConnectorConfig {

    public static Map<String, String> createElasticsearchSinkConfig() {
        Map<String, String> config = new HashMap<>();

        // Connector class
        config.put("connector.class",
            "io.confluent.connect.elasticsearch.ElasticsearchSinkConnector");

        // Connection properties
        config.put("connection.url", "http://localhost:9200");
        config.put("connection.username", "elastic");
        config.put("connection.password", "password");

        // Topics to consume
        config.put("topics", "postgres-users,postgres-orders");

        // Index configuration
        config.put("type.name", "_doc");
        config.put("key.ignore", "false");
        config.put("schema.ignore", "true");

        // Batching
        config.put("batch.size", "100");
        config.put("max.in.flight.requests", "5");

        // Error handling
        config.put("behavior.on.null.values", "ignore");
        config.put("behavior.on.malformed.documents", "warn");

        return config;
    }
}

```

## Standalone vs Distributed Mode

### Standalone Mode

```

public class StandaloneConnectExample {

    public static void main(String[] args) {
        // Standalone mode configuration
        Properties workerProps = new Properties();
        workerProps.put("bootstrap.servers", "localhost:9092");
        workerProps.put("key.converter",
            "org.apache.kafka.connect.json.JsonConverter");
        workerProps.put("value.converter",
            "org.apache.kafka.connect.json.JsonConverter");
    }
}

```

```
workerProps.put("key.converter.schemas.enable", false);
workerProps.put("value.converter.schemas.enable", false);

// Offset storage (file-based for standalone)
workerProps.put("offset.storage.file.filename", "/tmp/connect.offsets");
workerProps.put("offset.flush.interval.ms", 10000);

// Plugin path
workerProps.put("plugin.path", "/usr/share/java,/usr/share/confluent-hub-
components");

System.out.println("Standalone Connect Configuration:");
workerProps.forEach((key, value) ->
    System.out.println(key + " = " + value));
}

/*
 * Standalone Mode Characteristics:
 *
 * Pros:
 * - Simple setup and configuration
 * - Good for development and testing
 * - Single point of failure acceptable
 *
 * Cons:
 * - No fault tolerance
 * - Limited scalability
 * - Manual configuration management
 *
 * Use Cases:
 * - Development environments
 * - Simple integrations
 * - Proof of concepts
 */
}
```

## Distributed Mode

```
public class DistributedConnectExample {

    public static void main(String[] args) {
        // Distributed mode configuration
        Properties workerProps = new Properties();
        workerProps.put("bootstrap.servers", "localhost:9092");

        // Group and client identification
        workerProps.put("group.id", "connect-cluster");
        workerProps.put("client.id", "connect-worker-1");

        // Converters
        workerProps.put("key.converter",
```

```
"org.apache.kafka.connect.json.JsonConverter");
    workerProps.put("value.converter",
"org.apache.kafka.connect.json.JsonConverter");
    workerProps.put("key.converter.schemas.enable", true);
    workerProps.put("value.converter.schemas.enable", true);

    // Internal topic configuration
    workerProps.put("config.storage.topic", "connect-configs");
    workerProps.put("config.storage.replication.factor", 3);

    workerProps.put("offset.storage.topic", "connect-offsets");
    workerProps.put("offset.storage.replication.factor", 3);
    workerProps.put("offset.storage.partitions", 25);

    workerProps.put("status.storage.topic", "connect-status");
    workerProps.put("status.storage.replication.factor", 3);
    workerProps.put("status.storage.partitions", 5);

    // Worker configuration
    workerProps.put("offset.flush.interval.ms", 10000);
    workerProps.put("task.shutdown.graceful.timeout.ms", 30000);

    // REST API
    workerProps.put("rest.port", 8083);
    workerProps.put("rest.host.name", "localhost");

    // Plugin path
    workerProps.put("plugin.path", "/usr/share/java,/usr/share/confluent-hub-components");

    System.out.println("Distributed Connect Configuration:");
    workerProps.forEach((key, value) ->
        System.out.println(key + "=" + value));
}

/*
 * Distributed Mode Characteristics:
 *
 * Pros:
 * - Fault tolerant (automatic failover)
 * - Horizontally scalable
 * - Load balancing across workers
 * - REST API for management
 *
 * Cons:
 * - More complex setup
 * - Requires Kafka for coordination
 * - Network overhead
 *
 * Use Cases:
 * - Production environments
 * - High-availability requirements
 * - Large-scale data integration
*/
```

```
 */  
}
```

## Config Management via REST API

### REST API Client Example

```
import java.io.IOException;  
import java.net.URI;  
import java.net.http.HttpClient;  
import java.net.http.HttpRequest;  
import java.net.http.HttpResponse;  
import java.time.Duration;  
import com.fasterxml.jackson.databind.ObjectMapper;  
import com.fasterxml.jackson.databind.JsonNode;  
  
public class KafkaConnectRestClient {  
  
    private final HttpClient httpClient;  
    private final String connectUrl;  
    private final ObjectMapper objectMapper;  
  
    public KafkaConnectRestClient(String connectUrl) {  
        this.connectUrl = connectUrl;  
        this.httpClient = HttpClient.newBuilder()  
            .connectTimeout(Duration.ofSeconds(10))  
            .build();  
        this.objectMapper = new ObjectMapper();  
    }  
  
    // List all connectors  
    public String[] listConnectors() throws IOException, InterruptedException {  
        HttpRequest request = HttpRequest.newBuilder()  
            .uri(URI.create(connectUrl + "/connectors"))  
            .GET()  
            .build();  
  
        HttpResponse<String> response = httpClient.send(request,  
            HttpResponse.BodyHandlers.ofString());  
  
        if (response.statusCode() == 200) {  
            return objectMapper.readValue(response.body(), String[].class);  
        } else {  
            throw new RuntimeException("Failed to list connectors: " +  
                response.body());  
        }  
    }  
  
    // Create or update connector  
    public void createOrUpdateConnector(String name, Map<String, String> config)  
        throws IOException, InterruptedException {
```

```
Map<String, Object> connectorConfig = new HashMap<>();
connectorConfig.put("name", name);
connectorConfig.put("config", config);

String requestBody = objectMapper.writeValueAsString(connectorConfig);

HttpRequest request = HttpRequest.newBuilder()
    .uri(URI.create(connectUrl + "/connectors"))
    .header("Content-Type", "application/json")
    .POST(HttpRequest.BodyPublishers.ofString(requestBody))
    .build();

HttpResponse<String> response = httpClient.send(request,
    HttpResponse.BodyHandlers.ofString());

if (response.statusCode() != 201 && response.statusCode() != 200) {
    throw new RuntimeException("Failed to create connector: " +
response.body());
}

System.out.println("Connector created/updated successfully: " + name);
}

// Get connector status
public ConnectorStatus getConnectorStatus(String name)
    throws IOException, InterruptedException {

HttpRequest request = HttpRequest.newBuilder()
    .uri(URI.create(connectUrl + "/connectors/" + name + "/status"))
    .GET()
    .build();

HttpResponse<String> response = httpClient.send(request,
    HttpResponse.BodyHandlers.ofString());

if (response.statusCode() == 200) {
    JsonNode statusNode = objectMapper.readTree(response.body());
    return new ConnectorStatus(
        statusNode.get("name").asText(),
        statusNode.get("connector").get("state").asText(),
        statusNode.get("tasks").size()
    );
} else {
    throw new RuntimeException("Failed to get connector status: " +
response.body());
}
}

// Restart connector
public void restartConnector(String name) throws IOException,
InterruptedException {
    HttpRequest request = HttpRequest.newBuilder()
        .uri(URI.create(connectUrl + "/connectors/" + name + "/restart"))
}
```

```
.POST(HttpRequest.BodyPublishers.noBody())
.build();

HttpResponse<String> response = httpClient.send(request,
    HttpResponse.BodyHandlers.ofString());

if (response.statusCode() != 204) {
    throw new RuntimeException("Failed to restart connector: " +
response.body());
}

System.out.println("Connector restarted successfully: " + name);
}

// Pause connector
public void pauseConnector(String name) throws IOException,
InterruptedException {
    HttpRequest request = HttpRequest.newBuilder()
        .uri(URI.create(connectUrl + "/connectors/" + name + "/pause"))
        .PUT(HttpRequest.BodyPublishers.noBody())
        .build();

HttpResponse<String> response = httpClient.send(request,
    HttpResponse.BodyHandlers.ofString());

if (response.statusCode() != 202) {
    throw new RuntimeException("Failed to pause connector: " +
response.body());
}

System.out.println("Connector paused successfully: " + name);
}

// Resume connector
public void resumeConnector(String name) throws IOException,
InterruptedException {
    HttpRequest request = HttpRequest.newBuilder()
        .uri(URI.create(connectUrl + "/connectors/" + name + "/resume"))
        .PUT(HttpRequest.BodyPublishers.noBody())
        .build();

HttpResponse<String> response = httpClient.send(request,
    HttpResponse.BodyHandlers.ofString());

if (response.statusCode() != 202) {
    throw new RuntimeException("Failed to resume connector: " +
response.body());
}

System.out.println("Connector resumed successfully: " + name);
}

// Delete connector
public void deleteConnector(String name) throws IOException,
```

```
InterruptedException {
    HttpRequest request = HttpRequest.newBuilder()
        .uri(URI.create(connectUrl + "/connectors/" + name))
        .DELETE()
        .build();

    HttpResponse<String> response = httpClient.send(request,
        HttpResponse.BodyHandlers.ofString());

    if (response.statusCode() != 204) {
        throw new RuntimeException("Failed to delete connector: " +
response.body());
    }

    System.out.println("Connector deleted successfully: " + name);
}

// Get connector configuration
public Map<String, String> getConnectorConfig(String name)
    throws IOException, InterruptedException {

    HttpRequest request = HttpRequest.newBuilder()
        .uri(URI.create(connectUrl + "/connectors/" + name + "/config"))
        .GET()
        .build();

    HttpResponse<String> response = httpClient.send(request,
        HttpResponse.BodyHandlers.ofString());

    if (response.statusCode() == 200) {
        return objectMapper.readValue(response.body(),
            objectMapper.getTypeFactory().constructMapType(Map.class,
String.class, String.class));
    } else {
        throw new RuntimeException("Failed to get connector config: " +
response.body());
    }
}

// Usage example
public static void main(String[] args) {
    KafkaConnectRestClient client = new
KafkaConnectRestClient("http://localhost:8083");

    try {
        // List existing connectors
        String[] connectors = client.listConnectors();
        System.out.println("Existing connectors: " + String.join(", ", connectors));

        // Create JDBC source connector
        Map<String, String> jdbcConfig =
SourceConnectorConfig.createJDBCSourceConfig();
        client.createOrUpdateConnector("postgres-source", jdbcConfig);
    }
}
```

```
// Check status
ConnectorStatus status = client.getConnectorStatus("postgres-source");
System.out.println("Connector status: " + status);

// Create Elasticsearch sink connector
Map<String, String> esConfig =
SinkConnectorConfig.createElasticsearchSinkConfig();
client.createOrUpdateConnector("elasticsearch-sink", esConfig);

} catch (Exception e) {
    System.err.println("Error managing connectors: " + e.getMessage());
    e.printStackTrace();
}
}

// Helper class for connector status
public static class ConnectorStatus {
    private final String name;
    private final String state;
    private final int taskCount;

    public ConnectorStatus(String name, String state, int taskCount) {
        this.name = name;
        this.state = state;
        this.taskCount = taskCount;
    }

    @Override
    public String toString() {
        return String.format("ConnectorStatus{name='%s', state='%s',
tasks=%d}", name, state, taskCount);
    }

    // Getters
    public String getName() { return name; }
    public String getState() { return state; }
    public int getTaskCount() { return taskCount; }
}
}
```

---

## 🔗 Common Connectors

### JDBC Source/Sink

#### JDBC Source Connector - Database to Kafka

```
import java.util.Map;
import java.util.HashMap;
```

```
public class JDBCSourceConnectorExample {

    public static Map<String, String> createIncrementalSourceConfig() {
        Map<String, String> config = new HashMap<>();

        // Connector identification
        config.put("name", "postgres-users-source");
        config.put("connector.class",
                "io.confluent.connect.jdbc.JdbcSourceConnector");

        // Database connection
        config.put("connection.url",
                "jdbc:postgresql://localhost:5432/ecommerce");
        config.put("connection.user", "postgres");
        config.put("connection.password", "password");
        config.put("connection.attempts", "3");
        config.put("connection.backoff.ms", "10000");

        // Mode configuration
        config.put("mode", "incrementing");
        config.put("incrementing.column.name", "id");

        // Table configuration
        config.put("table.whitelist", "users");
        config.put("catalog.pattern", null);
        config.put("schema.pattern", "public");

        // Topic configuration
        config.put("topic.prefix", "postgres-");
        config.put("topic.creation.enable", "true");

        // Polling configuration
        config.put("poll.interval.ms", "5000");
        config.put("batch.max.rows", "100");
        config.put("table.poll.interval.ms", "60000");

        // Query customization
        config.put("query", "SELECT * FROM users WHERE id > ? ORDER BY id");
        config.put("validate.non.null", "false");

        return config;
    }

    public static Map<String, String> createTimestampSourceConfig() {
        Map<String, String> config = new HashMap<>();

        // Basic configuration
        config.put("name", "postgres-orders-source");
        config.put("connector.class",
                "io.confluent.connect.jdbc.JdbcSourceConnector");

        // Connection
        config.put("connection.url",
```

```
"jdbc:postgresql://localhost:5432/ecommerce");
    config.put("connection.user", "postgres");
    config.put("connection.password", "password");

    // Timestamp mode for CDC
    config.put("mode", "timestamp+incrementing");
    config.put("timestamp.column.name", "updated_at");
    config.put("incrementing.column.name", "id");

    // Table configuration
    config.put("table.whitelist", "orders");
    config.put("topic.prefix", "postgres-");

    // Polling configuration
    config.put("poll.interval.ms", "1000");
    config.put("batch.max.rows", "500");

    // Timestamp configuration
    config.put("timestamp.delay.interval.ms", "1000");
    config.put("db.timezone", "UTC");

    return config;
}

// Custom query example
public static Map<String, String> createCustomQuerySourceConfig() {
    Map<String, String> config = new HashMap<>();

    config.put("name", "postgres-analytics-source");
    config.put("connector.class",
"io.confluent.connect.jdbc.JdbcSourceConnector");

    // Connection
    config.put("connection.url",
"jdbc:postgresql://localhost:5432/analytics");
    config.put("connection.user", "analytics_user");
    config.put("connection.password", "password");

    // Custom query mode
    config.put("mode", "bulk");
    config.put("query",
"SELECT u.id, u.email, COUNT(o.id) as order_count, " +
"      SUM(o.total) as total_spent " +
"FROM users u " +
"LEFT JOIN orders o ON u.id = o.user_id " +
"WHERE u.updated_at > ? " +
"GROUP BY u.id, u.email");

    config.put("topic.prefix", "analytics-");
    config.put("poll.interval.ms", "300000"); // 5 minutes

    return config;
}
}
```

## JDBC Sink Connector - Kafka to Database

```
public class JDBCConnectorExample {

    public static Map<String, String> createPostgresSinkConfig() {
        Map<String, String> config = new HashMap<>();

        // Connector identification
        config.put("name", "postgres-users-sink");
        config.put("connector.class",
        "io.confluent.connect.jdbc.JdbcSinkConnector");

        // Database connection
        config.put("connection.url",
        "jdbc:postgresql://localhost:5432/warehouse");
        config.put("connection.user", "warehouse_user");
        config.put("connection.password", "password");

        // Topics configuration
        config.put("topics", "postgres-users,postgres-orders");
        config.put("topics.regex", "postgres-.*");

        // Table configuration
        config.put("table.name.format", "kafka_${topic}");
        config.put("auto.create", "true");
        config.put("auto.evolve", "true");

        // Insert mode
        config.put("insert.mode", "insert");
        config.put("batch.size", "3000");
        config.put("max.retries", "10");
        config.put("retry.backoff.ms", "3000");

        // Field mapping
        config.put("fields.whitelist", "id,name,email,created_at");
        config.put("pk.mode", "record_key");
        config.put("pk.fields", "id");

        // Data type mapping
        config.put("db.timezone", "UTC");

        return config;
    }

    public static Map<String, String> createUpsertSinkConfig() {
        Map<String, String> config = new HashMap<>();

        config.put("name", "postgres-upsert-sink");
        config.put("connector.class",
        "io.confluent.connect.jdbc.JdbcSinkConnector");
```

```
// Connection
config.put("connection.url",
"jdbc:postgresql://localhost:5432/warehouse");
config.put("connection.user", "warehouse_user");
config.put("connection.password", "password");

// Topics
config.put("topics", "user-updates");

// Upsert configuration
config.put("insert.mode", "upsert");
config.put("pk.mode", "record_value");
config.put("pk.fields", "id");

// Table configuration
config.put("table.name.format", "users");
config.put("auto.create", "false"); // Table should exist
config.put("auto.evolve", "true");

// Batching
config.put("batch.size", "1000");

return config;
}

public static void demonstrateJDBCConnectors() {
    KafkaConnectRestClient client = new
KafkaConnectRestClient("http://localhost:8083");

    try {
        // Create source connector
        Map<String, String> sourceConfig = createIncrementalSourceConfig();
        client.createOrUpdateConnector("postgres-users-source", sourceConfig);

        // Create sink connector
        Map<String, String> sinkConfig = createPostgresSinkConfig();
        client.createOrUpdateConnector("postgres-users-sink", sinkConfig);

        // Monitor status
        Thread.sleep(5000);
        ConnectorStatus sourceStatus = client.getConnectorStatus("postgres-
users-source");
        ConnectorStatus sinkStatus = client.getConnectorStatus("postgres-
users-sink");

        System.out.println("Source connector: " + sourceStatus);
        System.out.println("Sink connector: " + sinkStatus);

    } catch (Exception e) {
        System.err.println("Error setting up JDBC connectors: " +
e.getMessage());
    }
}
```

```
    }  
}
```

## Debezium CDC

### Debezium MySQL Connector

```
public class DebeziumMySQLConnectorExample {  
  
    public static Map<String, String> createMySQLCDCConfig() {  
        Map<String, String> config = new HashMap<>();  
  
        // Connector identification  
        config.put("name", "mysql-cdc-connector");  
        config.put("connector.class",  
"io.debezium.connector.mysql.MySqlConnector");  
  
        // Database connection  
        config.put("database.hostname", "localhost");  
        config.put("database.port", "3306");  
        config.put("database.user", "debezium");  
        config.put("database.password", "password");  
        config.put("database.server.id", "184054");  
        config.put("database.server.name", "ecommerce-mysql");  
  
        // Database selection  
        config.put("database.include.list", "ecommerce");  
        config.put("table.include.list",  
"ecommerce.users,ecommerce.orders,ecommerce.products");  
  
        // Binlog configuration  
        config.put("database.history.kafka.bootstrap.servers", "localhost:9092");  
        config.put("database.history.kafka.topic", "ecommerce-mysql-history");  
  
        // Topic routing  
        config.put("topic.prefix", "mysql-cdc");  
  
        // Snapshot configuration  
        config.put("snapshot.mode", "initial");  
        config.put("snapshot.locking.mode", "minimal");  
        config.put("snapshot.new.tables", "parallel");  
  
        // Schema changes  
        config.put("schema.history.internal.kafka.bootstrap.servers",  
"localhost:9092");  
        config.put("schema.history.internal.kafka.topic", "mysql-schema-changes");  
  
        // Message key configuration  
        config.put("message.key.columns",  
"ecommerce.users:id;ecommerce.orders:id");  
    }  
}
```

```
// Transforms (see SMT section)
config.put("transforms", "unwrap");
config.put("transforms.unwrap.type",
"io.debezium.transforms.ExtractNewRecordState");

return config;
}

public static Map<String, String> createPostgreSQLCDCConfig() {
Map<String, String> config = new HashMap<>();

// Connector identification
config.put("name", "postgres-cdc-connector");
config.put("connector.class",
"io.debezium.connector.postgresql.PostgresConnector");

// Database connection
config.put("database.hostname", "localhost");
config.put("database.port", "5432");
config.put("database.user", "postgres");
config.put("database.password", "password");
config.put("database dbname", "ecommerce");
config.put("database.server.name", "ecommerce-postgres");

// WAL configuration
config.put("slot.name", "debezium_slot");
config.put("plugin.name", "pgoutput");

// Schema and table filtering
config.put("schema.include.list", "public");
config.put("table.include.list",
"public.users,public.orders,public.order_items");

// Topic configuration
config.put("topic.prefix", "postgres-cdc");

// Snapshot configuration
config.put("snapshot.mode", "initial");

// Publication configuration (PostgreSQL 10+)
config.put("publication.name", "debezium_publication");
config.put("publication.autocreate.mode", "filtered");

return config;
}

// CDC Event Processing Example
public static void processCDCEvents() {
Properties props = new Properties();
props.put("bootstrap.servers", "localhost:9092");
props.put("group.id", "cdc-processor");
props.put("key.deserializer",
"org.apache.kafka.common.serialization.StringDeserializer");
props.put("value.deserializer",

```

```
"org.apache.kafka.common.serialization.StringDeserializer");
props.put("auto.offset.reset", "earliest");

try (KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props))
{
    consumer.subscribe(java.util.Arrays.asList(
        "mysql-cdc.ecommerce.users",
        "mysql-cdc.ecommerce.orders"
    ));

    while (true) {
        ConsumerRecords<String, String> records =
consumer.poll(Duration.ofMillis(1000));

        for (ConsumerRecord<String, String> record : records) {
            processCDCRecord(record);
        }
    }
} catch (Exception e) {
    System.err.println("Error processing CDC events: " + e.getMessage());
}

private static void processCDCRecord(ConsumerRecord<String, String> record) {
    try {
        ObjectMapper mapper = new ObjectMapper();
        JsonNode event = mapper.readTree(record.value());

        String operation = event.has("op") ? event.get("op").asText() :
"unknown";
        String table = extractTableName(record.topic());

        System.out.printf("CDC Event: table=%s, operation=%s, key=%s%n",
            table, operation, record.key());

        switch (operation) {
            case "c": // Create
                handleInsert(table, event.get("after"));
                break;
            case "u": // Update
                handleUpdate(table, event.get("before"), event.get("after"));
                break;
            case "d": // Delete
                handleDelete(table, event.get("before"));
                break;
            case "r": // Read (snapshot)
                handleSnapshot(table, event.get("after"));
                break;
            default:
                System.out.println("Unknown operation: " + operation);
        }

    } catch (Exception e) {
        System.err.println("Error processing CDC record: " + e.getMessage());
    }
}
```

```

        }

    }

    private static String extractTableName(String topic) {
        // Topic format: mysql-cdc.database.table
        String[] parts = topic.split("\\.");
        return parts.length >= 3 ? parts[2] : "unknown";
    }

    private static void handleInsert(String table, JsonNode after) {
        System.out.println("INSERT into " + table + ": " + after);
        // Implement your insert logic
    }

    private static void handleUpdate(String table, JsonNode before, JsonNode after) {
        System.out.println("UPDATE " + table + " from " + before + " to " + after);
        // Implement your update logic
    }

    private static void handleDelete(String table, JsonNode before) {
        System.out.println("DELETE from " + table + ": " + before);
        // Implement your delete logic
    }

    private static void handleSnapshot(String table, JsonNode data) {
        System.out.println("SNAPSHOT " + table + ": " + data);
        // Implement your snapshot logic
    }
}

```

## Elasticsearch/S3 Connectors

### Elasticsearch Sink Connector

```

public class ElasticsearchConnectorExample {

    public static Map<String, String> createElasticsearchSinkConfig() {
        Map<String, String> config = new HashMap<>();

        // Connector identification
        config.put("name", "elasticsearch-sink-connector");
        config.put("connector.class",
        "io.confluent.connect.elasticsearch.ElasticsearchSinkConnector");

        // Elasticsearch connection
        config.put("connection.url", "http://localhost:9200");
        config.put("connection.username", "elastic");
        config.put("connection.password", "changeme");
    }
}

```

```
// SSL configuration (if needed)
config.put("connection.ssl.enabled", "false");
config.put("connection.ssl.keystore.location", "/path/to/keystore.jks");
config.put("connection.ssl.keystore.password", "keystore-password");
config.put("connection.ssl.truststore.location",
"/path/to/truststore.jks");
config.put("connection.ssl.truststore.password", "truststore-password");

// Topics and routing
config.put("topics", "mysql-cdc.ecommerce.users,mysql-
cdc.ecommerce.orders");
config.put("topic.index.map",
"mysql-cdc.ecommerce.users:users,mysql-cdc.ecommerce.orders:orders");

// Document configuration
config.put("type.name", "_doc");
config.put("key.ignore", "false");
config.put("schema.ignore", "true");

// Batch configuration
config.put("batch.size", "2000");
config.put("max.in.flight.requests", "5");
config.put("max.buffered.records", "20000");
config.put("linger.ms", "1");
config.put("flush.timeout.ms", "10000");

// Error handling
config.put("behavior.on.null.values", "ignore");
config.put("behavior.on.malformed.documents", "warn");
config.put("drop.invalid.message", "true");

// Mapping configuration
config.put("schema.ignore", "true");
config.put("compact.map.entries", "true");

return config;
}

public static Map<String, String> createAdvancedElasticsearchConfig() {
Map<String, String> config = createElasticsearchSinkConfig();

// Advanced routing based on record content
config.put("topic.schema.ignore", "true");

// Custom transformations
config.put("transforms", "addTimestamp,extractKey");
config.put("transforms.addTimestamp.type",
"org.apache.kafka.connect.transforms.InsertField$Value");
config.put("transforms.addTimestamp.timestamp.field", "indexed_at");

config.put("transforms.extractKey.type",
"org.apache.kafka.connect.transforms.ValueToKey");
config.put("transforms.extractKey.fields", "id");
```

```

    // Dynamic index names based on timestamp
    config.put("topic.index.map",
        "user-events:user-events-{yyyy.MM.dd},order-events:order-events-
{yyyy.MM.dd}");

    // Retry configuration
    config.put("max.retries", "5");
    config.put("retry.backoff.ms", "5000");

    return config;
}

// Elasticsearch monitoring example
public static void monitorElasticsearchConnector(String connectorName) {
    KafkaConnectRestClient client = new
KafkaConnectRestClient("http://localhost:8083");

    try {
        // Get connector status
        ConnectorStatus status = client.getConnectorStatus(connectorName);
        System.out.println("Elasticsearch connector status: " + status);

        // Check if connector is healthy
        if (!"RUNNING".equals(status.getState())) {
            System.err.println("Connector is not running: " +
status.getState());

            // Try restarting
            client.restartConnector(connectorName);
            System.out.println("Attempted to restart connector");
        }
    } catch (Exception e) {
        System.err.println("Error monitoring Elasticsearch connector: " +
e.getMessage());
    }
}
}

```

## S3 Sink Connector

```

public class S3ConnectorExample {

    public static Map<String, String> createS3SinkConfig() {
        Map<String, String> config = new HashMap<>();

        // Connector identification
        config.put("name", "s3-sink-connector");
        config.put("connector.class", "io.confluent.connect.s3.S3SinkConnector");

        // AWS configuration

```

```
config.put("aws.access.key.id", "YOUR_ACCESS_KEY");
config.put("aws.secret.access.key", "YOUR_SECRET_KEY");
config.put("s3.region", "us-west-2");
config.put("s3.bucket.name", "kafka-connect-s3");

// Topics configuration
config.put("topics", "mysql-cdc.ecommerce.users,mysql-
cdc.ecommerce.orders");
config.put("topics.dir", "topics");

// Partitioning
config.put("s3.part.size", "5242880"); // 5MB
config.put("flush.size", "10000");
config.put("rotate.interval.ms", "3600000"); // 1 hour
config.put("rotate.schedule.interval.ms", "3600000");

// File format
config.put("format.class",
"io.confluent.connect.s3.format.json.JsonFormat");
config.put("partitioner.class",
"io.confluent.connect.storage.partitionер.TimeBasedPartitioner");
config.put("partition.duration.ms", "3600000"); // 1 hour partitions
config.put("path.format", "'year'=YYYY/'month'=MM/'day'=dd/'hour'=HH'");
config.put("locale", "US");
config.put("timezone", "UTC");

// Compression
config.put("s3.compression.type", "gzip");

// Schema registry (if using Avro)
config.put("schema.registry.url", "http://localhost:8081");

return config;
}

public static Map<String, String> createParquetS3SinkConfig() {
Map<String, String> config = new HashMap<>();

// Basic configuration
config.put("name", "s3-parquet-sink");
config.put("connector.class", "io.confluent.connect.s3.S3SinkConnector");

// AWS configuration
config.put("aws.access.key.id", "YOUR_ACCESS_KEY");
config.put("aws.secret.access.key", "YOUR_SECRET_KEY");
config.put("s3.region", "us-west-2");
config.put("s3.bucket.name", "data-lake-parquet");

// Topics
config.put("topics", "user-analytics,order-analytics");

// Parquet format
config.put("format.class",
"io.confluent.connect.s3.format.parquet.ParquetFormat");
```

```
config.put("schema.registry.url", "http://localhost:8081");
config.put("schema.compatibility", "BACKWARD");

// Partitioning for analytics
config.put("partitioner.class",
"io.confluent.connect.storage.partitioner.FieldPartitioner");
config.put("partition.field.name", "department,region");

// Larger flush sizes for analytics workloads
config.put("flush.size", "50000");
config.put("rotate.interval.ms", "7200000"); // 2 hours

// Compression
config.put("s3.compression.type", "snappy");

return config;
}

// S3 file organization example
public static void demonstrateS3Organization() {
/*
 * S3 Directory Structure Example:
 *
 * kafka-connect-s3/
 *   └── topics/
 *     ├── mysql-cdc.ecommerce.users/
 *     |   ├── year=2025/month=09/day=21/hour=10/
 *     |   |   ├── mysql-cdc.ecommerce.users+0+0000000000.json.gz
 *     |   |   └── mysql-cdc.ecommerce.users+0+0000010000.json.gz
 *     |   └── year=2025/month=09/day=21/hour=11/
 *     |       └── mysql-cdc.ecommerce.users+0+0000020000.json.gz
 *     └── mysql-cdc.ecommerce.orders/
 *         └── year=2025/month=09/day=21/hour=10/
 *             └── mysql-cdc.ecommerce.orders+0+0000000000.json.gz
 *
 * File naming convention:
 * {topic}+{partition}+{start_offset}.{format}.{compression}
 */

System.out.println("S3 connector will organize files by:");
System.out.println("1. Topic name");
System.out.println("2. Time-based partitioning (year/month/day/hour)");
System.out.println("3. Partition and offset information");
System.out.println("4. Configurable compression (gzip, snappy, etc.)");
}
}
```

---

## 🛠️ Customization

### Writing Custom Connectors

## Custom Source Connector Example

```
import org.apache.kafka.connect.source.SourceConnector;
import org.apache.kafka.connect.source.SourceTask;
import org.apache.kafka.connect.connector.Task;
import org.apache.kafka.common.config.ConfigDef;
import org.apache.kafka.common.config.AbstractConfig;
import java.util.Map;
import java.util.List;
import java.util.ArrayList;

public class CustomFileSourceConnector extends SourceConnector {

    private Map<String, String> configProps;

    @Override
    public String version() {
        return "1.0.0";
    }

    @Override
    public void start(Map<String, String> props) {
        this.configProps = props;

        // Validate configuration
        CustomFileSourceConfig config = new CustomFileSourceConfig(props);

        // Additional startup logic
        System.out.println("Starting Custom File Source Connector");
        System.out.println("Monitoring directory: " + config.getDirectoryPath());
        System.out.println("Topic: " + config.getTopic());
    }

    @Override
    public Class<? extends Task> taskClass() {
        return CustomFileSourceTask.class;
    }

    @Override
    public List<Map<String, String>> taskConfigs(int maxTasks) {
        List<Map<String, String>> configs = new ArrayList<>();

        // For simplicity, create one task
        // In production, you might split work across multiple tasks
        configs.add(configProps);

        return configs;
    }

    @Override
    public void stop() {
        System.out.println("Stopping Custom File Source Connector");
    }
}
```

```
// Cleanup logic
}

@Override
public ConfigDef config() {
    return CustomFileSourceConfig.CONFIG_DEF;
}

// Configuration class
public static class CustomFileSourceConfig extends AbstractConfig {

    public static final String DIRECTORY_PATH_CONFIG = "directory.path";
    public static final String TOPIC_CONFIG = "topic";
    public static final String POLL_INTERVAL_MS_CONFIG = "poll.interval.ms";
    public static final String FILE_PATTERN_CONFIG = "file.pattern";

    public static final ConfigDef CONFIG_DEF = new ConfigDef()
        .define(DIRECTORY_PATH_CONFIG,
            ConfigDef.Type.STRING,
            ConfigDef.Importance.HIGH,
            "Directory path to monitor for files")
        .define(TOPIC_CONFIG,
            ConfigDef.Type.STRING,
            ConfigDef.Importance.HIGH,
            "Kafka topic to publish messages to")
        .define(POLL_INTERVAL_MS_CONFIG,
            ConfigDef.Type.LONG,
            5000L,
            ConfigDef.Importance.MEDIUM,
            "Polling interval in milliseconds")
        .define(FILE_PATTERN_CONFIG,
            ConfigDef.Type.STRING,
            ".*\.\txt$",
            ConfigDef.Importance.MEDIUM,
            "Regex pattern for files to process");

    public CustomFileSourceConfig(Map<String, String> props) {
        super(CONFIG_DEF, props);
    }

    public String getDirectoryPath() {
        return getString(DIRECTORY_PATH_CONFIG);
    }

    public String getTopic() {
        return getString(TOPIC_CONFIG);
    }

    public Long getPollInterval() {
        return getLong(POLL_INTERVAL_MS_CONFIG);
    }

    public String getFilePattern() {
        return getString(FILE_PATTERN_CONFIG);
    }
}
```

```
        }
    }
}
```

## Custom Source Task

```
import org.apache.kafka.connect.source.SourceTask;
import org.apache.kafka.connect.source.SourceRecord;
import org.apache.kafka.connect.data.Schema;
import org.apache.kafka.connect.data.SchemaBuilder;
import org.apache.kafka.connect.data.Struct;

import java.io.*;
import java.nio.file.*;
import java.util.*;
import java.util.concurrent.ConcurrentHashMap;
import java.util.regex.Pattern;

public class CustomFileSourceTask extends SourceTask {

    private CustomFileSourceConnector.CustomFileSourceConfig config;
    private Path directoryPath;
    private Pattern filePattern;
    private Map<String, Long> fileOffsets;
    private long lastPollTime;

    private static final String FILENAME_FIELD = "filename";
    private static final String CONTENT_FIELD = "content";
    private static final String LINE_NUMBER_FIELD = "line_number";
    private static final String TIMESTAMP_FIELD = "timestamp";

    private static final Schema VALUE_SCHEMA = SchemaBuilder.struct()
        .name("file.record")
        .field(FILENAME_FIELD, Schema.STRING_SCHEMA)
        .field(CONTENT_FIELD, Schema.STRING_SCHEMA)
        .field(LINE_NUMBER_FIELD, Schema.INT64_SCHEMA)
        .field(TIMESTAMP_FIELD, Schema.INT64_SCHEMA)
        .build();

    @Override
    public String version() {
        return "1.0.0";
    }

    @Override
    public void start(Map<String, String> props) {
        config = new CustomFileSourceConnector.CustomFileSourceConfig(props);
        directoryPath = Paths.get(config.getDirectoryPath());
        filePattern = Pattern.compile(config.getFilePattern());
        fileOffsets = new ConcurrentHashMap<>();
        lastPollTime = System.currentTimeMillis();
    }

    @Override
    public void stop() {
        // Implementation
    }

    @Override
    public void seek(SourceRecord record, long offset) {
        fileOffsets.put(record.key().toString(), offset);
    }

    @Override
    public void poll() {
        // Implementation
    }

    @Override
    public void commit() {
        // Implementation
    }

    @Override
    public void close() {
        // Implementation
    }
}
```

```
// Load existing offsets from Kafka Connect offset storage
loadOffsets();

System.out.println("Started Custom File Source Task");
System.out.println("Monitoring: " + directoryPath);
}

@Override
public List<SourceRecord> poll() throws InterruptedException {
    List<SourceRecord> records = new ArrayList<>();

    try {
        // Sleep to avoid excessive polling
        Thread.sleep(config.getPollInterval());

        // Scan directory for files
        try (DirectoryStream<Path> stream =
Files.newDirectoryStream(directoryPath)) {
            for (Path filePath : stream) {
                if (shouldProcessFile(filePath)) {
                    records.addAll(processFile(filePath));
                }
            }
        }
    } catch (IOException e) {
        System.err.println("Error polling directory: " + e.getMessage());
    }

    return records;
}

private boolean shouldProcessFile(Path filePath) {
    if (Files.isDirectory(filePath)) {
        return false;
    }

    String filename = filePath.getFileName().toString();
    if (!filePattern.matcher(filename).matches()) {
        return false;
    }

    // Check if file has been modified since last poll
    try {
        long lastModified = Files.getLastModifiedTime(filePath).toMillis();
        return lastModified > lastPollTime;
    } catch (IOException e) {
        System.err.println("Error checking file modification time: " +
e.getMessage());
        return false;
    }
}
```

```
private List<SourceRecord> processFile(Path filePath) {
    List<SourceRecord> records = new ArrayList<>();
    String filename = filePath.getFileName().toString();

    try {
        Long currentOffset = fileOffsets.getOrDefault(filename, 0L);

        try (BufferedReader reader = Files.newBufferedReader(filePath)) {
            String line;
            long lineNumber = 0;

            // Skip lines we've already processed
            while (lineNumber < currentOffset && (line = reader.readLine()) != null) {
                lineNumber++;
            }

            // Process new lines
            while ((line = reader.readLine()) != null) {
                lineNumber++;

                // Create source record
                SourceRecord record = createSourceRecord(filename, line,
                    lineNumber);
                records.add(record);

                // Update offset
                fileOffsets.put(filename, lineNumber);
            }
        }
    } catch (IOException e) {
        System.err.println("Error processing file " + filename + ": " +
            e.getMessage());
    }

    return records;
}

private SourceRecord createSourceRecord(String filename, String content, long
lineNumber) {
    // Source partition - identifies the file
    Map<String, Object> sourcePartition = new HashMap<>();
    sourcePartition.put("filename", filename);

    // Source offset - line number in file
    Map<String, Object> sourceOffset = new HashMap<>();
    sourceOffset.put("position", lineNumber);

    // Record value
    Struct value = new Struct(VALUE_SCHEMA)
        .put(FILENAME_FIELD, filename)
        .put(CONTENT_FIELD, content)
        .put(LINE_NUMBER_FIELD, lineNumber)
```

```

    .put(TIMESTAMP_FIELD, System.currentTimeMillis());

    return new SourceRecord(
        sourcePartition,           // Source partition
        sourceOffset,             // Source offset
        config.getTopic(),        // Topic
        null,                    // Partition (let Kafka decide)
        Schema.STRING_SCHEMA,    // Key schema
        filename,                // Key (filename)
        VALUE_SCHEMA,            // Value schema
        value                    // Value
    );
}

private void loadOffsets() {
    // In a real implementation, you would load offsets from Connect's offset
    storage
    // For this example, we'll start from the beginning
    System.out.println("Loading offsets from Connect offset storage");
}

@Override
public void stop() {
    System.out.println("Stopping Custom File Source Task");
    // Cleanup resources
}
}

```

## Custom Sink Connector Example

```

import org.apache.kafka.connect.sink.SinkConnector;
import org.apache.kafka.connect.sink.SinkTask;
import org.apache.kafka.connect.connector.Task;
import org.apache.kafka.common.config.ConfigDef;
import org.apache.kafka.common.config.AbstractConfig;

public class CustomFileSinkConnector extends SinkConnector {

    private Map<String, String> configProps;

    @Override
    public String version() {
        return "1.0.0";
    }

    @Override
    public void start(Map<String, String> props) {
        this.configProps = props;

        CustomFileSinkConfig config = new CustomFileSinkConfig(props);
        System.out.println("Starting Custom File Sink Connector");
    }
}

```

```
        System.out.println("Output directory: " + config.getOutputDirectory());
    }

    @Override
    public Class<? extends Task> taskClass() {
        return CustomFileSinkTask.class;
    }

    @Override
    public List<Map<String, String>> taskConfigs(int maxTasks) {
        List<Map<String, String>> configs = new ArrayList<>();

        // Create task configurations
        for (int i = 0; i < maxTasks; i++) {
            Map<String, String> taskConfig = new HashMap<>(configProps);
            taskConfig.put("task.id", String.valueOf(i));
            configs.add(taskConfig);
        }

        return configs;
    }

    @Override
    public void stop() {
        System.out.println("Stopping Custom File Sink Connector");
    }

    @Override
    public ConfigDef config() {
        return CustomFileSinkConfig.CONFIG_DEF;
    }

    // Configuration class
    public static class CustomFileSinkConfig extends AbstractConfig {

        public static final String OUTPUT_DIRECTORY_CONFIG = "output.directory";
        public static final String FILE_EXTENSION_CONFIG = "file.extension";
        public static final String FLUSH_SIZE_CONFIG = "flush.size";
        public static final String ROTATE_INTERVAL_MS_CONFIG =
"rotate.interval.ms";

        public static final ConfigDef CONFIG_DEF = new ConfigDef()
            .define(OUTPUT_DIRECTORY_CONFIG,
                    ConfigDef.Type.STRING,
                    ConfigDef.Importance.HIGH,
                    "Directory to write output files")
            .define(FILE_EXTENSION_CONFIG,
                    ConfigDef.Type.STRING,
                    ".txt",
                    ConfigDef.Importance.MEDIUM,
                    "File extension for output files")
            .define(FLUSH_SIZE_CONFIG,
                    ConfigDef.Type.INT,
                    1000,
                    ConfigDef.Importance.LOW,
                    "Size at which to flush output files")
    }
}
```

```
        ConfigDef.Importance.MEDIUM,
        "Number of records before flushing to file")
    .define(ROTATE_INTERVAL_MS_CONFIG,
        ConfigDef.Type.LONG,
        3600000L, // 1 hour
        ConfigDef.Importance.MEDIUM,
        "File rotation interval in milliseconds");

    public CustomFileSinkConfig(Map<String, String> props) {
        super(CONFIG_DEF, props);
    }

    public String getOutputDirectory() {
        return getString(OUTPUT_DIRECTORY_CONFIG);
    }

    public String getFileExtension() {
        return getString(FILE_EXTENSION_CONFIG);
    }

    public Integer getFlushSize() {
        return getInt(FLUSH_SIZE_CONFIG);
    }

    public Long getRotateInterval() {
        return getLong(ROTATE_INTERVAL_MS_CONFIG);
    }
}
```

## Custom Sink Task

```
import org.apache.kafka.connect.sink.SinkTask;
import org.apache.kafka.connect.sink.SinkRecord;
import java.io.*;
import java.nio.file.*;
import java.util.*;

public class CustomFileSinkTask extends SinkTask {

    private CustomFileSinkConnector.CustomFileSinkConfig config;
    private Path outputDirectory;
    private Map<String, BufferedWriter> writers;
    private Map<String, Integer> recordCounts;
    private Map<String, Long> lastRotateTime;

    @Override
    public String version() {
        return "1.0.0";
    }
}
```

```
@Override
public void start(Map<String, String> props) {
    config = new CustomFileSinkConnector.CustomFileSinkConfig(props);
    outputDirectory = Paths.get(config.getOutputDirectory());
    writers = new HashMap<>();
    recordCounts = new HashMap<>();
    lastRotateTime = new HashMap<>();

    // Create output directory if it doesn't exist
    try {
        Files.createDirectories(outputDirectory);
    } catch (IOException e) {
        throw new RuntimeException("Failed to create output directory", e);
    }

    System.out.println("Started Custom File Sink Task");
    System.out.println("Output directory: " + outputDirectory);
}

@Override
public void put(Collection<SinkRecord> records) {
    for (SinkRecord record : records) {
        try {
            writeRecord(record);
        } catch (IOException e) {
            System.err.println("Error writing record: " + e.getMessage());
            // In production, you might want to implement retry logic
        }
    }

    // Check if any files need rotation
    checkFileRotation();
}

private void writeRecord(SinkRecord record) throws IOException {
    String topic = record.topic();
    String filename = generateFilename(topic);

    BufferedWriter writer = getWriter(topic, filename);

    // Write record to file
    String recordLine = formatRecord(record);
    writer.write(recordLine);
    writer.newLine();

    // Update record count
    recordCounts.put(topic, recordCounts.getOrDefault(topic, 0) + 1);

    // Check if we need to flush
    if (recordCounts.get(topic) >= config.getFlushSize()) {
        writer.flush();
        recordCounts.put(topic, 0);
    }
}
```

```
private BufferedWriter getWriter(String topic, String filename) throws
IOException {
    if (!writers.containsKey(topic)) {
        Path filePath = outputDirectory.resolve(filename);
        BufferedWriter writer = Files.newBufferedWriter(filePath,
            StandardOpenOption.CREATE, StandardOpenOption.APPEND);
        writers.put(topic, writer);
        lastRotateTime.put(topic, System.currentTimeMillis());
    }

    return writers.get(topic);
}

private String generateFilename(String topic) {
    long timestamp = System.currentTimeMillis();
    return String.format("%s-%d%s", topic, timestamp,
config.getFileExtension());
}

private String formatRecord(SinkRecord record) {
    // Format record as JSON or custom format
    return String.format(
 "{\"topic\":\"%s\", \"partition\":%d, \"offset\":%d, \"key\":\"%s\", \"value\":\"%s\", \
 \"timestamp\":%d}",
        record.topic(),
        record.kafkaPartition(),
        record.kafkaOffset(),
        record.key(),
        record.value(),
        record.timestamp());
}

private void checkFileRotation() {
    long currentTime = System.currentTimeMillis();

    for (Map.Entry<String, Long> entry : lastRotateTime.entrySet()) {
        String topic = entry.getKey();
        Long lastRotate = entry.getValue();

        if (currentTime - lastRotate > config.getRotateInterval()) {
            rotateFile(topic);
        }
    }
}

private void rotateFile(String topic) {
    try {
        BufferedWriter writer = writers.get(topic);
        if (writer != null) {
            writer.close();
            writers.remove(topic);
            lastRotateTime.put(topic, System.currentTimeMillis());
            recordCounts.put(topic, 0);
        }
    }
}
```

```

        System.out.println("Rotated file for topic: " + topic);
    }
} catch (IOException e) {
    System.err.println("Error rotating file for topic " + topic + ": " +
e.getMessage());
}
}

@Override
public void flush(Map<org.apache.kafka.common.TopicPartition,
org.apache.kafka.clients.consumer.OffsetAndMetadata> offsets) {
    // Flush all writers
    for (BufferedWriter writer : writers.values()) {
        try {
            writer.flush();
        } catch (IOException e) {
            System.err.println("Error flushing writer: " + e.getMessage());
        }
    }
}

@Override
public void stop() {
    // Close all writers
    for (BufferedWriter writer : writers.values()) {
        try {
            writer.close();
        } catch (IOException e) {
            System.err.println("Error closing writer: " + e.getMessage());
        }
    }

    writers.clear();
    recordCounts.clear();
    lastRotateTime.clear();

    System.out.println("Stopped Custom File Sink Task");
}
}
}

```

## Single Message Transforms (SMTs)

### Custom SMT Example

```

import org.apache.kafka.common.config.ConfigDef;
import org.apache.kafka.connect.connector.ConnectRecord;
import org.apache.kafka.connect.data.Schema;
import org.apache.kafka.connect.data.SchemaBuilder;
import org.apache.kafka.connect.data.Struct;
import org.apache.kafka.connect.transforms.Transformation;

```

```
import org.apache.kafka.connect.transforms.util.SimpleConfig;

import java.util.Map;
import java.util.regex.Pattern;

public class MaskSensitiveData<R extends ConnectRecord<R>> implements Transformation<R> {

    public static final String FIELD_NAME_CONFIG = "field.name";
    public static final String MASK_PATTERN_CONFIG = "mask.pattern";
    public static final String REPLACEMENT_CONFIG = "replacement";

    private static final ConfigDef CONFIG_DEF = new ConfigDef()
        .define(FIELD_NAME_CONFIG, ConfigDef.Type.STRING,
    ConfigDef.Importance.HIGH,
        "Field name to mask")
        .define(MASK_PATTERN_CONFIG, ConfigDef.Type.STRING, "\\d{4}-\\d{4}-\\d{4}-
\\d{4}",
            ConfigDef.Importance.MEDIUM, "Regex pattern to match sensitive data")
        .define(REPLACEMENT_CONFIG, ConfigDef.Type.STRING, "*****-*",
    ConfigDef.Importance.MEDIUM, "Replacement text for masked data");

    private String fieldName;
    private Pattern maskPattern;
    private String replacement;

    @Override
    public void configure(Map<String, ?> configs) {
        SimpleConfig config = new SimpleConfig(CONFIG_DEF, configs);
        fieldName = config.getString(FIELD_NAME_CONFIG);
        maskPattern = Pattern.compile(config.getString(MASK_PATTERN_CONFIG));
        replacement = config.getString(REPLACEMENT_CONFIG);
    }

    @Override
    public R apply(R record) {
        if (record.value() == null) {
            return record;
        }

        Object maskedValue = maskSensitiveData(record.value(),
    record.valueSchema());

        return record.newRecord(
            record.topic(),
            record.kafkaPartition(),
            record.keySchema(),
            record.key(),
            record.valueSchema(),
            maskedValue,
            record.timestamp()
        );
    }
}
```

```
private Object maskSensitiveData(Object value, Schema schema) {
    if (schema != null && schema.type() == Schema.Type.STRUCT) {
        return maskStructData((Struct) value, schema);
    } else if (value instanceof String) {
        return maskString((String) value);
    } else if (value instanceof Map) {
        return maskMapData((Map<String, Object>) value);
    }

    return value;
}

private Struct maskStructData(Struct struct, Schema schema) {
    Struct newStruct = new Struct(schema);

    for (org.apache.kafka.connect.data.Field field : schema.fields()) {
        Object fieldValue = struct.get(field);

        if (field.name().equals(fieldName) && fieldValue instanceof String) {
            String maskedValue = maskString((String) fieldValue);
            newStruct.put(field.name(), maskedValue);
        } else {
            newStruct.put(field.name(), fieldValue);
        }
    }

    return newStruct;
}

@SuppressWarnings("unchecked")
private Map<String, Object> maskMapData(Map<String, Object> map) {
    Map<String, Object> newMap = new HashMap<>(map);

    if (newMap.containsKey(fieldName)) {
        Object value = newMap.get(fieldName);
        if (value instanceof String) {
            newMap.put(fieldName, maskString((String) value));
        }
    }

    return newMap;
}

private String maskString(String input) {
    return maskPattern.matcher(input).replaceAll(replacement);
}

@Override
public ConfigDef config() {
    return CONFIG_DEF;
}

@Override
public void close() {
```

```
        // Cleanup if needed
    }
}
```

## Built-in SMTs Usage Examples

```
public class SMTExamples {

    // Example: Using SMTs in connector configuration
    public static Map<String, String> createConnectorWithSMTs() {
        Map<String, String> config = new HashMap<>();

        // Basic connector configuration
        config.put("name", "jdbc-source-with-transforms");
        config.put("connector.class",
        "io.confluent.connect.jdbc.JdbcSourceConnector");
        config.put("connection.url",
        "jdbc:postgresql://localhost:5432/ecommerce");
        config.put("table.whitelist", "users");
        config.put("topic.prefix", "postgres-");

        // Chain of Single Message Transforms
        config.put("transforms",
        "addTimestamp,maskSSN,renameFields,filterFields");

        // 1. Add timestamp field
        config.put("transforms.addTimestamp.type",
            "org.apache.kafka.connect.transforms.InsertField$Value");
        config.put("transforms.addTimestamp.timestamp.field", "extracted_at");

        // 2. Mask sensitive data (custom SMT)
        config.put("transforms.maskSSN.type", "com.example.MaskSensitiveData");
        config.put("transforms.maskSSN.field.name", "ssn");
        config.put("transforms.maskSSN.mask.pattern", "\d{3}-\d{2}-\d{4}");
        config.put("transforms.maskSSN.replacement", "***-*-*-*");

        // 3. Rename fields
        config.put("transforms.renameFields.type",
            "org.apache.kafka.connect.transforms.ReplaceField$Value");
        config.put("transforms.renameFields.renames",
        "first_name:firstName,last_name:lastName");

        // 4. Filter out fields
        config.put("transforms.filterFields.type",
            "org.apache.kafka.connect.transforms.ReplaceField$Value");
        config.put("transforms.filterFields.blacklist", "internal_id,temp_field");

        return config;
}

// Example: Complex transformation pipeline
```

```
public static Map<String, String> createComplexTransformPipeline() {
    Map<String, String> config = new HashMap<>();

    config.put("name", "complex-transform-pipeline");
    config.put("connector.class",
    "io.debezium.connector.mysql.MySqlConnector");

    // Database configuration
    config.put("database.hostname", "localhost");
    config.put("database.server.name", "ecommerce");
    config.put("table.include.list", "ecommerce.orders");

    // Complex transformation pipeline
    config.put("transforms",
        "unwrapCDC,extractKey,addMetadata,routeByRegion,formatCurrency");

    // 1. Unwrap Debezium CDC format
    config.put("transforms.unwrapCDC.type",
        "io.debezium.transforms.ExtractNewRecordState");
    config.put("transforms.unwrapCDC.drop.tombstones", "false");
    config.put("transforms.unwrapCDC.delete.handling.mode", "rewrite");

    // 2. Extract key from value
    config.put("transforms.extractKey.type",
        "org.apache.kafka.connect.transforms.ValueToKey");
    config.put("transforms.extractKey.fields", "order_id");

    // 3. Add metadata
    config.put("transforms.addMetadata.type",
        "org.apache.kafka.connect.transforms.InsertField$Value");
    config.put("transforms.addMetadata.static.field", "source");
    config.put("transforms.addMetadata.static.value", "mysql-orders");

    // 4. Route by region using header
    config.put("transforms.routeByRegion.type",
        "org.apache.kafka.connect.transforms.RegexRouter");
    config.put("transforms.routeByRegion.regex", ".*");
    config.put("transforms.routeByRegion.replacement",
    "orders-${header:region}");

    // 5. Format currency (custom SMT)
    config.put("transforms.formatCurrency.type",
    "com.example.FormatCurrency");
    config.put("transforms.formatCurrency.field.name", "total_amount");
    config.put("transforms.formatCurrency.currency.code", "USD");

    return config;
}

// Example: Conditional transformations
public static Map<String, String> createConditionalTransforms() {
    Map<String, String> config = new HashMap<>();

    config.put("name", "conditional-transforms");
```

```
    config.put("connector.class",
"io.confluent.connect.jdbc.JdbcSourceConnector");

    // Transformation chain with predicates
    config.put("transforms", "filterCustomers,addRegion");

    // 1. Filter only customers (with predicate)
    config.put("transforms.filterCustomers.type",
        "org.apache.kafka.connect.transforms.Filter");
    config.put("transforms.filterCustomers.predicate", "isCustomer");

    // 2. Add region field
    config.put("transforms.addRegion.type",
        "org.apache.kafka.connect.transforms.InsertField$Value");
    config.put("transforms.addRegion.static.field", "region");
    config.put("transforms.addRegion.static.value", "north-america");

    // Predicate configuration
    config.put("predicates", "isCustomer");
    config.put("predicates.isCustomer.type",
        "org.apache.kafka.connect.transforms.predicates.HasHeaderKey");
    config.put("predicates.isCustomer.name", "record_type");

    return config;
}

// Custom SMT for currency formatting
public static class FormatCurrency<R extends ConnectRecord<R>> implements
Transformation<R> {

    private String fieldName;
    private String currencyCode;

    @Override
    public void configure(Map<String, ?> configs) {
        fieldName = (String) configs.get("field.name");
        currencyCode = (String) configs.get("currency.code");
    }

    @Override
    public R apply(R record) {
        if (record.value() instanceof Struct) {
            Struct struct = (Struct) record.value();
            Object amount = struct.get(fieldName);

            if (amount instanceof Number) {
                String formattedAmount = String.format("%s %.2f",
                    currencyCode, ((Number) amount).doubleValue());

                Struct newStruct = new Struct(struct.schema());
                for (org.apache.kafka.connect.data.Field field :
struct.schema().fields()) {
                    if (field.name().equals(fieldName)) {
                        newStruct.put(field.name(), formattedAmount);
                    }
                }
                record.set_value(newStruct);
            }
        }
        return record;
    }
}
```

```
        } else {
            newStruct.put(field.name(), struct.get(field));
        }
    }

    return record.newRecord(
        record.topic(), record.kafkaPartition(),
        record.keySchema(), record.key(),
        record.valueSchema(), newStruct,
        record.timestamp()
    );
}

return record;
}

@Override
public ConfigDef config() {
    return new ConfigDef()
        .define("field.name", ConfigDef.Type.STRING,
ConfigDef.Importance.HIGH, "Field to format")
        .define("currency.code", ConfigDef.Type.STRING, "USD",
ConfigDef.Importance.MEDIUM, "Currency code");
}

@Override
public void close() {}
}
```

## ⌚ Comprehensive Java Examples

### Production Kafka Connect Management System

```
import java.util.concurrent.*;
import java.util.concurrent.atomic.AtomicBoolean;

/**
 * Production-ready Kafka Connect Management System
 * Handles connector lifecycle, monitoring, and health checks
 */
public class KafkaConnectManager {

    private final KafkaConnectRestClient restClient;
    private final ScheduledExecutorService scheduler;
    private final AtomicBoolean running;
    private final Map<String, ConnectorHealth> connectorHealth;

    public KafkaConnectManager(String connectUrl) {
```

```
this.restClient = new KafkaConnectRestClient(connectUrl);
this.scheduler = Executors.newScheduledThreadPool(3);
this.running = new AtomicBoolean(false);
this.connectorHealth = new ConcurrentHashMap<>();
}

public void start() {
    if (running.compareAndSet(false, true)) {
        // Start health monitoring
        scheduler.scheduleAtFixedRate(this::monitorConnectorHealth,
            30, 30, TimeUnit.SECONDS);

        // Start metrics collection
        scheduler.scheduleAtFixedRate(this::collectMetrics,
            60, 60, TimeUnit.SECONDS);

        // Start configuration backup
        scheduler.scheduleAtFixedRate(this::backupConfigurations,
            0, 1, TimeUnit.HOURS);

        System.out.println("Kafka Connect Manager started");
    }
}

public void stop() {
    if (running.compareAndSet(true, false)) {
        scheduler.shutdown();
        try {
            if (!scheduler.awaitTermination(30, TimeUnit.SECONDS)) {
                scheduler.shutdownNow();
            }
        } catch (InterruptedException e) {
            scheduler.shutdownNow();
            Thread.currentThread().interrupt();
        }
        System.out.println("Kafka Connect Manager stopped");
    }
}

// Deploy multiple connectors with dependency management
public void deployConnectorPipeline(List<ConnectorDeployment> deployments) {
    System.out.println("Deploying connector pipeline...");

    // Sort by dependency order
    deployments.sort((a, b) -> a.getOrder() - b.getOrder());

    for (ConnectorDeployment deployment : deployments) {
        try {
            deployConnectorWithRetry(deployment);

            // Wait for connector to be ready before deploying next
            waitForConnectorReady(deployment.getName(),
Duration.ofMinutes(5));
        }
    }
}
```

```
        } catch (Exception e) {
            System.err.printf("Failed to deploy connector %s: %s%n",
                deployment.getName(), e.getMessage());

            // Decide whether to continue or stop pipeline deployment
            if (deployment.isCritical()) {
                throw new RuntimeException("Critical connector deployment
failed", e);
            }
        }
    }

    System.out.println("Connector pipeline deployment completed");
}

private void deployConnectorWithRetry(ConnectorDeployment deployment)
    throws Exception {

    int maxRetries = 3;
    int attempt = 0;

    while (attempt < maxRetries) {
        try {
            restClient.createOrUpdateConnector(deployment.getName(),
                deployment.getConfig());

            System.out.printf("Successfully deployed connector: %s%n",
                deployment.getName());
            return;
        } catch (Exception e) {
            attempt++;
            if (attempt >= maxRetries) {
                throw e;
            }

            System.err.printf("Deployment attempt %d failed for %s: %s%n",
                attempt, deployment.getName(), e.getMessage());

            // Exponential backoff
            Thread.sleep(1000 * (1L << attempt));
        }
    }
}

private void waitForConnectorReady(String connectorName, Duration timeout)
    throws Exception {

    long endTime = System.currentTimeMillis() + timeout.toMillis();

    while (System.currentTimeMillis() < endTime) {
        try {
            ConnectorStatus status =
restClient.getConnectorStatus(connectorName);
```

```
        if ("RUNNING".equals(status.getState()) && status.getTaskCount() >
0) {
    System.out.printf("Connector %s is ready%n", connectorName);
    return;
}

Thread.sleep(5000);

} catch (Exception e) {
    System.err.printf("Error checking connector %s status: %s%n",
                      connectorName, e.getMessage());
    Thread.sleep(5000);
}
}

throw new TimeoutException("Connector " + connectorName +
    " did not become ready within " + timeout);
}

private void monitorConnectorHealth() {
    try {
        String[] connectors = restClient.listConnectors();

        for (String connectorName : connectors) {
            checkConnectorHealth(connectorName);
        }
    } catch (Exception e) {
        System.err.println("Error during health monitoring: " +
e.getMessage());
    }
}

private void checkConnectorHealth(String connectorName) {
    try {
        ConnectorStatus status = restClient.getConnectorStatus(connectorName);
        ConnectorHealth health =
connectorHealth.computeIfAbsent(connectorName,
    k -> new ConnectorHealth());

        health.updateStatus(status.getState(), status.getTaskCount());

        // Check for issues
        if (!"RUNNING".equals(status.getState())) {
            health.recordFailure();
            handleUnhealthyConnector(connectorName, status);
        } else {
            health.recordSuccess();
        }

        // Log health summary
        if (health.getConsecutiveFailures() > 0) {
            System.err.printf("Connector %s health: %d consecutive

```

```
failures%n",
        connectorName, health.getConsecutiveFailures());
    }

} catch (Exception e) {
    System.err.printf("Error checking health for connector %s: %s%n",
        connectorName, e.getMessage());
}
}

private void handleUnhealthyConnector(String connectorName, ConnectorStatus
status) {
    ConnectorHealth health = connectorHealth.get(connectorName);

    // Auto-restart after 3 consecutive failures
    if (health.getConsecutiveFailures() >= 3) {
        try {
            System.out.printf("Auto-restarting unhealthy connector: %s%n",
connectorName);
            restClient.restartConnector(connectorName);
            health.recordRestart();

        } catch (Exception e) {
            System.err.printf("Failed to restart connector %s: %s%n",
                connectorName, e.getMessage());
        }
    }

    // Alert after 5 consecutive failures
    if (health.getConsecutiveFailures() >= 5) {
        sendAlert(String.format("Connector %s has failed %d times
consecutively",
            connectorName, health.getConsecutiveFailures()));
    }
}

private void collectMetrics() {
    try {
        String[] connectors = restClient.listConnectors();

        System.out.println("\n==== Connector Metrics Summary ====");
        System.out.printf("Total Connectors: %d%n", connectors.length);

        int runningCount = 0;
        int failedCount = 0;

        for (String connectorName : connectors) {
            try {
                ConnectorStatus status =
restClient.getConnectorStatus(connectorName);

                if ("RUNNING".equals(status.getState())) {
                    runningCount++;
                } else {

```

```
        failedCount++;
    }

    System.out.printf(" %s: %s (%d tasks)%n",
                      connectorName, status.getState(), status.getTaskCount());

} catch (Exception e) {
    failedCount++;
    System.err.printf(" %s: ERROR - %s%n", connectorName,
e.getMessage());
}
}

System.out.printf("Running: %d, Failed: %d%n", runningCount,
failedCount);

} catch (Exception e) {
    System.err.println("Error collecting metrics: " + e.getMessage());
}
}

private void backupConfigurations() {
try {
    String[] connectors = restClient.listConnectors();

    System.out.println("Backing up connector configurations...");

    for (String connectorName : connectors) {
        try {
            Map<String, String> config =
restClient.getConnectorConfig(connectorName);
            saveConfigurationBackup(connectorName, config);

        } catch (Exception e) {
            System.err.printf("Failed to backup config for %s: %s%n",
connectorName, e.getMessage());
        }
    }
}

} catch (Exception e) {
    System.err.println("Error during configuration backup: " +
e.getMessage());
}
}

private void saveConfigurationBackup(String connectorName, Map<String, String>
config) {
    // In production, save to file system, database, or configuration
management system
    System.out.printf("Backed up configuration for %s (%d properties)%n",
                      connectorName, config.size());
}

private void sendAlert(String message) {
```

```
// In production, integrate with alerting system (Slack, PagerDuty, etc.)
System.err.println("ALERT: " + message);
}

// Helper classes
public static class ConnectorDeployment {
    private final String name;
    private final Map<String, String> config;
    private final int order;
    private final boolean critical;

    public ConnectorDeployment(String name, Map<String, String> config,
                              int order, boolean critical) {
        this.name = name;
        this.config = config;
        this.order = order;
        this.critical = critical;
    }

    // Getters
    public String getName() { return name; }
    public Map<String, String> getConfig() { return config; }
    public int getOrder() { return order; }
    public boolean isCritical() { return critical; }
}

public static class ConnectorHealth {
    private String lastState = "UNKNOWN";
    private int taskCount = 0;
    private int consecutiveFailures = 0;
    private int consecutiveSuccesses = 0;
    private long lastUpdate = System.currentTimeMillis();
    private int restartCount = 0;

    public void updateStatus(String state, int tasks) {
        this.lastState = state;
        this.taskCount = tasks;
        this.lastUpdate = System.currentTimeMillis();
    }

    public void recordFailure() {
        consecutiveFailures++;
        consecutiveSuccesses = 0;
    }

    public void recordSuccess() {
        consecutiveSuccesses++;
        consecutiveFailures = 0;
    }

    public void recordRestart() {
        restartCount++;
    }
}
```

```

// Getters
public String getLastState() { return lastState; }
public int getTaskCount() { return taskCount; }
public int getConsecutiveFailures() { return consecutiveFailures; }
public int getConsecutiveSuccesses() { return consecutiveSuccesses; }
public long getLastUpdate() { return lastUpdate; }
public int getRestartCount() { return restartCount; }
}

// Usage example
public static void main(String[] args) {
    KafkaConnectManager manager = new
KafkaConnectManager("http://localhost:8083");

    try {
        manager.start();

        // Create deployment pipeline
        List<ConnectorDeployment> pipeline = Arrays.asList(
            new ConnectorDeployment("postgres-source",
                JDBCSourceConnectorExample.createIncrementalSourceConfig(), 1,
true),
            new ConnectorDeployment("elasticsearch-sink",
                ElasticsearchConnectorExample.createElasticsearchSinkConfig(),
2, false),
            new ConnectorDeployment("s3-backup",
                S3ConnectorExample.createS3SinkConfig(), 3, false)
        );

        // Deploy pipeline
        manager.deployConnectorPipeline(pipeline);

        // Keep running
        Thread.sleep(Long.MAX_VALUE);

    } catch (Exception e) {
        System.err.println("Manager failed: " + e.getMessage());
    } finally {
        manager.stop();
    }
}
}

```

## ⚖️ Comparisons & Trade-offs

### Deployment Mode Comparison

Feature	Standalone Mode	Distributed Mode
<b>Setup Complexity</b>	Simple	Complex

Feature	Standalone Mode	Distributed Mode
<b>Fault Tolerance</b>	None	High
<b>Scalability</b>	Limited	Horizontal
<b>Configuration</b>	File-based	REST API + Kafka topics
<b>Use Case</b>	Development, Testing	Production
<b>Resource Usage</b>	Low	Higher
<b>Management</b>	Manual	Automated

## Connector Types Comparison

Aspect	Source Connectors	Sink Connectors
<b>Data Flow</b>	External System → Kafka	Kafka → External System
<b>Offset Management</b>	Track external system position	Track Kafka offsets
<b>Error Handling</b>	Retry reading from source	DLQ, ignore, or fail
<b>Backpressure</b>	Kafka topic capacity	External system capacity
<b>Exactly-Once</b>	Depends on source capabilities	Idempotent writes needed

## Common Connectors Comparison

Connector	Type	Use Case	Complexity	Performance
<b>JDBC</b>	Source/Sink	Database integration	Medium	Good
<b>Debezium</b>	Source	Change Data Capture	High	Excellent
<b>Elasticsearch</b>	Sink	Search indexing	Low	Good
<b>S3</b>	Sink	Data lake/backup	Low	Excellent
<b>File</b>	Source/Sink	File system integration	Low	Good

## 💡 Common Pitfalls & Best Practices

### 1. Configuration Issues

#### ✗ Incorrect Offset Storage Configuration

```
// DON'T - Wrong offset storage for distributed mode
Map<String, String> badConfig = new HashMap<>();
badConfig.put("offset.storage.file.filename", "/tmp/connect.offsets"); // File
storage in distributed mode!
badConfig.put("group.id", "connect-cluster");
```

```
// DO - Proper distributed mode configuration
Map<String, String> goodConfig = new HashMap<>();
goodConfig.put("group.id", "connect-cluster");
goodConfig.put("offset.storage.topic", "connect-offsets");
goodConfig.put("offset.storage.replication.factor", "3");
goodConfig.put("offset.storage.partitions", "25");
```

## ✖ Missing Replication Factor

```
// DON'T - Default replication factor (might be 1)
config.put("config.storage.topic", "connect-configs");
```

```
// DO - Explicit replication factor for production
config.put("config.storage.topic", "connect-configs");
config.put("config.storage.replication.factor", "3");
```

## 2. Resource Management Issues

### ✖ No Resource Limits

```
// DON'T - Unbounded resource usage
Map<String, String> connectorConfig = new HashMap<>();
connectorConfig.put("tasks.max", "100"); // Too many tasks
connectorConfig.put("batch.size", "100000"); // Too large batches
```

```
// DO - Reasonable resource limits
Map<String, String> connectorConfig = new HashMap<>();
connectorConfig.put("tasks.max", "4"); // Based on partition count
connectorConfig.put("batch.size", "2000"); // Reasonable batch size
connectorConfig.put("poll.interval.ms", "5000"); // Don't poll too frequently
```

## 3. Error Handling Problems

### ✖ No Error Handling Strategy

```
// DON'T - Default error handling (might stop connector)
Map<String, String> config = new HashMap<>();
config.put("connector.class",
"io.confluent.connect.elasticsearch.ElasticsearchSinkConnector");
// No error handling configuration
```

```
// DO - Comprehensive error handling
Map<String, String> config = new HashMap<>();
config.put("connector.class",
"io.confluent.connect.elasticsearch.ElasticsearchSinkConnector");

// Error tolerance
config.put("errors.tolerance", "all");
config.put("errors.log.enable", "true");
config.put("errors.log.include.messages", "true");

// Dead letter queue
config.put("errors.deadletterqueue.topic.name", "connect-dlq");
config.put("errors.deadletterqueue.topic.replication.factor", "3");
config.put("errors.deadletterqueue.context.headers.enable", "true");

// Retry configuration
config.put("errors.retry.timeout", "300000"); // 5 minutes
config.put("errors.retry.delay.max.ms", "60000"); // 1 minute max delay
```

## 4. Security Oversights

### ✗ No Security Configuration

```
// DON'T - Plain text passwords and no SSL
config.put("connection.password", "plaintext-password");
config.put("connection.url", "jdbc:postgresql://localhost:5432/db");
```

```
// DO - Proper security configuration
// Use ConfigProvider for sensitive data
config.put("config.providers", "file");
config.put("config.providers.file.class",
"org.apache.kafka.common.config.provider.FileConfigProvider");

// Reference encrypted passwords
config.put("connection.password", "${file:/etc/kafka-
connect/secrets.properties:db.password}");

// SSL configuration
config.put("connection.url", "jdbc:postgresql://localhost:5432/db?
ssl=true&sslfactory=org.postgresql.ssl.NonValidatingFactory");

// Kafka security
config.put("bootstrap.servers", "localhost:9093");
config.put("security.protocol", "SSL");
config.put("ssl.keystore.location", "/etc/kafka-connect/keystore.jks");
```

```
config.put("ssl.keystore.password", "${file:/etc/kafka-connect/secrets.properties:keystore.password}");
```

## Best Practices Summary

### Configuration Best Practices

1. **Use distributed mode for production** - Better fault tolerance and scalability
2. **Set appropriate replication factors** - Ensure high availability for internal topics
3. **Configure error handling** - Use DLQ and error tolerance for resilience
4. **Secure sensitive data** - Use ConfigProvider for passwords and keys
5. **Monitor resource usage** - Set reasonable limits on tasks, batch sizes, and polling

### Operational Best Practices

1. **Monitor connector health** - Implement automated health checks and alerting
2. **Backup configurations** - Store connector configs in version control
3. **Use proper naming conventions** - Include environment and purpose in names
4. **Implement gradual rollouts** - Deploy connectors incrementally
5. **Plan for disaster recovery** - Document recovery procedures

### Development Best Practices

1. **Test with realistic data volumes** - Performance characteristics change with scale
2. **Implement custom SMTs carefully** - Test thoroughly and handle edge cases
3. **Use schema registry** - For better data governance and evolution
4. **Monitor offset lag** - Ensure connectors keep up with data flow
5. **Document dependencies** - Clearly specify connector deployment order

## Real-World Use Cases

### 1. E-commerce Data Pipeline

```
public class EcommerceDataPipeline {  
  
    public static List<ConnectorDeployment> createEcommercePipeline() {  
        List<ConnectorDeployment> pipeline = new ArrayList<>();  
  
        // 1. Source: MySQL database (orders, customers, products)  
        Map<String, String> mysqlSource = new HashMap<>();  
        mysqlSource.put("name", "mysql-ecommerce-source");  
        mysqlSource.put("connector.class",  
"io.debezium.connector.mysql.MySqlConnector");  
        mysqlSource.put("database.hostname", "mysql.internal");  
        mysqlSource.put("database.server.name", "ecommerce");  
        mysqlSource.put("table.include.list",  
"ecommerce.orders,ecommerce.customers,ecommerce.products");  
    }  
}
```

```
mysqlSource.put("transforms", "unwrap");
mysqlSource.put("transforms.unwrap.type",
"io.debezium.transforms.ExtractNewRecordState");

pipeline.add(new ConnectorDeployment("mysql-source", mysqlSource, 1,
true));

// 2. Sink: Elasticsearch for search and analytics
Map<String, String> esSink = new HashMap<>();
esSink.put("name", "elasticsearch-ecommerce-sink");
esSink.put("connector.class",
"io.confluent.connect.elasticsearch.ElasticsearchSinkConnector");
esSink.put("topics",
"ecommerce.orders,ecommerce.customers,ecommerce.products");
esSink.put("connection.url", "http://elasticsearch.internal:9200");
esSink.put("type.name", "_doc");
esSink.put("topic.index.map",
"ecommerce.orders:orders,ecommerce.customers:customers,ecommerce.products:products");
esSink.put("topic.index.map",
"ecommerce.orders:orders,ecommerce.customers:customers,ecommerce.products:products");

pipeline.add(new ConnectorDeployment("elasticsearch-sink", esSink, 2,
false));

// 3. Sink: S3 for data lake
Map<String, String> s3Sink = new HashMap<>();
s3Sink.put("name", "s3-datalake-sink");
s3Sink.put("connector.class", "io.confluent.connect.s3.S3SinkConnector");
s3Sink.put("topics",
"ecommerce.orders,ecommerce.customers,ecommerce.products");
s3Sink.put("s3.bucket.name", "ecommerce-datalake");
s3Sink.put("format.class",
"io.confluent.connect.s3.format.parquet.ParquetFormat");
s3Sink.put("partitioner.class",
"io.confluent.connect.storage.partition.TimeBasedPartitioner");
s3Sink.put("path.format", "'year'=YYYY/'month'=MM/'day'=dd'");
s3Sink.put("rotate.interval.ms", "3600000"); // Hourly rotation

pipeline.add(new ConnectorDeployment("s3-datalake-sink", s3Sink, 3,
false));

// 4. Sink: PostgreSQL data warehouse
Map<String, String> postgresWarehouse = new HashMap<>();
postgresWarehouse.put("name", "postgres-warehouse-sink");
postgresWarehouse.put("connector.class",
"io.confluent.connect.jdbc.JdbcSinkConnector");
postgresWarehouse.put("topics", "ecommerce.orders,ecommerce.customers");
postgresWarehouse.put("connection.url",
"jdbc:postgresql://warehouse.internal:5432/analytics");
postgresWarehouse.put("auto.create", "true");
postgresWarehouse.put("auto.evolve", "true");
postgresWarehouse.put("insert.mode", "upsert");
postgresWarehouse.put("pk.mode", "record_value");
postgresWarehouse.put("pk.fields", "id");
```

```
        pipeline.add(new ConnectorDeployment("postgres-warehouse-sink",
postgresWarehouse, 4, false));

        return pipeline;
    }
}
```

## 2. IoT Sensor Data Processing

```
public class IoTSensorPipeline {

    public static List<ConnectorDeployment> createIoTPipeline() {
        List<ConnectorDeployment> pipeline = new ArrayList<>();

        // 1. Source: File-based sensor data ingestion
        Map<String, String> fileSource = new HashMap<>();
        fileSource.put("name", "iot-file-source");
        fileSource.put("connector.class", "com.custom.IoTFileSourceConnector");
        fileSource.put("directory.path", "/data/sensors");
        fileSource.put("file.pattern", ".*\\".sensor$");
        fileSource.put("topic", "raw-sensor-data");
        fileSource.put("poll.interval.ms", "1000");

        // Transform sensor data
        fileSource.put("transforms", "parseJson,addTimestamp,routeByType");
        fileSource.put("transforms.parseJson.type", "com.custom.ParseIoTData");
        fileSource.put("transforms.addTimestamp.type",
"org.apache.kafka.connect.transforms.InsertField$Value");
        fileSource.put("transforms.addTimestamp.timestamp.field", "processed_at");
        fileSource.put("transforms.routeByType.type",
"org.apache.kafka.connect.transforms.RegexRouter");
        fileSource.put("transforms.routeByType.regex", "raw-sensor-data");
        fileSource.put("transforms.routeByType.replacement",
"sensors-${header:sensor_type}");

        pipeline.add(new ConnectorDeployment("iot-file-source", fileSource, 1,
true));

        // 2. Sink: InfluxDB for time series data
        Map<String, String> influxSink = new HashMap<>();
        influxSink.put("name", "influxdb-timeseries-sink");
        influxSink.put("connector.class",
"com.influxdb.kafka.InfluxDBSinkConnector");
        influxSink.put("topics.regex", "sensors-.*");
        influxSink.put("influxdb.url", "http://influxdb.internal:8086");
        influxSink.put("influxdb.db", "iot_sensors");
        influxSink.put("measurement.name.format", "${topic}");

        pipeline.add(new ConnectorDeployment("influxdb-sink", influxSink, 2,
false));
    }
}
```

```

    // 3. Sink: Alert system for anomalies
    Map<String, String> alertSink = new HashMap<>();
    alertSink.put("name", "alert-webhook-sink");
    alertSink.put("connector.class", "com.custom.WebhookSinkConnector");
    alertSink.put("topics", "sensor-alerts");
    alertSink.put("webhook.url", "http://alert-service.internal/webhook");
    alertSink.put("http.headers", "Content-
Type:application/json,Authorization:Bearer ${file:/secrets:webhook.token}");
    pipeline.add(new ConnectorDeployment("alert-webhook-sink", alertSink, 3,
true));

    return pipeline;
}
}

```

### 3. Financial Data Compliance Pipeline

```

public class FinancialCompliancePipeline {

    public static List<ConnectorDeployment> createCompliancePipeline() {
        List<ConnectorDeployment> pipeline = new ArrayList<>();

        // 1. Source: Core banking system
        Map<String, String> bankingSource = new HashMap<>();
        bankingSource.put("name", "core-banking-source");
        bankingSource.put("connector.class",
"io.confluent.connect.jdbc.JdbcSourceConnector");
        bankingSource.put("connection.url", "${file:/secrets:banking.url}");
        bankingSource.put("connection.user", "${file:/secrets:banking.user}");
        bankingSource.put("connection.password",
"${file:/secrets:banking.password}");
        bankingSource.put("table.whitelist", "transactions,accounts,customers");
        bankingSource.put("mode", "timestamp+incrementing");
        bankingSource.put("timestamp.column.name", "updated_at");
        bankingSource.put("incrementing.column.name", "id");
        bankingSource.put("poll.interval.ms", "30000"); // 30 seconds for near
real-time

        // Privacy transformations
        bankingSource.put("transforms",
"maskPII,encryptSensitive,addAuditFields");
        bankingSource.put("transforms.maskPII.type", "com.bank.MaskPersonalData");
        bankingSource.put("transforms.maskPII.fields",
"ssn,account_number,routing_number");
        bankingSource.put("transforms.encryptSensitive.type",
"com.bank.EncryptField");
        bankingSource.put("transforms.encryptSensitive.fields",
"customer_id,transaction_amount");
        bankingSource.put("transforms.addAuditFields.type",
"org.apache.kafka.connect.transforms.InsertField$Value");

```

```
    bankingSource.put("transforms.addAuditFields.static.field",
"audit_source");
    bankingSource.put("transforms.addAuditFields.static.value", "core-
banking");

    pipeline.add(new ConnectorDeployment("banking-source", bankingSource, 1,
true));

    // 2. Sink: Compliance data warehouse
    Map<String, String> complianceSink = new HashMap<>();
    complianceSink.put("name", "compliance-warehouse-sink");
    complianceSink.put("connector.class",
"io.confluent.connect.jdbc.JdbcSinkConnector");
    complianceSink.put("topics", "banking-transactions,banking-
accounts,banking-customers");
    complianceSink.put("connection.url", "${file:/secrets:warehouse.url}");
    complianceSink.put("connection.user", "${file:/secrets:warehouse.user}");
    complianceSink.put("connection.password",
"${file:/secrets:warehouse.password}");
    complianceSink.put("table.name.format", "compliance_${topic}");
    complianceSink.put("auto.create", "true");
    complianceSink.put("auto.evolve", "false"); // Strict schema control
    complianceSink.put("insert.mode", "insert"); // Append-only for audit
trail

    // Enhanced error handling for compliance
    complianceSink.put("errors.tolerance", "none"); // No error tolerance for
compliance
    complianceSink.put("errors.log.enable", "true");
    complianceSink.put("errors.log.include.messages", "true");

    pipeline.add(new ConnectorDeployment("compliance-warehouse-sink",
complianceSink, 2, true));

    // 3. Sink: Encrypted backup to S3
    Map<String, String> backupSink = new HashMap<>();
    backupSink.put("name", "encrypted-backup-sink");
    backupSink.put("connector.class",
"io.confluent.connect.s3.S3SinkConnector");
    backupSink.put("topics", "banking-transactions,banking-accounts,banking-
customers");
    backupSink.put("s3.bucket.name", "financial-compliance-backup");
    backupSink.put("format.class",
"io.confluent.connect.s3.format.json.JsonFormat");

    // Encryption and retention
    backupSink.put("s3.ssea.name", "AES256");
    backupSink.put("s3.canned.acl", "private");
    backupSink.put("rotate.interval.ms", "86400000"); // Daily rotation

    // Compliance partitioning
    backupSink.put("partitioner.class",
"io.confluent.connect.storage.partition.TimeBasedPartitioner");
    backupSink.put("path.format", "'year'=YYYY/'month'=MM/'day'=dd');
```

```

        backupSink.put("timezone", "UTC");

        pipeline.add(new ConnectorDeployment("encrypted-backup-sink", backupSink,
3, true));

        return pipeline;
    }
}

```

## Version Highlights

### Kafka Connect Evolution Timeline

Version	Release Date	Major Features
<b>4.0</b>	September 2025	Enhanced REST API, improved error handling
<b>3.0</b>	September 2021	Better connector management, KRaft support
<b>2.8</b>	April 2021	Incremental cooperative rebalancing
<b>2.4</b>	December 2019	Enhanced SMT support, better monitoring
<b>2.0</b>	July 2018	Header support, improved error handling
<b>1.0</b>	October 2017	Schema evolution support
<b>0.11</b>	June 2017	Exactly-once semantics support
<b>0.10</b>	May 2016	<b>Initial Kafka Connect release</b>

### Key Features by Version

#### Kafka Connect 4.0 (September 2025)

- ◆ **Enhanced REST API:** Better error responses and bulk operations
- ◆ **Improved Error Handling:** More granular error tolerance configuration
- ◆ **Better Monitoring:** Enhanced JMX metrics and health checks
- ◆ **Security Improvements:** Better integration with OAuth and RBAC

#### Kafka Connect 3.x Series

- 3.6** (October 2023): Enhanced connector plugin management
- 3.5** (June 2023): Improved offset management and recovery
- 3.4** (February 2023): Better resource management and scaling
- 3.3** (October 2022): Enhanced transform pipeline performance
- 3.2** (May 2022): Improved configuration validation
- 3.1** (January 2022): Better dead letter queue handling
- 3.0** (September 2021): **Major stability improvements**

#### Kafka Connect 2.x Highlights

- **2.8** (April 2021): Incremental cooperative rebalancing for Connect
- **2.7** (December 2020): Enhanced connector restart capabilities
- **2.6** (August 2020): Better offset commit strategies
- **2.5** (April 2020): Improved exactly-once support
- **2.4** (December 2019): **Enhanced SMT framework**

## Notable Connector Ecosystem Growth

- **Debezium** (2016-present): Leading CDC connector platform
- **Confluent Hub** (2018-present): Connector distribution platform
- **Community Connectors** (2017-present): 100+ community-developed connectors

## Current Recommendations (2025)

```
// Modern Kafka Connect configuration (4.0+)
public static Properties modernConnectConfig() {
    Properties props = new Properties();

    // Core configuration
    props.put("bootstrap.servers", "localhost:9092");
    props.put("group.id", "connect-cluster");
    props.put("client.id", "connect-worker-1");

    // Enhanced security (4.0+)
    props.put("security.protocol", "SASL_SSL");
    props.put("sasl.mechanism", "OAUTHBEARER");
    props.put("sasl.login.callback.handler.class",
        "org.apache.kafka.common.security.oauthbearer.secured.OAuthBearerLoginCallbackHandler");

    // Improved error handling (4.0+)
    props.put("errors.tolerance", "all");
    props.put("errors.log.enable", "true");
    props.put("errors.deadletterqueue.topic.name", "connect-dlq");
    props.put("errors.deadletterqueue.context.headers.enable", "true");

    // Enhanced monitoring (4.0+)
    props.put("metrics.reporters", "org.apache.kafka.common.metrics.JmxReporter");
    props.put("metrics.jmx.prefix", "kafka.connect");

    // Internal topics with high availability
    props.put("config.storage.topic", "connect-configs");
    props.put("config.storage.replication.factor", 3);

    props.put("offset.storage.topic", "connect-offsets");
    props.put("offset.storage.replication.factor", 3);
    props.put("offset.storage.partitions", 25);

    props.put("status.storage.topic", "connect-status");
    props.put("status.storage.replication.factor", 3);
```

```
    return props;  
}
```

---

## 🔗 Additional Resources

### 📘 Official Documentation

- [Kafka Connect Documentation](#)
- [Kafka Connect REST API](#)
- [Connector Development Guide](#)

### 🎓 Learning Resources

- [Confluent Connect Tutorial](#)
- [Building Kafka Connect Connectors](#)
- [Kafka Connect Examples](#)

### 🔧 Connector Ecosystem

- [Confluent Hub](#) - Official connector marketplace
- [Debezium](#) - CDC connectors
- [Kafka Connect File Pulse](#)

### 📊 Monitoring & Operations

- [Kafka Connect Monitoring](#)
- [Connect Metrics](#)
- [Operational Best Practices](#)

### 🛠 Troubleshooting

- [Common Connect Issues](#)
- [Connect FAQ](#)
- [Debugging Connect](#)

---

**Last Updated:** September 2025

**Kafka Version:** 4.0.0

**Connect Framework:** Production-ready since 2.0+

**💡 Pro Tip:** Start with well-tested connectors from Confluent Hub, use distributed mode for production, implement comprehensive error handling with DLQ, and monitor connector health continuously. Custom connectors should follow the single responsibility principle and be thoroughly tested.