

Spring Kafka Error Handling & Retry: Part 2 - Dead Letter Topics & Legacy Handlers

Continuation of the comprehensive guide covering Dead Letter Topics configuration, retry patterns, and legacy SeekToCurrentErrorHandler migration strategies.

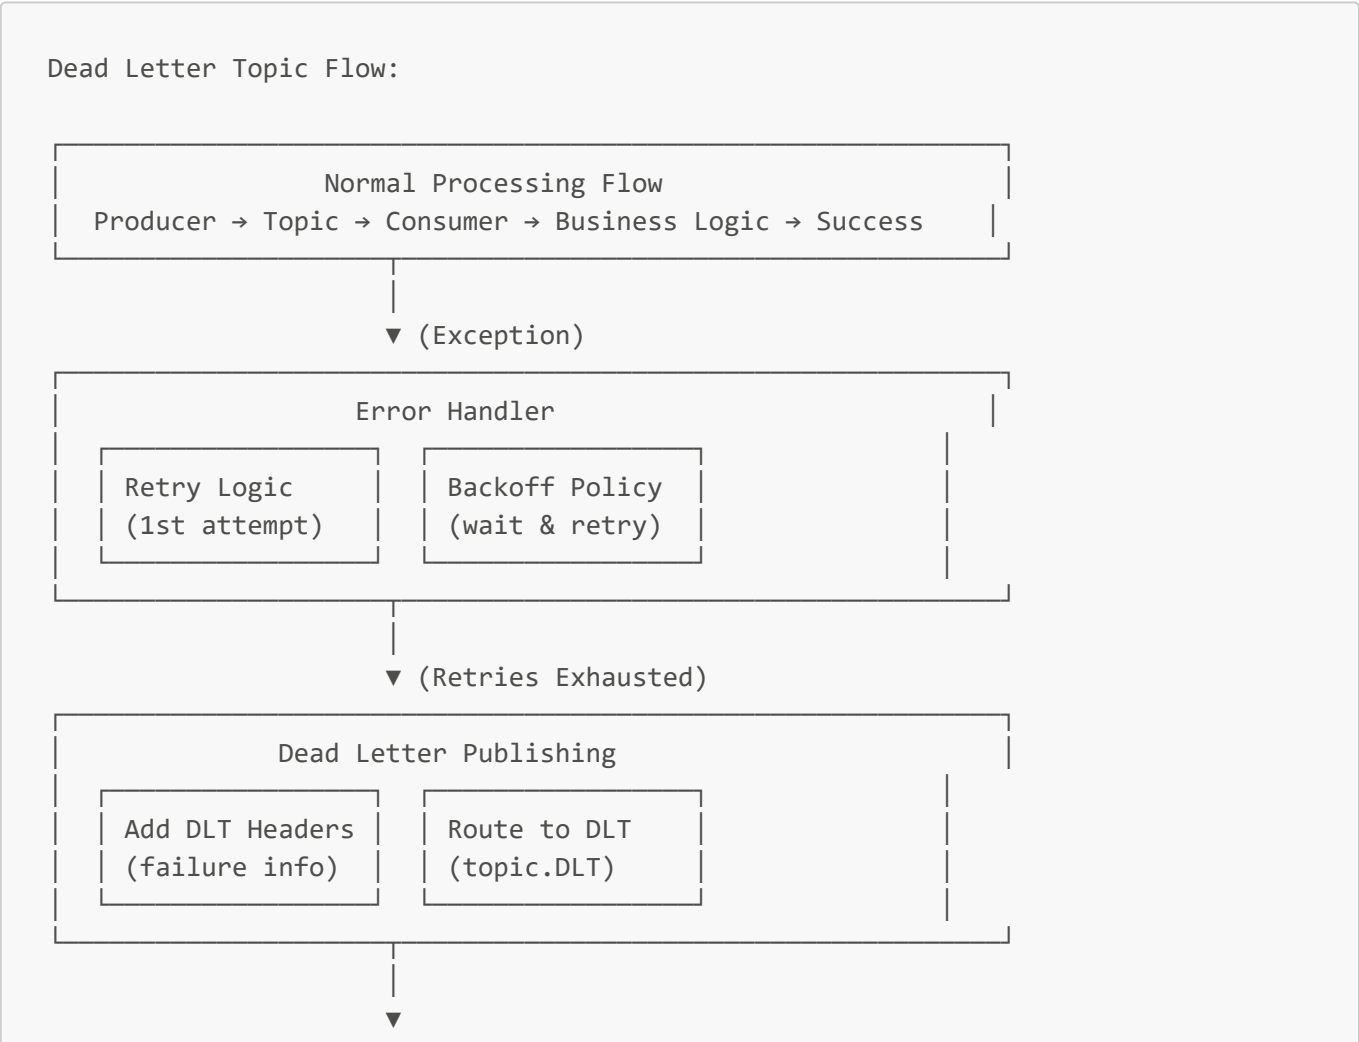
🦴 Dead Letter Topics (DLT)

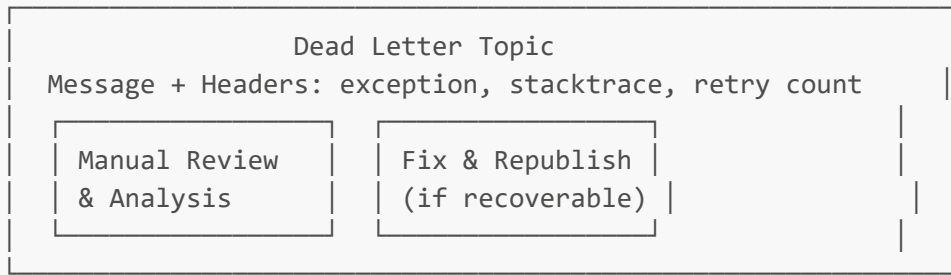
Simple Explanation: Dead Letter Topics (DLT) are special Kafka topics where failed messages are sent after all retry attempts are exhausted. They act as a "safety net" to prevent message loss and provide a mechanism for manual inspection and recovery of problematic messages.

Why Dead Letter Topics are Essential:

- **Message Preservation:** Ensure no messages are lost even after processing failures
- **Problem Isolation:** Separate failed messages from healthy processing flow
- **Manual Recovery:** Allow operators to inspect, fix, and reprocess failed messages
- **Debugging:** Provide visibility into failure patterns and root causes
- **System Reliability:** Prevent poison pills from blocking entire processing pipelines

Dead Letter Topic Architecture:





Configuring DLT Publishing

Advanced Dead Letter Topic Configuration

```

import org.springframework.context.annotation.Bean;
import org.springframework.context.annotation.Configuration;
import org.springframework.kafka.listener.DeadLetterPublishingRecoverer;
import org.springframework.kafka.listener.DefaultErrorHandler;
import org.springframework.kafka.core.KafkaTemplate;
import org.springframework.kafka.support.KafkaHeaders;

import org.apache.kafka.clients.consumer.ConsumerRecord;
import org.apache.kafka.clients.producer.ProducerRecord;
import org.apache.kafka.common.TopicPartition;

import java.util.function.BiFunction;

/**
 * Advanced Dead Letter Topic configuration and management
 */
@Configuration
@lombok.extern.slf4j.Slf4j
public class DeadLetterTopicConfiguration {

    @Autowired
    private KafkaTemplate<String, Object> kafkaTemplate;

    @Autowired
    private MeterRegistry meterRegistry;

    /**
     * Basic Dead Letter Topic configuration
     */
    @Bean("basicDltErrorHandler")
    public CommonErrorHandler basicDltErrorHandler() {

        // Simple DLT recoverer - sends to topic.DLT with same partition
        DeadLetterPublishingRecoverer recoverer = new
        DeadLetterPublishingRecoverer(kafkaTemplate);

        // Fixed backoff: 3 retries with 2 second intervals
        FixedBackOff backOff = new FixedBackOff(2000L, 3L);
  
```

```

        DefaultErrorHandler errorHandler = new DefaultErrorHandler(recoverer,
backOff);

        // Configure retryable/non-retryable exceptions
        configureBasicExceptionClassification(errorHandler);

        log.info("Configured basic DLT error handler: 3 retries, 2s interval,
default DLT routing");

        return errorHandler;
    }

    /**
     * Advanced DLT configuration with custom routing and headers
     */
    @Bean("advancedDltErrorHandler")
    public CommonErrorHandler advancedDltErrorHandler() {

        // Custom destination resolver for intelligent DLT routing
        BiFunction<ConsumerRecord<?, ?>, Exception, TopicPartition>
destinationResolver =
            this::resolveDltDestination;

        DeadLetterPublishingRecoverer recoverer = new
DeadLetterPublishingRecoverer(
            kafkaTemplate, destinationResolver);

        // Enhanced header function to add diagnostic information
        recoverer.setHeadersFunction(this::enhanceDltHeaders);

        // Exponential backoff with max retries
        ExponentialBackOffWithMaxRetries backOff = new
ExponentialBackOffWithMaxRetries(5);
        backOff.setInitialInterval(1000L);
        backOff.setMultiplier(2.0);
        backOff.setMaxInterval(30000L);

        DefaultErrorHandler errorHandler = new DefaultErrorHandler(recoverer,
backOff);

        // Advanced exception classification
        configureAdvancedExceptionClassification(errorHandler);

        // Add retry listeners for monitoring
        errorHandler.setRetryListeners(createDltRetryListener());

        log.info("Configured advanced DLT error handler with custom routing and
enhanced headers");

        return errorHandler;
    }

    /**

```

```

    * Multi-tier DLT configuration for different failure types
    */
    @Bean("multiTierDltErrorHandler")
    public CommonErrorHandler multiTierDltErrorHandler() {

        // Multi-tier destination resolver
        BiFunction<ConsumerRecord<?, ?>, Exception, TopicPartition>
multiTierResolver =
            (record, exception) -> {
                String originalTopic = record.topic();
                int originalPartition = record.partition();

                // Determine DLT tier based on exception severity and type
                if (isPoisonPill(exception)) {
                    return new TopicPartition(originalTopic + ".poison.DLT",
originalPartition);

                } else if (isCriticalFailure(exception)) {
                    return new TopicPartition(originalTopic + ".critical.DLT",
originalPartition);

                } else if (isBusinessLogicFailure(exception)) {
                    return new TopicPartition(originalTopic + ".business.DLT",
originalPartition);

                } else if (isTransientFailure(exception)) {
                    return new TopicPartition(originalTopic + ".transient.DLT",
originalPartition);

                } else {
                    return new TopicPartition(originalTopic + ".unknown.DLT",
originalPartition);
                }
            };

        DeadLetterPublishingRecoverer recoverer = new
DeadLetterPublishingRecoverer(
            kafkaTemplate, multiTierResolver);

        // Add tier-specific headers
        recoverer.setHeadersFunction((record, exception) -> {
            Map<String, Object> headers = enhanceDltHeaders(record, exception);
            headers.put("dlt-tier", determineDltTier(exception));
            headers.put("severity", determineSeverity(exception));
            headers.put("recovery-strategy",
determineRecoveryStrategy(exception));
            return headers;
        });

        // Different backoff based on exception type
        DefaultErrorHandler errorHandler = new DefaultErrorHandler(recoverer);

        errorHandler.setBackOffFunction((record, ex) -> {
            if (isPoisonPill(ex)) {

```

```

        return new FixedBackOff(0L, 0L); // No retry for poison pills
    } else if (isCriticalFailure(ex)) {
        return new FixedBackOff(5000L, 2L); // Quick retry for critical
    } else if (isTransientFailure(ex)) {
        return new ExponentialBackOff(1000L, 2.0); // Exponential for
transient
    } else {
        return new FixedBackOff(2000L, 3L); // Default
    }
});

configureAdvancedExceptionClassification(errorHandler);

log.info("Configured multi-tier DLT error handler with intelligent
routing");

return errorHandler;
}

/**
 * Conditional DLT configuration - some messages go to retry, others to DLT
 */
@Bean("conditionalDltErrorHandler")
public CommonErrorHandler conditionalDltErrorHandler() {

    // Conditional recoverer
    ConsumerRecordRecoverer conditionalRecoverer = (record, exception) -> {

        // Check if message should be retried instead of going to DLT
        if (shouldRetryInsteadOfDlt(record, exception)) {
            sendToRetryTopic(record, exception);

        } else if (shouldQuarantine(record, exception)) {
            sendToQuarantineTopic(record, exception);

        } else {
            // Use standard DLT
            DeadLetterPublishingRecoverer standardRecoverer =
                new DeadLetterPublishingRecoverer(kafkaTemplate);
            standardRecoverer.accept(record, exception);
        }
    };

    DefaultErrorHandler errorHandler = new
DefaultErrorHandler(conditionalRecoverer);

    // Very limited retries since we have conditional logic
    errorHandler.setBackOff(new FixedBackOff(1000L, 2L));

    configureAdvancedExceptionClassification(errorHandler);

    log.info("Configured conditional DLT error handler with retry/quarantine
logic");

```

```

        return errorHandler;
    }

    /**
     * Batched DLT configuration for high-volume scenarios
     */
    @Bean("batchedDltErrorHandler")
    public CommonErrorHandler batchedDltErrorHandler() {

        // Custom recoverer that batches DLT messages
        ConsumerRecordRecoverer batchedRecoverer = new
        BatchedDltRecoverer(kafkaTemplate);

        DefaultErrorHandler errorHandler = new
        DefaultErrorHandler(batchedRecoverer);

        // Shorter retries since we're batching
        errorHandler.setBackOff(new FixedBackOff(500L, 2L));

        configureAdvancedExceptionClassification(errorHandler);

        log.info("Configured batched DLT error handler for high-volume
        processing");

        return errorHandler;
    }

    // Helper methods for DLT configuration
    private TopicPartition resolveDltDestination(ConsumerRecord<?, ?> record,
    Exception exception) {
        String originalTopic = record.topic();
        int originalPartition = record.partition();

        // Custom routing logic based on message content and exception
        if (exception instanceof ValidationException) {
            return new TopicPartition(originalTopic + ".validation.DLT",
            originalPartition);

        } else if (exception instanceof ExternalServiceException) {
            return new TopicPartition(originalTopic + ".external.DLT",
            originalPartition);

        } else if (exception instanceof DatabaseException) {
            return new TopicPartition(originalTopic + ".database.DLT",
            originalPartition);

        } else if (exception.getCause() instanceof
        org.springframework.kafka.support.serializer.DeserializationException) {
            return new TopicPartition(originalTopic + ".poison.DLT",
            originalPartition);

        } else {
            // Check message content for additional routing
            Object value = record.value();

```

```

        if (value instanceof OrderEvent) {
            return new TopicPartition(originalTopic + ".orders.DLT",
originalPartition);
        } else if (value instanceof PaymentEvent) {
            return new TopicPartition(originalTopic + ".payments.DLT",
originalPartition);
        } else {
            return new TopicPartition(originalTopic + ".DLT",
originalPartition);
        }
    }
}

private Map<String, Object> enhanceDltHeaders(ConsumerRecord<?, ?> record,
Exception exception) {
    Map<String, Object> headers = new HashMap<>();

    // Standard DLT headers
    headers.put(KafkaHeaders.DLT_ORIGINAL_TOPIC, record.topic());
    headers.put(KafkaHeaders.DLT_ORIGINAL_PARTITION, record.partition());
    headers.put(KafkaHeaders.DLT_ORIGINAL_OFFSET, record.offset());
    headers.put(KafkaHeaders.DLT_ORIGINAL_TIMESTAMP, record.timestamp());
    headers.put(KafkaHeaders.DLT_EXCEPTION_FQCN,
exception.getClass().getName());
    headers.put(KafkaHeaders.DLT_EXCEPTION_MESSAGE, exception.getMessage());

    // Enhanced diagnostic headers
    headers.put("dlt-timestamp", System.currentTimeMillis());
    headers.put("dlt-hostname", getHostname());
    headers.put("dlt-application", getApplicationName());
    headers.put("dlt-version", getApplicationVersion());
    headers.put("retry-count", getRetryCount(record));
    headers.put("processing-duration", calculateProcessingDuration(record));

    // Business context headers
    Object value = record.value();
    if (value instanceof OrderEvent order) {
        headers.put("business-order-id", order.getOrderid());
        headers.put("business-customer-id", order.getCustomerId());
        headers.put("business-amount", order.getAmount().toString());
    } else if (value instanceof PaymentEvent payment) {
        headers.put("business-payment-id", payment.getPaymentId());
        headers.put("business-payment-method", payment.getMethod());
        headers.put("business-amount", payment.getAmount().toString());
    }

    // Exception chain analysis
    Throwable cause = exception.getCause();
    if (cause != null) {
        headers.put(KafkaHeaders.DLT_EXCEPTION_CAUSE_FQCN,
cause.getClass().getName());
        headers.put("dlt-root-cause",
getRootCause(exception).getClass().getName());
    }
}

```

```

    }

    // Stack trace (compressed for large stacks)
    String stackTrace = getStackTraceString(exception);
    if (stackTrace.length() > 4000) { // Kafka header limit consideration
        stackTrace = stackTrace.substring(0, 4000) + "...[truncated]";
    }
    headers.put(KafkaHeaders.DLT_EXCEPTION_STACKTRACE, stackTrace);

    return headers;
}

private RetryListener createDltRetryListener() {
    return new RetryListener() {
        @Override
        public void failedDelivery(ConsumerRecord<?, ?> record, Exception ex,
int deliveryAttempt) {
            log.warn("DLT retry attempt {} failed: topic={}, partition={},
offset={}, error={}",
                deliveryAttempt, record.topic(), record.partition(),
record.offset(), ex.getMessage());

            // Update retry metrics by exception type
            meterRegistry.counter("kafka.dlt.retry.attempts",
                Tags.of(
                    "topic", record.topic(),
                    "exception", ex.getClass().getSimpleName(),
                    "attempt", String.valueOf(deliveryAttempt)
                )).increment();
        }

        @Override
        public void recovered(ConsumerRecord<?, ?> record, Exception ex) {
            log.info("Record recovered before DLT: topic={}, partition={},
offset={}",
                record.topic(), record.partition(), record.offset());

            meterRegistry.counter("kafka.dlt.recovered",
                Tags.of("topic", record.topic(), "exception",
ex.getClass().getSimpleName()))
                .increment();
        }

        @Override
        public void recoveryFailed(ConsumerRecord<?, ?> record, Exception
original, Exception failure) {
            log.error("Record sent to DLT: topic={}, partition={}, offset={},
original={}, recovery={}",
                record.topic(), record.partition(), record.offset(),
                original.getMessage(), failure.getMessage());

            meterRegistry.counter("kafka.dlt.published",
                Tags.of(
                    "topic", record.topic(),

```



```

        "exception", original.getClass().getSimpleName(),
        "dlt-destination", resolveDltDestination(record,
original).topic()
    )).increment();
    }
};
}

// Exception classification methods
private boolean isPoisonPill(Exception exception) {
    return exception instanceof
org.springframework.kafka.support.serializer.DeserializationException ||
        exception instanceof
org.springframework.messaging.converter.MessageConversionException ||
        exception instanceof ClassCastException ||
        exception instanceof IllegalArgumentException;
}

private boolean isCriticalFailure(Exception exception) {
    return exception instanceof DatabaseException ||
        exception instanceof SecurityException ||
        (exception instanceof RuntimeException &&
            exception.getMessage() != null &&
            exception.getMessage().contains("CRITICAL"));
}

private boolean isBusinessLogicFailure(Exception exception) {
    return exception instanceof ValidationException ||
        exception instanceof BusinessException ||
        exception instanceof WorkflowException;
}

private boolean isTransientFailure(Exception exception) {
    return exception instanceof ExternalServiceException ||
        exception instanceof java.util.concurrent.TimeoutException ||
        exception instanceof
org.springframework.dao.TransientDataAccessException ||
        exception instanceof java.net.ConnectException;
}

private String determineDltTier(Exception exception) {
    if (isPoisonPill(exception)) return "POISON";
    if (isCriticalFailure(exception)) return "CRITICAL";
    if (isBusinessLogicFailure(exception)) return "BUSINESS";
    if (isTransientFailure(exception)) return "TRANSIENT";
    return "UNKNOWN";
}

private String determineSeverity(Exception exception) {
    if (isCriticalFailure(exception)) return "HIGH";
    if (isBusinessLogicFailure(exception)) return "MEDIUM";
    if (isTransientFailure(exception)) return "LOW";
    return "MEDIUM";
}

```

```

private String determineRecoveryStrategy(Exception exception) {
    if (isPoisonPill(exception)) return "MANUAL_INSPECTION";
    if (isCriticalFailure(exception)) return "IMMEDIATE_ATTENTION";
    if (isBusinessLogicFailure(exception)) return "BUSINESS_REVIEW";
    if (isTransientFailure(exception)) return "RETRY_LATER";
    return "INVESTIGATE";
}

private boolean shouldRetryInsteadOfDlt(ConsumerRecord<?, ?> record, Exception
exception) {
    // Only retry transient failures and only for certain message types
    if (!isTransientFailure(exception)) {
        return false;
    }

    // Check retry count
    int retryCount = getRetryCount(record);
    if (retryCount >= 3) {
        return false; // Already retried enough
    }

    // Check message age
    long messageAge = System.currentTimeMillis() - record.timestamp();
    if (messageAge > Duration.ofHours(1).toMillis()) {
        return false; // Too old to retry
    }

    return true;
}

private boolean shouldQuarantine(ConsumerRecord<?, ?> record, Exception
exception) {
    return isPoisonPill(exception) ||
        (exception instanceof SecurityException) ||
        (exception.getMessage() != null &&
            exception.getMessage().contains("QUARANTINE"));
}

private void sendToRetryTopic(ConsumerRecord<?, ?> record, Exception
exception) {
    String retryTopic = record.topic() + ".retry";

    ProducerRecord<Object, Object> retryRecord = new ProducerRecord<>(
        retryTopic, record.partition(), record.key(), record.value());

    // Add retry headers
    retryRecord.headers().add("original-topic", record.topic().getBytes());
    retryRecord.headers().add("retry-count",
String.valueOf(getRetryCount(record) + 1).getBytes());
    retryRecord.headers().add("retry-reason",
exception.getMessage().getBytes());
    retryRecord.headers().add("retry-timestamp",
String.valueOf(System.currentTimeMillis()).getBytes());

```

```

        retryRecord.headers().add("scheduled-retry-time",
            String.valueOf(System.currentTimeMillis() +
                Duration.ofMinutes(5).toMillis()).getBytes());

        kafkaTemplate.send(retryRecord);

        log.info("Sent to retry topic: original={}, retry={}, attempt={}",
            record.topic(), retryTopic, getRetryCount(record) + 1);
    }

    private void sendToQuarantineTopic(ConsumerRecord<?, ?> record, Exception
exception) {
        String quarantineTopic = record.topic() + ".quarantine";

        ProducerRecord<Object, Object> quarantineRecord = new ProducerRecord<>(
            quarantineTopic, record.partition(), record.key(), record.value());

        // Add quarantine headers
        quarantineRecord.headers().add("quarantine-reason",
            determinedDltTier(exception).getBytes());
        quarantineRecord.headers().add("quarantine-timestamp",
            String.valueOf(System.currentTimeMillis()).getBytes());
        quarantineRecord.headers().add("requires-manual-review",
            "true".getBytes());

        kafkaTemplate.send(quarantineRecord);

        log.warn("Sent to quarantine topic: original={}, quarantine={}, reason=
{}",
            record.topic(), quarantineTopic, determinedDltTier(exception));
    }

    // Utility methods
    private void configureBasicExceptionClassification(DefaultErrorHandler
errorHandler) {
        errorHandler.addNotRetryableExceptions(
            IllegalArgumentException.class,
            ClassCastException.class,

org.springframework.kafka.support.serializer.DeserializationException.class
        );

        errorHandler.addRetryableExceptions(
            ExternalServiceException.class,
            java.util.concurrent.TimeoutException.class
        );
    }

    private void configureAdvancedExceptionClassification(DefaultErrorHandler
errorHandler) {
        // Fatal exceptions - immediate DLT
        errorHandler.addNotRetryableExceptions(
            IllegalArgumentException.class,
            NullPointerException.class,

```

```

        ClassCastException.class,

org.springframework.kafka.support.serializer.DeserializationException.class,

org.springframework.messaging.converter.MessageConversionException.class,
        SecurityException.class
    );

    // Retryable exceptions
    errorHandler.addRetryableExceptions(
        ValidationException.class,
        ExternalServiceException.class,
        DatabaseException.class,
        java.util.concurrent.TimeoutException.class,
        org.springframework.dao.TransientDataAccessException.class,
        java.net.ConnectException.class
    );

    // Reset retry state on exception type change
    errorHandler.setResetStateOnExceptionChange(true);
}

private int getRetryCount(ConsumerRecord<?, ?> record) {
    Header retryHeader = record.headers().lastHeader("retry-count");
    if (retryHeader != null) {
        return Integer.parseInt(new String(retryHeader.value()));
    }
    return 0;
}

private long calculateProcessingDuration(ConsumerRecord<?, ?> record) {
    Header startHeader = record.headers().lastHeader("processing-start");
    if (startHeader != null) {
        long startTime = Long.parseLong(new String(startHeader.value()));
        return System.currentTimeMillis() - startTime;
    }
    return 0;
}

private String getHostname() {
    try {
        return InetAddress.getLocalHost().getHostName();
    } catch (Exception e) {
        return "unknown";
    }
}

private String getApplicationName() {
    return "kafka-consumer-app"; // Could read from properties
}

private String getApplicationVersion() {
    return "1.0.0"; // Could read from manifest
}

```

```

    private Throwable getRootCause(Throwable throwable) {
        Throwable cause = throwable;
        while (cause.getCause() != null) {
            cause = cause.getCause();
        }
        return cause;
    }

    private String getStackTraceString(Exception exception) {
        StringWriter sw = new StringWriter();
        PrintWriter pw = new PrintWriter(sw);
        exception.printStackTrace(pw);
        return sw.toString();
    }
}

/**
 * Batched DLT recoverer for high-volume scenarios
 */
@Component
@lombok.extern.slf4j.Slf4j
public class BatchedDltRecoverer implements ConsumerRecordRecoverer {

    private final KafkaTemplate<String, Object> kafkaTemplate;
    private final Map<String, List<ProducerRecord<Object, Object>>> batchedRecords
= new ConcurrentHashMap<>();
    private final ScheduledExecutorService scheduler =
Executors.newScheduledThreadPool(2);

    public BatchedDltRecoverer(KafkaTemplate<String, Object> kafkaTemplate) {
        this.kafkaTemplate = kafkaTemplate;

        // Schedule batch flushing every 5 seconds
        scheduler.scheduleAtFixedRate(this::flushBatches, 5, 5, TimeUnit.SECONDS);

        // Schedule batch flushing when batch size reaches threshold
        scheduler.scheduleAtFixedRate(this::flushLargeBatches, 1, 1,
TimeUnit.SECONDS);
    }

    @Override
    public void accept(ConsumerRecord<?, ?> record, Exception exception) {
        String dltTopic = record.topic() + ".DLT";

        ProducerRecord<Object, Object> dltRecord = new ProducerRecord<>(
            dltTopic, record.partition(), record.key(), record.value());

        // Add DLT headers
        dltRecord.headers().add("dlt-batch-timestamp",
String.valueOf(System.currentTimeMillis()).getBytes());
        dltRecord.headers().add("dlt-exception",
exception.getClass().getSimpleName().getBytes());
    }
}

```

```

        // Add to batch
        batchedRecords.computeIfAbsent(dltTopic, k -> new ArrayList<>
()).add(dltRecord);

        log.debug("Added record to DLT batch: topic={}, batchSize={}",
            dltTopic, batchedRecords.get(dltTopic).size());
    }

    private void flushBatches() {
        batchedRecords.forEach((topic, records) -> {
            if (!records.isEmpty()) {
                log.info("Flushing DLT batch: topic={}, size={}", topic,
records.size());

                records.forEach(kafkaTemplate::send);
                records.clear();
            }
        });
    }

    private void flushLargeBatches() {
        batchedRecords.forEach((topic, records) -> {
            if (records.size() >= 100) { // Threshold for large batches
                log.info("Flushing large DLT batch: topic={}, size={}", topic,
records.size());

                records.forEach(kafkaTemplate::send);
                records.clear();
            }
        });
    }

    @PreDestroy
    public void shutdown() {
        log.info("Shutting down batched DLT recoverer");
        flushBatches(); // Final flush
        scheduler.shutdown();
    }
}

// Custom exception classes for demonstration
class BusinessException extends Exception {
    public BusinessException(String message) { super(message); }
}

class WorkflowException extends Exception {
    public WorkflowException(String message) { super(message); }
}

```

Retrying from DLT

DLT Processing and Retry Patterns

```

import org.springframework.kafka.annotation.KafkaListener;
import org.springframework.kafka.annotation.DltHandler;
import org.springframework.kafka.support.KafkaHeaders;
import org.springframework.messaging.handler.annotation.Header;
import org.springframework.messaging.handler.annotation.Payload;

/**
 * Dead Letter Topic processing and retry patterns
 */
@Component
@lombok.extern.slf4j.Slf4j
public class DeadLetterTopicProcessor {

    @Autowired
    private KafkaTemplate<String, Object> kafkaTemplate;

    @Autowired
    private DltAnalysisService dltAnalysisService;

    @Autowired
    private MessageRepairService messageRepairService;

    /**
     * Manual DLT processing - inspect and potentially reprocess messages
     */
    @KafkaListener(
        topics = "orders.DLT",
        groupId = "dlt-manual-processor"
    )
    public void processOrdersDlt(
        @Payload OrderEvent order,
        @Header(KafkaHeaders.RECEIVED_TOPIC) String dltTopic,
        @Header(name = KafkaHeaders.DLT_ORIGINAL_TOPIC, required = false)
String originalTopic,
        @Header(name = KafkaHeaders.DLT_EXCEPTION_FQCN, required = false)
String exceptionClass,
        @Header(name = KafkaHeaders.DLT_EXCEPTION_MESSAGE, required = false)
String exceptionMessage,
        @Header(name = "retry-count", required = false) String
retryCountHeader,
        ConsumerRecord<String, OrderEvent> record) {

        log.info("Processing DLT message: orderId={}, originalTopic={}, exception=
{}",
            order.getOrderId(), originalTopic, exceptionClass);

        try {
            // Analyze the failure
            DltAnalysisResult analysis = dltAnalysisService.analyzeDltMessage(
                record, exceptionClass, exceptionMessage);

            if (analysis.isRecoverable()) {

```

```

        log.info("DLT message is recoverable: orderId={}, strategy={}",
            order.getOrderId(), analysis.getRecoveryStrategy());

        // Attempt to fix and reprocess
        handleRecoverableDltMessage(order, analysis, originalTopic);

    } else {
        log.warn("DLT message is not recoverable: orderId={}, reason={}",
            order.getOrderId(), analysis.getFailureReason());

        // Log for manual intervention
        handleNonRecoverableDltMessage(order, analysis);
    }

} catch (Exception e) {
    log.error("Failed to process DLT message: orderId={}",
        order.getOrderId(), e);

    // Send to manual review queue
    sendToManualReview(record, e);
}

}

/**
 * Automated DLT retry processor
 */
@KafkaListener(
    topics = {"orders.transient.DLT", "payments.transient.DLT"},
    groupId = "dlt-auto-retry-processor"
)
public void processTransientDlt(
    @Payload Object message,
    @Header(KafkaHeaders.RECEIVED_TOPIC) String dltTopic,
    @Header(name = KafkaHeaders.DLT_ORIGINAL_TOPIC) String originalTopic,
    @Header(name = "dlt-timestamp") long dltTimestamp,
    ConsumerRecord<String, Object> record) {

    // Only retry if message is not too old
    long messageAge = System.currentTimeMillis() - dltTimestamp;
    if (messageAge > Duration.ofHours(24).toMillis()) {
        log.info("DLT message too old, skipping retry: age={}h, topic={}",
            messageAge / 3600000, dltTopic);

        sendToExpiredDlt(record);
        return;
    }

    log.info("Attempting automated retry for transient DLT: originalTopic={},
age={}min",
        originalTopic, messageAge / 60000);

    try {
        // Check if the transient issue might be resolved
        if (isSystemHealthy() && canRetryMessage(message)) {

```



```

        // Add retry tracking headers
        ProducerRecord<Object, Object> retryRecord = new ProducerRecord<>(
            originalTopic, record.key(), message);

        retryRecord.headers().add("dlt-retry", "true".getBytes());
        retryRecord.headers().add("dlt-retry-timestamp",
            String.valueOf(System.currentTimeMillis()).getBytes());
        retryRecord.headers().add("original-dlt-topic",
            dltTopic.getBytes());

        kafkaTemplate.send(retryRecord);

        log.info("Successfully sent DLT message back to original topic: {}
-> {}",
            dltTopic, originalTopic);

    } else {
        log.info("System not ready for DLT retry, will try later: topic=
{}", dltTopic);

        // Send back to DLT with delay
        sendToDltWithDelay(record, Duration.ofMinutes(30));
    }

} catch (Exception e) {
    log.error("Failed to process transient DLT message", e);

    // Move to permanent DLT
    sendToPermanentDlt(record, e);
}

}

/**
 * DLT handler using @DltHandler annotation (Spring Kafka 2.5+)
 */
@Component
public static class AnnotationBasedDltHandler {

    @Autowired
    private DltRepairService repairService;

    /**
     * Primary message listener
     */
    @KafkaListener(
        topics = "annotated-orders",
        groupId = "annotated-order-processor"
    )
    public void processOrder(OrderEvent order) {
        log.info("Processing order: {}", order.getOrderid());

        // Simulate processing that might fail
        if (order.getOrderid().contains("FAIL")) {

```

```

        throw new ValidationException("Simulated validation failure");
    }

    log.info("Successfully processed order: {}", order.getOrderId());
}

/**
 * DLT handler for the same listener group
 */
@DltHandler
public void handleDltOrder(
    OrderEvent order,
    @Header(KafkaHeaders.DLT_EXCEPTION_MESSAGE) String
exceptionMessage,
    @Header(KafkaHeaders.DLT_EXCEPTION_FQCN) String exceptionClass) {

    log.warn("Handling DLT order: orderId={}, exception={}, message={}",
        order.getOrderId(), exceptionClass, exceptionMessage);

    try {
        // Attempt repair based on exception type
        if (exceptionClass.contains("ValidationException")) {
            OrderEvent repairedOrder =
repairService.repairValidation(order);
            if (repairedOrder != null) {
                log.info("Repaired and reprocessing order: {}",
order.getOrderId());
                processOrder(repairedOrder);
                return;
            }
        }

        // If repair fails, log for manual intervention
        log.error("Could not repair DLT order, requires manual
intervention: {}",
            order.getOrderId());

    } catch (Exception e) {
        log.error("DLT handler failed: orderId={}", order.getOrderId(),
e);
    }
}

/**
 * Scheduled DLT reprocessing job
 */
@Scheduled(fixedDelay = 3600000) // Every hour
public void scheduledDltReprocessing() {
    log.info("Starting scheduled DLT reprocessing job");

    try {
        // Get list of DLT topics to process
        List<String> dltTopics = getDltTopicsForReprocessing();

```

```

        for (String dltTopic : dltTopics) {
            processDltTopicForRetry(dltTopic);
        }

    } catch (Exception e) {
        log.error("Scheduled DLT reprocessing failed", e);
    }
}

/**
 * Batch DLT processing for efficiency
 */
@KafkaListener(
    topics = "high-volume.DLT",
    groupId = "dlt-batch-processor",
    containerFactory = "batchListenerContainerFactory"
)
public void processDltBatch(List<ConsumerRecord<String, Object>> records) {
    log.info("Processing DLT batch: size={}", records.size());

    Map<String, List<ConsumerRecord<String, Object>>> groupedByException =
records.stream()
    .collect(Collectors.groupingBy(record -> {
        Header exceptionHeader =
record.headers().lastHeader(KafkaHeaders.DLT_EXCEPTION_FQCN);
        return exceptionHeader != null ? new
String(exceptionHeader.value()) : "Unknown";
    })));

    // Process each exception type separately
    groupedByException.forEach(this::processDltByExceptionType);
}

// Helper methods for DLT processing
private void handleRecoverableDltMessage(OrderEvent order, DltAnalysisResult
analysis, String originalTopic) {
    try {
        switch (analysis.getRecoveryStrategy()) {
            case "FIX_AND_RETRY" -> {
                OrderEvent fixedOrder =
messageRepairService.repairOrder(order, analysis);
                if (fixedOrder != null) {
                    republishToOriginalTopic(fixedOrder, originalTopic);
                }
            }
            case "RETRY_AFTER_DELAY" -> {
                scheduleDelayedRetry(order, originalTopic,
Duration.ofMinutes(30));
            }
            case "MANUAL_REVIEW" -> {
                sendToManualReviewQueue(order, analysis);
            }
            default -> {

```

```
        log.warn("Unknown recovery strategy: {}",
analysis.getRecoveryStrategy());
    }
}
} catch (Exception e) {
    log.error("Recovery failed for order: {}", order.getOrderid(), e);
}
}

private void handleNonRecoverableDltMessage(OrderEvent order,
DltAnalysisResult analysis) {
    // Log to metrics and alerting systems
    meterRegistry.counter("kafka.dlt.nonrecoverable",
        Tags.of("reason", analysis.getFailureReason()))
        .increment();

    // Store in database for audit trail
    dltAnalysisService.recordNonRecoverableFailure(order, analysis);

    // Alert operations team if needed
    if (analysis.requiresImmedateAttention()) {
        alertOperationsTeam(order, analysis);
    }
}

private void sendToManualReview(ConsumerRecord<String, OrderEvent> record,
Exception e) {
    ManualReviewTask reviewTask = ManualReviewTask.builder()
        .originalTopic(record.topic())
        .messageKey(record.key())
        .messageValue(record.value())
        .processingException(e.getMessage())
        .timestamp(System.currentTimeMillis())
        .priority("HIGH")
        .build();

    kafkaTemplate.send("manual-review-queue", reviewTask);
}

private boolean isSystemHealthy() {
    // Check system health indicators
    return true; // Placeholder
}

private boolean canRetryMessage(Object message) {
    // Business logic to determine if message can be retried
    return true; // Placeholder
}

private void sendToExpiredDlt(ConsumerRecord<String, Object> record) {
    String expiredDltTopic = record.topic().replace(".DLT", ".expired.DLT");
    kafkaTemplate.send(expiredDltTopic, record.key(), record.value());
}
```

```
private void sendToDltWithDelay(ConsumerRecord<String, Object> record,
Duration delay) {
    ProducerRecord<Object, Object> delayedRecord = new ProducerRecord<>(
        record.topic(), record.key(), record.value());

    delayedRecord.headers().add("retry-after",
        String.valueOf(System.currentTimeMillis() +
delay.toMillis()).getBytes());

    kafkaTemplate.send(delayedRecord);
}

private void sendToPermanentDlt(ConsumerRecord<String, Object> record,
Exception e) {
    String permanentDltTopic = record.topic().replace(".transient.DLT",
".permanent.DLT");

    ProducerRecord<Object, Object> permanentRecord = new ProducerRecord<>(
        permanentDltTopic, record.key(), record.value());

    permanentRecord.headers().add("permanent-failure-reason",
e.getMessage().getBytes());
    permanentRecord.headers().add("permanent-failure-timestamp",
        String.valueOf(System.currentTimeMillis()).getBytes());

    kafkaTemplate.send(permanentRecord);
}

private void republishToOriginalTopic(OrderEvent order, String originalTopic)
{
    ProducerRecord<Object, Object> republishRecord = new ProducerRecord<>(
        originalTopic, order.getOrderid(), order);

    republishRecord.headers().add("dlt-recovered", "true".getBytes());
    republishRecord.headers().add("recovery-timestamp",
        String.valueOf(System.currentTimeMillis()).getBytes());

    kafkaTemplate.send(republishRecord);

    log.info("Republished recovered message: orderId={}, topic={}",
        order.getOrderid(), originalTopic);
}

private void scheduleDelayedRetry(OrderEvent order, String originalTopic,
Duration delay) {
    // Implementation would use a scheduler or delayed message system
    log.info("Scheduled delayed retry: orderId={}, delay={}min",
        order.getOrderid(), delay.toMinutes());
}

private void sendToManualReviewQueue(OrderEvent order, DltAnalysisResult
analysis) {
    ManualReviewTask task = ManualReviewTask.builder()
        .orderid(order.getOrderid())
```

```

        .analysisResult(analysis)
        .priority(analysis.requiresImmediateAttention() ? "HIGH" : "MEDIUM")
        .build();

    kafkaTemplate.send("manual-review-tasks", task);
}

private List<String> getDltTopicsForReprocessing() {
    // Implementation would query Kafka admin client for DLT topics
    return Arrays.asList("orders.transient.DLT", "payments.transient.DLT");
}

private void processDltTopicForRetry(String dltTopic) {
    log.info("Processing DLT topic for retry: {}", dltTopic);
    // Implementation would consume from DLT topic and apply retry logic
}

private void processDltByExceptionType(String exceptionType,
                                       List<ConsumerRecord<String, Object>>
records) {
    log.info("Processing DLT batch by exception type: type={}, count={}",
            exceptionType, records.size());

    // Apply exception-specific recovery strategies
    for (ConsumerRecord<String, Object> record : records) {
        try {
            // Process based on exception type
            processRecordByExceptionType(record, exceptionType);
        } catch (Exception e) {
            log.error("Failed to process DLT record: topic={}, offset={}",
                    record.topic(), record.offset(), e);
        }
    }
}

private void processRecordByExceptionType(ConsumerRecord<String, Object>
record, String exceptionType) {
    switch (exceptionType) {
        case "ValidationException" -> attemptValidationFix(record);
        case "ExternalServiceException" -> attemptServiceRetry(record);
        case "DatabaseException" -> attemptDatabaseRetry(record);
        default -> logUnknownExceptionType(record, exceptionType);
    }
}

private void attemptValidationFix(ConsumerRecord<String, Object> record) {
    // Attempt to fix validation issues
    log.debug("Attempting validation fix for record: topic={}, offset={}",
            record.topic(), record.offset());
}

private void attemptServiceRetry(ConsumerRecord<String, Object> record) {
    // Retry external service calls
    log.debug("Attempting service retry for record: topic={}, offset={}",

```

```

        record.topic(), record.offset());
    }

    private void attemptDatabaseRetry(ConsumerRecord<String, Object> record) {
        // Retry database operations
        log.debug("Attempting database retry for record: topic={}, offset={}",
            record.topic(), record.offset());
    }

    private void logUnknownExceptionType(ConsumerRecord<String, Object> record,
        String exceptionType) {
        log.warn("Unknown exception type in DLT: type={}, topic={}, offset={}",
            exceptionType, record.topic(), record.offset());
    }

    private void alertOperationsTeam(OrderEvent order, DltAnalysisResult analysis)
    {
        log.error("🚨 ALERT: Non-recoverable DLT message requires immediate
        attention: orderId={}, reason={}",
            order.getOrderid(), analysis.getFailureReason());
        // Integration with alerting systems (PagerDuty, Slack, etc.)
    }

    @Autowired
    private MeterRegistry meterRegistry;
}

// Supporting classes for DLT processing
@Service
public class DltAnalysisService {

    public DltAnalysisResult analyzeDltMessage(ConsumerRecord<String, OrderEvent>
        record,
                                           String exceptionClass, String
        exceptionMessage) {

        DltAnalysisResult.DltAnalysisResultBuilder builder =
        DltAnalysisResult.builder()
            .messageId(record.key())
            .exceptionType(exceptionClass)
            .exceptionMessage(exceptionMessage);

        // Analyze based on exception type
        if ("ValidationException".equals(exceptionClass)) {
            return builder
                .recoverable(true)
                .recoveryStrategy("FIX_AND_RETRY")
                .confidence(0.8)
                .build();
        } else if ("ExternalServiceException".equals(exceptionClass)) {
            return builder
                .recoverable(true)
                .recoveryStrategy("RETRY_AFTER_DELAY")

```

```

        .confidence(0.9)
        .build();

    } else if ("DeserializationException".equals(exceptionClass)) {
        return builder
            .recoverable(false)
            .failureReason("POISON_PILL")
            .requiresImmediateAttention(true)
            .build();

    } else {
        return builder
            .recoverable(false)
            .failureReason("UNKNOWN_EXCEPTION")
            .requiresImmediateAttention(false)
            .build();
    }
}

public void recordNonRecoverableFailure(OrderEvent order, DltAnalysisResult
analysis) {
    log.info("Recording non-recoverable failure: orderId={}, reason={}",
        order.getOrderId(), analysis.getFailureReason());
    // Store in database or metrics system
}

}

@Service
public class MessageRepairService {

    public OrderEvent repairOrder(OrderEvent order, DltAnalysisResult analysis) {
        try {
            OrderEvent.OrderEventBuilder builder = order.toBuilder();

            // Apply repairs based on analysis
            if ("ValidationException".equals(analysis.getExceptionType())) {
                // Fix common validation issues
                if (order.getCustomerId() == null) {
                    builder.customerId("UNKNOWN_CUSTOMER");
                }

                if (order.getAmount() == null) {
                    builder.amount(java.math.BigDecimal.ZERO);
                }

                if (order.getStatus() == null) {
                    builder.status("PENDING_REPAIR");
                }
            }

            OrderEvent repaired = builder.build();
            log.info("Repaired order: original={}, repaired={}", order, repaired);

            return repaired;
        }
    }
}

```



```

        } catch (Exception e) {
            log.error("Failed to repair order: {}", order.getOrderid(), e);
            return null;
        }
    }
}

@Service
public class DltRepairService {

    public OrderEvent repairValidation(OrderEvent order) {
        // Implement validation repair logic
        return order.toBuilder()
            .status("REPAIRED")
            .build();
    }
}

// Supporting data structures
@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
class DltAnalysisResult {
    private String messageId;
    private String exceptionType;
    private String exceptionMessage;
    private boolean recoverable;
    private String recoveryStrategy;
    private String failureReason;
    private double confidence;
    private boolean requiresImmediateAttention;
}

@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
class ManualReviewTask {
    private String originalTopic;
    private String messageKey;
    private Object messageValue;
    private String processingException;
    private long timestamp;
    private String priority;
    private String orderId;
    private DltAnalysisResult analysisResult;
}

```

This comprehensive section covers Dead Letter Topics (DLT) configuration, publishing, and retry patterns with extensive Java examples for production scenarios. The guide continues with the legacy

SeekToCurrentErrorHandler section next.