Spring Kafka Transactions: Part 3 - Outbox Processor, Best Practices & Production Guide

Final part of the comprehensive guide covering outbox event processing, comparisons, best practices, CLI operations, and version highlights.

Outbox Event Processor Implementation

Simple Explanation: The Outbox Event Processor is a separate component that monitors the outbox table for new events and publishes them to Kafka. It ensures reliable, at-least-once delivery of events while maintaining the transactional consistency provided by the outbox pattern.

Advanced Outbox Event Processor

```
/**
 * Comprehensive outbox event processor with retry and monitoring
@Component
@lombok.extern.slf4j.Slf4j
public class OutboxEventProcessor {
    @Autowired
    private OutboxEventRepository outboxEventRepository;
    @Autowired
    private KafkaTemplate<String, Object> kafkaTemplate;
    @Autowired
    private ObjectMapper objectMapper;
    @Autowired
    private MeterRegistry meterRegistry;
    private final int maxRetryAttempts = 5;
    private final Duration processingInterval = Duration.ofSeconds(10);
    private final int batchSize = 100;
    /**
     * Scheduled outbox event processing with batch optimization
    @Scheduled(fixedDelayString = "#
{@outboxProcessorConfig.getProcessingIntervalMs()}")
    @Transactional
    public void processOutboxEvents() {
        log.debug("Starting outbox event processing batch");
        Timer.Sample sample = Timer.start(meterRegistry);
```

```
try {
            // Fetch pending events in batches
            Pageable pageable = PageRequest.of(∅, batchSize);
            List<OutboxEvent> pendingEvents = outboxEventRepository
.findByStatusOrderByCreatedAt(OutboxEvent.OutboxEventStatus.PENDING, pageable);
            if (pendingEvents.isEmpty()) {
                log.trace("No pending outbox events to process");
                return;
            }
            log.info("Processing outbox events batch: size={}",
pendingEvents.size());
            int successCount = 0;
            int failureCount = 0;
            // Process each event
            for (OutboxEvent event : pendingEvents) {
                try {
                    boolean processed = processOutboxEvent(event);
                    if (processed) {
                        successCount++;
                    } else {
                        failureCount++;
                } catch (Exception e) {
                    log.error("Error processing outbox event: eventId={}",
event.getId(), e);
                    failureCount++;
                }
            }
            // Update metrics
meterRegistry.counter("outbox.events.processed.success").increment(successCount);
meterRegistry.counter("outbox.events.processed.failure").increment(failureCount);
            meterRegistry.gauge("outbox.events.pending.count",
outboxEventRepository.countByStatus(OutboxEvent.OutboxEventStatus.PENDING));
            log.info("Outbox event processing completed: success={}, failures={}",
successCount, failureCount);
        } catch (Exception e) {
            log.error("Error in outbox event processing batch", e);
            meterRegistry.counter("outbox.events.batch.error").increment();
        } finally {
            sample.stop(Timer.builder("outbox.events.processing.duration")
                .register(meterRegistry));
```

```
}
     * Process individual outbox event with retry logic
    @Transactional
    public boolean processOutboxEvent(OutboxEvent event) {
        log.debug("Processing outbox event: eventId={}, eventType={}",
event.getId(), event.getEventType());
        try {
            // Mark as processing
            event.setStatus(OutboxEvent.OutboxEventStatus.PROCESSING);
            event.setProcessingAttempts(event.getProcessingAttempts() + 1);
            outboxEventRepository.save(event);
            // Deserialize event data
            Object eventData = deserializeEventData(event.getEventData(),
event.getEventType());
            // Create Kafka producer record
            ProducerRecord<String, Object> producerRecord = new ProducerRecord<>(
                event.getKafkaTopic(),
                event.getKafkaKey(),
                eventData
            );
            // Add headers for traceability
            addEventHeaders(producerRecord, event);
            // Send to Kafka
            ListenableFuture<SendResult<String, Object>> future =
kafkaTemplate.send(producerRecord);
            // Wait for result with timeout
            SendResult<String, Object> result = future.get(30, TimeUnit.SECONDS);
            // Mark as processed
            event.setStatus(OutboxEvent.OutboxEventStatus.PROCESSED);
            event.setProcessedAt(Instant.now());
            event.setLastError(null);
            outboxEventRepository.save(event);
            log.info("Outbox event published successfully: eventId={}, topic={},
partition={}, offset={}",
                event.getId(), event.getKafkaTopic(),
                result.getRecordMetadata().partition(),
                result.getRecordMetadata().offset());
            return true;
        } catch (Exception e) {
```

```
log.error("Failed to process outbox event: eventId={}", event.getId(),
e);
            // Handle failure
            return handleEventProcessingFailure(event, e);
        }
    }
     * Handle event processing failures with retry logic
    private boolean handleEventProcessingFailure(OutboxEvent event, Exception
exception) {
        String errorMessage = exception.getMessage();
        if (errorMessage != null && errorMessage.length() > 500) {
            errorMessage = errorMessage.substring(0, 500);
        event.setLastError(errorMessage);
        if (event.getProcessingAttempts() >= maxRetryAttempts) {
            // Max retries exceeded - mark as failed
            event.setStatus(OutboxEvent.OutboxEventStatus.FAILED);
            outboxEventRepository.save(event);
            log.error("Outbox event failed after {} attempts: eventId={}, error=
{}",
                maxRetryAttempts, event.getId(), errorMessage);
            // Send to dead letter queue or alert system
            handleFailedEvent(event, exception);
            return false;
        } else {
            // Reset to pending for retry
            event.setStatus(OutboxEvent.OutboxEventStatus.PENDING);
            outboxEventRepository.save(event);
            log.warn("Outbox event will be retried: eventId={}, attempt={}, error=
{}",
                event.getId(), event.getProcessingAttempts(), errorMessage);
            return false;
        }
    }
     * Retry failed events that haven't exceeded max attempts
    @Scheduled(fixedDelayString = "#
{@outboxProcessorConfig.getRetryIntervalMs()}")
    @Transactional
```

```
public void retryFailedEvents() {
        log.debug("Starting failed event retry processing");
        try {
            Pageable pageable = PageRequest.of(0, batchSize / 2); // Smaller batch
for retries
            List<OutboxEvent> failedEvents = outboxEventRepository
                .findFailedEventsForRetry(maxRetryAttempts, pageable);
            if (failedEvents.isEmpty()) {
                log.trace("No failed events to retry");
                return;
            }
            log.info("Retrying failed outbox events: count={}",
failedEvents.size());
            for (OutboxEvent event : failedEvents) {
                try {
                    // Add exponential backoff delay
                    long delayMs =
calculateRetryDelay(event.getProcessingAttempts());
(event.getCreatedAt().plus(Duration.ofMillis(delayMs)).isAfter(Instant.now())) {
                        continue; // Too early to retry
                    }
                    processOutboxEvent(event);
                } catch (Exception e) {
                    log.error("Error retrying failed event: eventId={}",
event.getId(), e);
                }
            }
        } catch (Exception e) {
            log.error("Error in failed event retry processing", e);
        }
    }
    /**
     * Cleanup processed events to prevent table growth
    @Scheduled(cron = "#{@outboxProcessorConfig.getCleanupCron()}")
    @Transactional
    public void cleanupProcessedEvents() {
        log.info("Starting cleanup of processed outbox events");
        try {
            // Delete events processed more than configured retention period ago
            Instant cutoffTime = Instant.now().minus(Duration.ofDays(7)); // 7
days retention
```

```
int deletedCount =
outboxEventRepository.deleteProcessedEventsOlderThan(cutoffTime);
meterRegistry.counter("outbox.events.cleanup.deleted").increment(deletedCount);
            log.info("Cleanup completed: deleted {} processed events older than
{}",
                deletedCount, cutoffTime);
        } catch (Exception e) {
            log.error("Error during outbox event cleanup", e);
            meterRegistry.counter("outbox.events.cleanup.error").increment();
        }
    }
     * Monitor outbox table health and metrics
    @Scheduled(fixedDelayString = "#
{@outboxProcessorConfig.getHealthCheckIntervalMs()}")
    public void monitorOutboxHealth() {
        try {
            // Count events by status
            long pendingCount =
outboxEventRepository.countByStatus(OutboxEvent.OutboxEventStatus.PENDING);
            long processingCount =
outboxEventRepository.countByStatus(OutboxEvent.OutboxEventStatus.PROCESSING);
            long failedCount =
outboxEventRepository.countByStatus(OutboxEvent.OutboxEventStatus.FAILED);
            long processedCount =
outboxEventRepository.countByStatus(OutboxEvent.OutboxEventStatus.PROCESSED);
            // Update metrics
            meterRegistry.gauge("outbox.events.pending.count", pendingCount);
            meterRegistry.gauge("outbox.events.processing.count",
processingCount);
            meterRegistry.gauge("outbox.events.failed.count", failedCount);
            meterRegistry.gauge("outbox.events.processed.count", processedCount);
            // Check for concerning conditions
            if (pendingCount > 1000) {
                log.warn("High number of pending outbox events: {}",
pendingCount);
            }
            if (failedCount > 100) {
                log.warn("High number of failed outbox events: {}", failedCount);
            }
            if (processingCount > 50) {
                log.warn("High number of events stuck in processing: {}",
```

```
processingCount);
            log.debug("Outbox health check: pending={}, processing={}, failed={},
processed={}",
                pendingCount, processingCount, failedCount, processedCount);
        } catch (Exception e) {
            log.error("Error during outbox health monitoring", e);
        }
    }
    // Helper methods
    private Object deserializeEventData(String eventData, String eventType) throws
Exception {
        switch (eventType) {
            case "OrderCreated":
                return objectMapper.readValue(eventData, OrderCreatedEvent.class);
            case "OrderStatusChanged":
                return objectMapper.readValue(eventData,
OrderStatusChangedEvent.class);
            case "PaymentConfirmed":
                return objectMapper.readValue(eventData,
PaymentConfirmedEvent.class);
            case "ShippingRequested":
                return objectMapper.readValue(eventData,
ShippingRequestedEvent.class);
            case "RefundRequested":
                return objectMapper.readValue(eventData,
RefundRequestedEvent.class);
            case "InventoryRestoration":
                return objectMapper.readValue(eventData,
InventoryRestorationEvent.class);
            case "BulkOrderProcessingCompleted":
                return objectMapper.readValue(eventData,
BulkOrderProcessingCompletedEvent.class);
            default:
                // Generic deserialization for unknown event types
                return objectMapper.readValue(eventData, Map.class);
        }
    }
    private void addEventHeaders(ProducerRecord<String, Object> producerRecord,
OutboxEvent event) {
        Headers headers = producerRecord.headers();
        // Add traceability headers
        headers.add("event-id", event.getId().toString().getBytes());
        headers.add("event-type", event.getEventType().getBytes());
        headers.add("aggregate-id", event.getAggregateId().getBytes());
        headers.add("aggregate-type", event.getAggregateType().getBytes());
        headers.add("created-at", event.getCreatedAt().toString().getBytes());
        headers.add("processing-attempt",
```

```
String.valueOf(event.getProcessingAttempts()).getBytes());
        // Add correlation headers for tracing
        headers.add("correlation-id", UUID.randomUUID().toString().getBytes());
        headers.add("source", "outbox-processor".getBytes());
    }
    private void handleFailedEvent(OutboxEvent event, Exception exception) {
        log.error("Handling permanently failed outbox event: eventId={}",
event.getId());
        try {
            // Create failure notification
            OutboxFailureNotification failureNotification =
OutboxFailureNotification.builder()
                .eventId(event.getId())
                .aggregateId(event.getAggregateId())
                .eventType(event.getEventType())
                .failureReason(exception.getMessage())
                .attempts(event.getProcessingAttempts())
                .failedAt(Instant.now())
                .build();
            // Send to dead letter queue or monitoring system
            kafkaTemplate.send("outbox-failures", event.getId().toString(),
failureNotification);
            meterRegistry.counter("outbox.events.permanently.failed",
                "event-type", event.getEventType()).increment();
        } catch (Exception e) {
            log.error("Failed to send failure notification for eventId={}",
event.getId(), e);
        }
    }
    private long calculateRetryDelay(int attemptNumber) {
        // Exponential backoff: 2^attempt * 1000ms
        return (long) (Math.pow(2, attemptNumber) * 1000);
    }
}
 * Outbox processor configuration
@Component
@ConfigurationProperties(prefix = "outbox.processor")
@lombok.Data
public class OutboxProcessorConfig {
    private long processingIntervalMs = 10000; // 10 seconds
    private long retryIntervalMs = 30000; // 30 seconds
    private long healthCheckIntervalMs = 60000; // 1 minute
```

```
private String cleanupCron = "0 0 2 * * *"; // 2 AM daily
    private int batchSize = 100;
    private int maxRetryAttempts = 5;
    private int retentionDays = 7;
}
 * Event data classes for outbox pattern
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class OrderCreatedEvent {
    private String orderId;
    private String customerId;
    private String productId;
    private Integer quantity;
    private BigDecimal totalAmount;
    private Instant createdAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class OrderStatusChangedEvent {
    private String orderId;
    private String previousStatus;
    private String newStatus;
    private String reason;
    private Instant changedAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class PaymentConfirmedEvent {
    private String orderId;
    private BigDecimal amount;
    private Instant confirmedAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class ShippingRequestedEvent {
    private String orderId;
    private String customerId;
    private String productId;
    private Integer quantity;
    private Instant requestedAt;
```

```
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class RefundRequestedEvent {
    private String orderId;
    private BigDecimal amount;
    private String reason;
    private Instant requestedAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class InventoryRestorationEvent {
    private String orderId;
    private String productId;
    private Integer quantity;
    private String reason;
    private Instant requestedAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class BulkOrderProcessingCompletedEvent {
    private String batchId;
    private Integer orderCount;
    private List<String> orderIds;
    private Instant completedAt;
}
@lombok.Data
@lombok.Builder
@lombok.NoArgsConstructor
@lombok.AllArgsConstructor
class OutboxFailureNotification {
    private UUID eventId;
    private String aggregateId;
    private String eventType;
    private String failureReason;
    private Integer attempts;
    private Instant failedAt;
}
```

Ⅲ Comparisons & Trade-offs

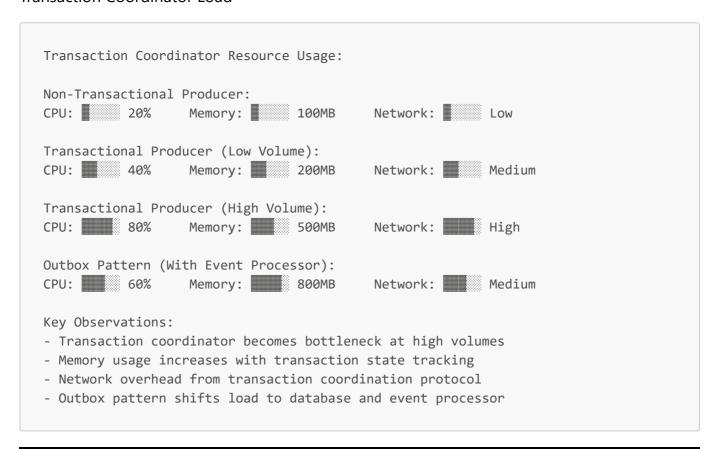
Transaction Approaches Comparison

Approach	Consistency	Performance	Complexity	Use Case
At-Most- Once	X Low	★★★★ Highest	★★★★ Simple	Logging, Analytics
At-Least- Once	★★★ Medium	★★★ High	★ ★ Medium	Most Applications
Exactly-Once	★ ★ ★ ★ Perfect	★ ★ ★ Medium	★ ★ Complex	Financial, Critical
Outbox Pattern	★★★★ Perfect	★ ★ Lower	★ Most Complex	Microservices

Performance Impact Analysis

Configuration	Throughput (msg/sec)	Latency (ms)	Resource Usage
Non-Transactional	100,000	1	Low
Transactional EOS	15,000	10	Medium
Outbox Pattern	8,000	25	High

Transaction Coordinator Load



Common Pitfalls & Best Practices

Critical Anti-Patterns to Avoid

X Transaction Configuration Mistakes

```
// DON'T - Missing or incorrect transactional.id configuration
@Bean
public ProducerFactory<String, Object> badTransactionalProducerFactory() {
    Map<String, Object> props = new HashMap<>();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    // BAD: Missing TRANSACTIONAL_ID_CONFIG
    // BAD: This won't enable transactions!
    props.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
    return new DefaultKafkaProducerFactory<>(props);
}
// DON'T - Non-unique transactional.id across instances
@Bean
public ProducerFactory<String, Object> badNonUniqueTransactionalId() {
    Map<String, Object> props = new HashMap<>();
    props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
    // BAD: Same transactional.id used by multiple instances
    props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "fixed-id");
    // This will cause zombie producer fencing issues
    return new DefaultKafkaProducerFactory<>(props);
}
// DON'T - Wrong isolation level for transactional consumers
@Bean
public ConsumerFactory<String, Object> badConsumerIsolation() {
    Map<String, Object> props = new HashMap<>();
    props.put(ConsumerConfig.BOOTSTRAP SERVERS CONFIG, "localhost:9092");
    // BAD: Default isolation level won't see transactional guarantees
    // props.put(ConsumerConfig.ISOLATION LEVEL CONFIG, "read committed");
    return new DefaultKafkaConsumerFactory<>(props);
}
```

X Error Handling Anti-Patterns

```
// DON'T - Catching transaction exceptions incorrectly
@Service
public class BadTransactionalService {

    @KafkaListener(topics = "orders")
    public void badErrorHandling(@Payload OrderMessage order) {
        try {
            processOrder(order);
            kafkaTemplate.send("processed-orders", order);
        } catch (Exception e) {
            // BAD: Swallowing exceptions prevents transaction rollback
```

```
log.error("Error processing order", e);
    // Transaction will commit instead of rolling back
}

@Transactional("kafkaTransactionManager")
public void badRetryLogic(OrderMessage order) {
    try {
        processOrder(order);
    } catch (ProducerFencedException e) {
        // BAD: Trying to retry a non-recoverable exception
        retryProcessing(order); // This will always fail
    }
}
```

X Outbox Pattern Anti-Patterns

```
// DON'T - Missing transactional boundary in outbox pattern
@Service
public class BadOutboxService {
    public void badOutboxImplementation(CreateOrderRequest request) {
        // BAD: No transaction boundary - dual write problem still exists
        Order order = orderRepository.save(createOrder(request));
        // If this fails, order is saved but event is not
        OutboxEvent event = createOutboxEvent(order);
        outboxEventRepository.save(event);
    }
    // DON'T - Processing outbox events without proper concurrency control
    @Scheduled(fixedDelay = 1000)
    public void badOutboxProcessor() {
        List<OutboxEvent> events = outboxEventRepository.findPendingEvents();
        // BAD: No concurrency control - multiple instances will process same
events
        events.forEach(event -> {
            kafkaTemplate.send(event.getTopic(), event.getData());
            event.setStatus(OutboxEventStatus.PROCESSED);
            outboxEventRepository.save(event);
        });
    }
}
```

Production Best Practices

☑ Optimal Transaction Configuration

```
* ✓ GOOD - Production-ready transactional configuration
 */
@Configuration
@lombok.extern.slf4j.Slf4j
public class ProductionTransactionalConfiguration {
    @Value("${spring.application.name:default-app}")
    private String applicationName;
    @Value("${server.port:8080}")
    private String serverPort;
    @Bean
    public ProducerFactory<String, Object>
productionTransactionalProducerFactory() {
        Map<String, Object> props = new HashMap<>();
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "${kafka.bootstrap-
servers}");
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
StringSerializer.class);
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
JsonSerializer.class);
        // GOOD: Unique transactional.id per application instance
        String transactionalId = applicationName + "-" + serverPort + "-" +
UUID.randomUUID().toString();
        props.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, transactionalId);
        // GOOD: Optimized transaction timeout
        props.put(ProducerConfig.TRANSACTION_TIMEOUT_CONFIG, 30000); // 30 seconds
        // GOOD: Production-optimized settings
        props.put(ProducerConfig.ACKS_CONFIG, "all");
        props.put(ProducerConfig.RETRIES_CONFIG, Integer.MAX_VALUE);
        props.put(ProducerConfig.BATCH_SIZE_CONFIG, 65536);
        props.put(ProducerConfig.LINGER_MS_CONFIG, 10);
        props.put(ProducerConfig.COMPRESSION TYPE CONFIG, "snappy");
        // GOOD: Connection pooling optimization
        props.put(ProducerConfig.CONNECTIONS_MAX_IDLE_MS_CONFIG, 600000);
        log.info("Configured production transactional producer: transactionalId=
{}", transactionalId);
        return new DefaultKafkaProducerFactory<>(props);
    }
    @Bean
    public ConsumerFactory<String, Object>
productionTransactionalConsumerFactory() {
        Map<String, Object> props = new HashMap<>();
```

```
props.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "${kafka.bootstrap-
servers}");
        props.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
StringDeserializer.class);
        props.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
JsonDeserializer.class);
        props.put(ConsumerConfig.GROUP ID CONFIG, "${kafka.consumer.group-id}");
        // GOOD: Essential for exactly-once semantics
        props.put(ConsumerConfig.ISOLATION_LEVEL_CONFIG, "read_committed");
        props.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);
        // GOOD: Production-optimized consumer settings
        props.put(ConsumerConfig.SESSION_TIMEOUT_MS_CONFIG, 30000);
        props.put(ConsumerConfig.HEARTBEAT_INTERVAL_MS_CONFIG, 10000);
        props.put(ConsumerConfig.MAX_POLL_INTERVAL_MS_CONFIG, 300000);
        props.put(ConsumerConfig.MAX POLL RECORDS CONFIG, 100); // Smaller batches
for EOS
        return new DefaultKafkaConsumerFactory<>(props);
    }
}
/**
 * ✓ GOOD - Robust error handling for transactions
 */
@Service
@lombok.extern.slf4j.Slf4j
public class RobustTransactionalService {
    @Autowired
    private KafkaTemplate<String, Object> kafkaTemplate;
    @Autowired
    private MeterRegistry meterRegistry;
    @KafkaListener(
        topics = "orders",
        containerFactory = "transactionalListenerContainerFactory"
    public void processWithRobustErrorHandling(@Payload OrderMessage order,
                                             @Header(KafkaHeaders.OFFSET) long
offset) {
        Timer.Sample sample = Timer.start(meterRegistry);
        try {
            log.info("Processing order transactionally: orderId={}, offset={}",
order.getOrderId(), offset);
            // Business logic processing
            ProcessedOrder processedOrder = processOrderSafely(order);
```

```
// Send to output topic within same transaction
            kafkaTemplate.send("processed-orders", order.getOrderId(),
processedOrder);
            // Update success metrics
            meterRegistry.counter("orders.processed.success").increment();
            log.info("Order processed successfully: orderId={}",
order.getOrderId());
        } catch (NonRetryableException e) {
            // Send to DLQ for manual inspection
            sendToDLQ(order, e);
            meterRegistry.counter("orders.processed.dlq").increment();
            log.error("Order sent to DLQ: orderId={}", order.getOrderId(), e);
        } catch (RetryableException e) {
            // Let transaction rollback for automatic retry
            meterRegistry.counter("orders.processed.retry").increment();
            log.warn("Order processing will be retried: orderId={}",
order.getOrderId(), e);
            throw e;
        } catch (ProducerFencedException | FencedInstanceIdException |
OutOfOrderSequenceException e) {
            // Fatal exceptions - requires application restart
            meterRegistry.counter("orders.processed.fatal").increment();
            log.error("Fatal transaction error - application should restart:
orderId={}", order.getOrderId(), e);
            // Gracefully shutdown application
            initiateGracefulShutdown();
            throw e;
        } catch (Exception e) {
            // Unexpected errors
            meterRegistry.counter("orders.processed.error").increment();
            log.error("Unexpected error processing order: orderId={}",
order.getOrderId(), e);
            throw e;
        } finally {
            sample.stop(Timer.builder("orders.processing.duration")
                .tag("result", "completed")
                .register(meterRegistry));
        }
    }
```

```
* ✓ GOOD - Graceful handling of processing errors
    private ProcessedOrder processOrderSafely(OrderMessage order) {
        try {
            // Validate order
            validateOrder(order);
            // Process business logic
            return executeBusinessLogic(order);
        } catch (ValidationException e) {
            // Validation errors are not retryable
            throw new NonRetryableException("Order validation failed", e);
        } catch (ExternalServiceException e) {
            // External service errors may be retryable
            if (e.isRetryable()) {
                throw new RetryableException("External service temporarily
unavailable", e);
            } else {
                throw new NonRetryableException("External service permanently
failed", e);
            }
        } catch (Exception e) {
            // Unknown errors are retryable by default
            throw new RetryableException("Unknown processing error", e);
        }
    }
    private void sendToDLQ(OrderMessage order, Exception exception) {
        try {
            DLQMessage dlqMessage = DLQMessage.builder()
                .originalMessage(order)
                .errorMessage(exception.getMessage())
                .errorClass(exception.getClass().getName())
                .timestamp(Instant.now())
                .build();
            kafkaTemplate.send("orders-dlq", order.getOrderId(), dlqMessage);
        } catch (Exception e) {
            log.error("Failed to send message to DLQ: orderId={}",
order.getOrderId(), e);
        }
    }
    private void validateOrder(OrderMessage order) {
        if (order.getOrderId() == null || order.getOrderId().isEmpty()) {
            throw new ValidationException("Order ID is required");
        if (order.getAmount() == null ||
order.getAmount().compareTo(BigDecimal.ZERO) <= ∅) {
```

```
throw new ValidationException("Order amount must be positive");
        }
    }
    private ProcessedOrder executeBusinessLogic(OrderMessage order) {
        // Business logic implementation
        return ProcessedOrder.builder()
            .orderId(order.getOrderId())
            .status("PROCESSED")
            .processedAt(Instant.now())
            .build();
    }
    private void initiateGracefulShutdown() {
        // Implementation would trigger application shutdown
        log.error("Initiating graceful application shutdown due to fatal Kafka
error");
   }
}
```

% CLI Commands and Monitoring

Kafka CLI Commands for Transaction Monitoring

```
#!/bin/bash
# Monitor transaction coordinator
kafka-log-dirs.sh \
  --bootstrap-server localhost:9092 \
  --topic-list __transaction_state \
  --describe
# Check transaction state topic
kafka-topics.sh \
  --bootstrap-server localhost:9092 \
  --describe \
  --topic __transaction_state
# Monitor consumer group with exactly-once semantics
kafka-consumer-groups.sh \
  --bootstrap-server localhost:9092 \
  --group exactly-once-group \
  --describe
# Check producer metrics for transactional producer
kafka-run-class.sh kafka.tools.JmxTool \
  --object-name kafka.producer:type=producer-metrics,client-id=transactional-
producer \
  --jmx-url service:jmx:rmi:///jndi/rmi://localhost:9999/jmxrmi
```

```
# Monitor transaction coordinator metrics
kafka-run-class.sh kafka.tools.JmxTool \
  --object-name kafka.coordinator.transaction:type=TransactionCoordinator \
  --jmx-url service:jmx:rmi:///jndi/rmi://localhost:9999/jmxrmi
# Test transactional producer
kafka-console-producer.sh \
  --bootstrap-server localhost:9092 \
  --topic test-transactions \
  --producer-property transactional.id=test-tx-1
# Test read committed consumer
kafka-console-consumer.sh \
  --bootstrap-server localhost:9092 \
  --topic test-transactions \
  --from-beginning \
  --isolation-level read_committed
# Monitor outbox table (PostgreSQL)
psql -h localhost -d myapp -c "
SELECT
    status.
    COUNT(*) as count,
   MIN(created_at) as oldest,
   MAX(created_at) as newest
FROM outbox_events
GROUP BY status;
```

JMX Metrics for Transaction Monitoring

```
# Key transaction metrics to monitor

# Producer transaction metrics
kafka.producer.transaction-send-rate: Rate of transactional sends
kafka.producer.transaction-duration-avg: Average transaction duration
kafka.producer.transaction-duration-max: Maximum transaction duration

# Consumer transaction metrics
kafka.consumer.transaction-read-rate: Rate of transactional reads
kafka.consumer.committed-time-ns-total: Time spent committing

# Coordinator metrics
kafka.coordinator.transaction.transaction-abort-rate: Transaction abort rate
kafka.coordinator.transaction.transaction-commit-rate: Transaction commit rate
kafka.coordinator.transaction.active-transactions: Current active transactions

# Application metrics (custom)
outbox.events.pending.count: Pending outbox events
outbox.events.processing.duration: Processing time per event
```

orders.processed.success: Successfully processed orders orders.processed.error: Failed order processing

∀ersion Highlights

Spring Kafka Transaction Evolution

Version	Release	Key Transaction Features
3.2.x	2024	Enhanced EOS performance, batch transaction improvements
3.1.x	2024	KRaft transaction coordinator support, observability enhancements
3.0.x	2023	EOSMode.V2 only, native compilation support
2.9.x	2022	KafkaAwareTransactionManager improvements, batch EOS
2.8.x	2022	Enhanced exactly-once semantics, transaction error handling
2.7.x	2021	@Transactional improvements, ChainedKafkaTransactionManager
2.6.x	2021	KIP-447 fetch-offset-request fencing (EOSMode.V2)
2.5.x	2020	EOSMode.V2 introduction, improved transaction coordinator
2.4.x	2020	KafkaAwareTransactionManager, better JTA integration
2.3.x	2019	Initial transaction support, basic exactly-once semantics

Critical Apache Kafka Version Dependencies

Kafka 3.5+ (2023):

- KRaft Transaction Coordinator: Native ZooKeeper-free transaction coordination
- Enhanced Performance: Improved transaction throughput and latency
- Better Observability: Enhanced metrics and monitoring capabilities

Kafka 2.8+ (2021):

- **KRaft Mode Support**: Transition away from ZooKeeper dependency
- Transaction Performance: Significant improvements in transaction processing

Kafka 2.5+ (2020):

- EOSMode.V2: Fetch-offset-request fencing eliminates need for multiple producers
- **Producer Scalability**: Single producer can handle multiple input partitions
- Performance Improvement: Reduced overhead in exactly-once processing

Spring Boot Compatibility Matrix:

Spring Boot	Spring Kafka	Min Kafka Broker	EOS Support
3.2.x	3.1.x	2.8+	✓ V2 Only

Spring Boot	Spring Kafka	Min Kafka Broker	EOS Support
3.1.x	3.0.x	2.8+	✓ V2 Only
3.0.x	2.9.x	2.8+	✓ V2 Only
2.7.x	2.8.x	2.5+	✓ V2 Only
2.6.x	2.7.x	2.5+	✓ V1 + V2

& Production Checklist

Essential Configuration Checklist

- **Unique transactional.id** per application instance
- **Solution.level=read_committed** for transactional consumers
- **enable.auto.commit=false** for exactly-once consumers
- Appropriate transaction.timeout configuration
- **WafkaAwareTransactionManager** for consumer offset management
- Proper error handling for ProducerFencedException
- Dead Letter Topic strategy for non-retryable failures
- Monitoring and alerting for transaction metrics
- W Health checks for transaction coordinator connectivity
- Graceful shutdown handling for transactional applications

Outbox Pattern Checklist

- Single database transaction for business logic + outbox event
- Idempotent event processor with proper concurrency control
- Retry logic with exponential backoff for failed events
- Dead letter handling for permanently failed events
- **Event cleanup** strategy to prevent table growth
- Monitoring for outbox table health and processing metrics
- **Backup and recovery** procedures for outbox events

Key Takeaways

When to Use Each Approach

- 1. **Producer Transactions**: Single-service, multi-topic publishing with exactly-once guarantees
- 2. Consumer Transactions: Read-process-write patterns requiring exactly-once semantics
- 3. Outbox Pattern: Cross-service consistency with database + Kafka atomicity
- 4. Chain Transactions: Complex multi-step workflows with rollback requirements

Performance Considerations

- Exactly-Once Semantics: 5-10x throughput reduction vs at-least-once
- Transaction Coordinator: Can become bottleneck at high volumes
- Outbox Pattern: Additional database load + eventual consistency delay

• Memory Usage: Increased with transaction state tracking

Critical Success Factors

- 1. Unique Transaction IDs: Prevent zombie producer issues
- 2. **Proper Error Handling**: Distinguish retryable vs fatal exceptions
- 3. Monitoring: Track transaction metrics and failures
- 4. **Testing**: Validate failure scenarios and recovery procedures
- 5. Capacity Planning: Account for transaction coordinator overhead

Last Updated: September 2025

Spring Kafka Version Coverage: 3.2.x Apache Kafka Compatibility: 3.5.x

Spring Boot Version: 3.2.x

This comprehensive guide provides production-ready patterns for implementing exactly-once semantics in Spring Kafka applications, from basic transactions to advanced outbox patterns, ensuring data consistency and reliability in distributed systems.

[613] [614] [615]