

Kafka Internals Cheat Sheet - Master Level

6.1 Storage Engine

Log Segments, Index Files

Definition Log segments are the fundamental storage units in Kafka, representing immutable append-only files that store record batches with configurable size and time-based rolling policies, accompanied by index files enabling efficient offset-based lookups and time-based queries. Each partition maintains multiple active and historical segments with specialized index structures including offset indexes (.index), timestamp indexes (.timeindex), and transaction indexes (.txnindex) optimizing various access patterns and query requirements.

Key Highlights Segments roll over based on configurable size thresholds (default 1GB), time intervals (default 7 days), or administrative operations with automatic index file creation and maintenance for each segment. Index files use sparse indexing with configurable index intervals reducing storage overhead while maintaining logarithmic lookup performance for random access patterns. Multiple index types enable optimized access patterns including offset-based seeks, timestamp-based queries, and transaction boundary identification supporting various consumer and administrative operations.

Responsibility / Role Log segments handle persistent storage of record batches with immutable append-only semantics, managing disk I/O optimization through sequential writes and memory-mapped file access for read operations. Index files coordinate efficient record lookup operations eliminating need for sequential scans during consumer seeks, administrative queries, and retention policy enforcement procedures. Segment management includes automatic rolling, cleanup coordination with retention policies, and integration with replication systems for follower synchronization and recovery operations.

Underlying Data Structures / Mechanism Segment files use binary format with record batches containing headers, compression metadata, and variable-length record arrays with CRC checksums for corruption detection and prevention. Index files implement sparse binary search structures with fixed-size entries mapping logical offsets to physical file positions with configurable density affecting storage overhead and lookup performance. Memory-mapped file access provides efficient random read access leveraging operating system page cache while append operations use buffered writes with configurable flush policies for durability guarantees.

Advantages Sequential write patterns optimize disk I/O performance achieving near-hardware limits for write throughput while memory-mapped reads leverage operating system page cache for efficient random access patterns. Immutable segment design enables safe concurrent access between writers and readers without locking overhead, supporting high-concurrency scenarios with predictable performance characteristics. Sparse indexing provides efficient lookup capabilities with minimal storage overhead, typically consuming less than 1% of total log storage for index metadata and structures.

Disadvantages / Trade-offs Large segment files can delay retention policy enforcement and increase recovery time during broker startup as segments must be validated and indexed before becoming available for serving requests. Index file corruption requires segment rebuilding causing temporary unavailability and increased I/O overhead during recovery procedures affecting broker performance and availability. Memory-mapped file limitations on 32-bit systems and large partition counts can cause virtual memory exhaustion requiring careful system configuration and monitoring.

Corner Cases Segment rolling during high-throughput scenarios can cause temporary write pauses as new segments initialize and index files create, potentially affecting producer latency and throughput during rolling operations. Filesystem limitations including maximum file count per directory and file size limitations can affect segment management requiring careful directory structure design and monitoring. Index file corruption during unclean shutdown can cause segment unavailability until index rebuilding completes, potentially affecting consumer access and partition availability.

Limits / Boundaries Maximum segment size typically ranges from 100MB to 10GB with 1GB default balancing retention granularity with file management overhead and recovery time characteristics. Index interval configuration affects memory usage and lookup performance, typically ranging from 1KB to 64KB with 4KB default balancing index accuracy with storage overhead. Maximum segments per partition limited by filesystem constraints and memory usage for segment metadata, typically supporting thousands of segments per partition depending on configuration.

Default Values Segment size defaults to 1GB (`log.segment.bytes=1073741824`), segment rolling time is 7 days (`log.roll.hours=168`), and index interval is 4KB (`log.index.interval.bytes=4096`). Index file size defaults to 10MB (`log.index.size.max.bytes=10485760`) with automatic growth as needed for segment coverage.

Best Practices Configure segment sizes based on retention policies and recovery time objectives, with smaller segments enabling more granular retention at cost of increased file management overhead. Monitor segment rolling frequency and index file health as indicators of broker performance and storage system health, implementing alerting for abnormal rolling patterns or index corruption events. Optimize filesystem configuration including directory structure, file system type, and mount options for sequential write performance and memory-mapped file efficiency supporting high-throughput Kafka deployments.

Retention Policies (Time, Size, Log Compaction)

Definition Retention policies define data lifecycle management strategies including time-based retention (deleting data older than specified duration), size-based retention (maintaining maximum partition size), and log compaction (retaining latest value per key) with configurable cleanup intervals and enforcement mechanisms. These policies operate independently or in combination providing flexible data management strategies optimized for various use cases including event streaming, state management, and regulatory compliance requirements.

Key Highlights Time-based retention uses record timestamps (`CreateTime` or `LogAppendTime`) with configurable retention periods ranging from minutes to years, automatically deleting eligible segments during cleanup cycles. Size-based retention monitors partition size with configurable thresholds triggering segment deletion to maintain storage limits while log compaction implements key-based deduplication preserving latest values indefinitely. Multiple retention policies can operate simultaneously with cleanup coordination ensuring consistent application across different retention strategies and policy combinations.

Responsibility / Role Retention enforcement coordinates segment deletion, compaction scheduling, and cleanup timing with broker resource management to minimize performance impact during retention operations. Time and size-based policies handle segment-level deletion decisions while log compaction manages record-level deduplication requiring sophisticated coordination between cleanup processes and active partition usage. Policy enforcement integrates with replication systems ensuring retention consistency across replica sets and coordinating cleanup timing with follower synchronization requirements.

Underlying Data Structures / Mechanism Cleanup processes use background threads with configurable scheduling intervals scanning partition segments for retention eligibility using segment metadata and timestamp information. Log compaction implements two-phase cleaning with dirty segment identification and clean segment merging using efficient merge algorithms preserving latest values while maintaining offset sequences. Retention coordination uses partition-level locking and segment reference counting preventing cleanup of actively used segments while enabling concurrent cleanup across multiple partitions.

Advantages Flexible retention policies enable optimization for various use cases from short-term event processing to long-term state storage without requiring external data management systems or complex application logic. Automatic cleanup eliminates manual data management overhead while maintaining predictable storage usage and performance characteristics across varying data volumes and retention requirements. Log compaction provides efficient state storage with automatic deduplication enabling indefinite retention of latest state without unbounded storage growth for key-value use cases.

Disadvantages / Trade-offs Retention policy enforcement can cause I/O overhead during cleanup cycles potentially affecting broker performance and increasing latency for concurrent producer and consumer operations. Log compaction overhead includes CPU usage for merge operations and temporary storage requirements during cleaning cycles potentially consuming significant resources during high-activity periods. Policy coordination complexity increases with multiple retention strategies requiring careful tuning to prevent conflicts and ensure consistent cleanup behavior across different policy types.

Corner Cases Clock skew between producers can affect time-based retention accuracy as records with future timestamps may be retained longer than intended while records with past timestamps may be deleted prematurely. Log compaction with high update rates can cause cleaning to fall behind data production requiring aggressive compaction scheduling or configuration tuning to maintain effectiveness. Retention policy changes during runtime can cause unexpected data deletion or retention behavior requiring careful coordination during policy updates and potential data migration procedures.

Limits / Boundaries Minimum retention periods typically range from 1 minute to prevent accidental data loss while maximum retention periods are limited by available storage capacity and policy enforcement overhead. Log compaction effectiveness depends on key distribution and update patterns with high-cardinality key spaces potentially causing reduced compaction efficiency and increased storage usage. Cleanup thread count and scheduling intervals affect retention enforcement latency and resource usage, typically ranging from 1-10 cleanup threads depending on partition count and retention activity levels.

Default Values Time-based retention defaults to 7 days (`log.retention.hours=168`), size-based retention is unlimited by default (`log.retention.bytes=-1`), and cleanup interval is 5 minutes (`log.retention.check.interval.ms=300000`). Log compaction is disabled by default (`log.cleanup.policy=delete`) requiring explicit configuration for key-value use cases, and compaction lag threshold defaults to 50% (`log.cleaner.min.cleanable.ratio=0.5`).

Best Practices Configure retention policies based on business requirements, regulatory compliance, and storage capacity with monitoring for cleanup effectiveness and resource usage during retention operations. Use time-based retention for event streaming use cases and log compaction for state management scenarios, avoiding overly aggressive retention that could cause data loss or performance issues. Monitor retention metrics including cleanup frequency, deleted segment counts, and compaction effectiveness to optimize policy configuration and prevent storage issues or performance degradation during cleanup operations.

Tiered Storage (New Feature for Cloud/Offload)

Definition Tiered storage enables automatic offloading of historical log segments to low-cost remote storage systems including cloud object storage (S3, GCS, Azure Blob) while maintaining local hot storage for recent data and active processing workloads. This feature introduces storage tier coordination between local disk storage for performance-critical operations and remote storage for cost-efficient long-term retention with transparent access patterns for consumers and administrative operations.

Key Highlights Automatic tiering policies use configurable thresholds including segment age, local storage utilization, and access patterns to determine eligible segments for remote offloading with background transfer processes. Remote storage integration supports various cloud providers and storage systems with pluggable implementations enabling custom storage backends and optimization for specific deployment environments. Transparent access provides seamless consumer experience with automatic segment retrieval from remote storage when accessing historical data beyond local storage boundaries.

Responsibility / Role Tiered storage coordination manages segment lifecycle including local-to-remote migration decisions, transfer scheduling, and retention policy enforcement across storage tiers with integration to existing cleanup and retention mechanisms. Remote storage operations handle segment upload, download, and caching with optimization for network efficiency, retry logic for transient failures, and coordination with local storage management for capacity planning. Access pattern optimization includes prefetching strategies, caching policies, and intelligent segment placement balancing cost optimization with performance requirements for various consumer access patterns.

Underlying Data Structures / Mechanism Segment metadata tracking extends existing segment management with remote storage location information, transfer status, and access pattern statistics enabling intelligent tiering decisions and efficient retrieval operations. Remote storage clients implement pluggable interfaces supporting various cloud storage APIs with consistent retry logic, authentication handling, and performance optimization strategies across different storage backends. Local cache management uses configurable LRU policies for remote segment caching with automatic eviction based on local storage pressure and access pattern analysis.

Advantages Significant cost reduction for long-term data retention by leveraging low-cost cloud storage while maintaining high-performance local storage for active workloads and recent data access patterns. Scalable storage capacity beyond local disk limitations enabling indefinite retention policies and compliance requirements without proportional infrastructure cost increases. Simplified capacity planning and storage management with automatic tiering reducing operational overhead for storage lifecycle management and capacity forecasting across varying data retention requirements.

Disadvantages / Trade-offs Remote storage access introduces latency penalties for historical data queries potentially affecting consumer performance when accessing data beyond local storage boundaries requiring careful access pattern optimization. Network bandwidth and transfer costs for remote storage operations can create operational overhead and cost implications depending on data access patterns and cloud provider pricing structures. Complexity increases for backup, disaster recovery, and cross-region replication scenarios requiring coordination between local and remote storage systems across multiple availability zones and regions.

Corner Cases Network failures during segment transfer can cause partial uploads requiring cleanup procedures and transfer retry logic to prevent storage inconsistencies and orphaned data in remote storage systems. Remote storage service outages can cause historical data unavailability affecting consumers requiring historical data access, potentially requiring fallback strategies or temporary service degradation. Tiering policy

misconfiguration can cause excessive local storage usage or premature remote storage migration affecting performance and cost optimization objectives.

Limits / Boundaries Remote storage transfer throughput is limited by network bandwidth and cloud storage API rate limits, typically supporting hundreds of MB/s to GB/s depending on deployment configuration and provider capabilities. Local cache capacity for remote segments depends on available disk storage with typical configurations ranging from GB to TB depending on access pattern requirements and cost optimization objectives. Maximum remote storage capacity is generally unlimited but subject to cloud provider quotas and cost considerations for large-scale deployments.

Default Values Tiered storage is disabled by default requiring explicit configuration and setup of remote storage backends, local cache sizing, and tiering policies based on deployment requirements. Transfer policies, retention coordination, and access pattern optimization require application-specific configuration with no universal defaults provided by the framework.

Best Practices Configure tiered storage based on access pattern analysis and cost optimization objectives, implementing monitoring for transfer performance, access latency, and cost metrics to validate tiering effectiveness. Design retention policies considering both local and remote storage costs with optimization for typical consumer access patterns and regulatory compliance requirements. Implement comprehensive monitoring for remote storage operations including transfer success rates, access latency, and network utilization to identify optimization opportunities and potential issues affecting consumer experience and cost efficiency.

6.2 Replication

ISR (In-Sync Replicas)

Definition In-Sync Replicas represent the set of replica brokers that are actively synchronized with the partition leader, maintaining current data state within configurable lag thresholds and participating in availability and consistency decisions for partition operations. ISR membership dynamically adjusts based on replica performance, network conditions, and broker health with automatic promotion and demotion procedures ensuring optimal balance between availability and consistency guarantees.

Key Highlights ISR membership uses configurable lag thresholds including maximum time lag (`replica.lag.time.max.ms`) and fetch frequency requirements determining replica eligibility for ISR participation with automatic adjustment during performance variations. Minimum ISR count (`min.insync.replicas`) controls availability versus consistency trade-offs by requiring specified ISR count for write acknowledgments while ISR shrinking and expansion procedures handle dynamic membership changes during broker failures and recovery scenarios. ISR state coordination uses leader-follower protocols with periodic health checks, lag monitoring, and automatic membership adjustment based on performance characteristics and configured thresholds.

Responsibility / Role ISR management coordinates replica health monitoring including lag measurement, fetch performance tracking, and network connectivity assessment to determine replica eligibility for consistency participation in partition operations. Dynamic membership adjustment handles replica promotion to ISR when synchronization improves and demotion when replicas fall behind configured thresholds with coordination across cluster members for consistent ISR state. Availability decisions use ISR membership for write acknowledgment requirements, partition availability determination during broker failures, and leader election candidate selection ensuring optimal balance between performance and consistency.

Underlying Data Structures / Mechanism ISR tracking uses high-water mark coordination between leader and follower replicas with periodic synchronization status updates and lag measurement based on fetch request timing and offset progression. Membership decisions implement hysteresis algorithms preventing rapid ISR membership changes during temporary performance variations while ensuring responsive adjustment to persistent performance issues. ISR state propagation uses controller coordination and metadata updates ensuring consistent ISR membership information across cluster members and client metadata for optimal routing and acknowledgment decisions.

Advantages Dynamic ISR membership provides automatic adaptation to changing network conditions, broker performance, and infrastructure variations ensuring optimal availability while maintaining consistency guarantees. Flexible consistency models through `min.insync.replicas` configuration enable application-specific trade-offs between availability and durability with real-time adjustment based on cluster health and replica performance. Automatic failure detection and recovery through ISR membership changes eliminate manual intervention requirements during common failure scenarios while maintaining data protection and availability characteristics.

Disadvantages / Trade-offs ISR membership instability during network issues or performance variations can cause availability fluctuations and inconsistent acknowledgment behavior requiring careful threshold tuning to balance responsiveness with stability. Minimum ISR requirements can cause write unavailability during broker failures if insufficient replicas remain in-sync, potentially requiring manual intervention or configuration adjustment during extended outages. ISR coordination overhead increases with replica count and cluster size potentially affecting metadata operation performance and leader election timing during failure scenarios.

Corner Cases Network partitions can cause ISR membership disputes between cluster segments potentially leading to split-brain scenarios requiring careful configuration of session timeouts and failure detection parameters. Broker performance degradation can cause ISR demotion even during successful replication creating availability issues despite functional replica sets requiring performance monitoring and capacity planning. ISR membership changes during high-throughput scenarios can cause temporary acknowledgment delays as membership stabilizes affecting producer performance and latency characteristics.

Limits / Boundaries Maximum replica count per partition typically ranges from 3-7 replicas with ISR membership limited by network capacity, broker performance, and coordination overhead affecting cluster scalability and performance. ISR lag thresholds typically range from seconds to minutes balancing failure detection responsiveness with stability during performance variations and network congestion scenarios. Minimum ISR count typically ranges from 1-3 depending on consistency requirements with higher counts providing stronger consistency at cost of reduced availability during failures.

Default Values ISR lag threshold defaults to 30 seconds (`replica.lag.time.max.ms=30000`), minimum ISR count defaults to 1 (`min.insync.replicas=1`), and ISR membership evaluation occurs every 5 seconds (`replica.lag.time.max.ms`). Default replication factor is 1 providing no fault tolerance requiring explicit configuration for production deployments.

Best Practices Configure ISR parameters based on network characteristics, broker performance, and consistency requirements with monitoring for ISR stability and membership changes as cluster health indicators. Set minimum ISR count based on availability requirements and expected failure scenarios, typically using 2 for balanced deployments providing single-broker fault tolerance. Monitor ISR membership patterns, lag distribution, and stability metrics to identify performance issues, capacity constraints, and optimization opportunities for replica coordination and cluster health.

Replica Fetcher Threads

Definition Replica fetcher threads are specialized background threads running on follower brokers responsible for continuously synchronizing partition data from leader brokers through optimized fetch protocols, batch processing, and connection management. These threads coordinate replication lag minimization, bandwidth optimization, and failure recovery while maintaining order preservation and exactly-once replication semantics across distributed broker infrastructure.

Key Highlights Each follower broker maintains configurable pools of fetcher threads (`num.replica.fetchers`) with intelligent partition assignment, connection pooling to leader brokers, and adaptive fetch sizing based on throughput patterns and network characteristics. Fetch optimization includes request batching, compression coordination, and adaptive fetch sizing balancing latency reduction with throughput maximization while maintaining memory efficiency and network utilization optimization. Failure handling includes automatic retry logic, leader discovery updates, and graceful degradation during network issues or leader broker unavailability with minimal impact on replication progress and cluster coordination.

Responsibility / Role Fetcher threads handle continuous data synchronization from leader brokers including fetch request optimization, batch processing coordination, and offset management ensuring minimal replication lag and optimal network utilization. They coordinate with ISR management systems providing lag metrics, synchronization status, and health indicators enabling dynamic ISR membership decisions and cluster availability coordination. Critical responsibilities include connection management across multiple leader brokers, request pipelining for throughput optimization, and failure recovery procedures maintaining replication progress during various infrastructure failure scenarios.

Underlying Data Structures / Mechanism Fetcher thread pools use work-stealing algorithms for partition assignment optimization with connection reuse and request pipelining maximizing throughput while minimizing connection overhead across distributed leader brokers. Fetch request coordination uses adaptive sizing algorithms adjusting request size based on throughput patterns, network characteristics, and memory pressure with intelligent batching across multiple partitions sharing same leader broker. Replication offset tracking uses local checkpointing with coordination to ISR management and high-water mark progression ensuring consistent replication progress and recovery capabilities.

Advantages Optimized replication throughput through connection pooling, request batching, and adaptive fetch sizing achieving near-network-limit performance for replication traffic while minimizing resource overhead. Automatic failure recovery and leader discovery enable resilient replication during broker failures, network issues, and cluster topology changes without manual intervention or configuration updates. Scalable thread pool architecture provides linear performance improvements with thread count increases enabling optimization for various cluster sizes and replication requirements.

Disadvantages / Trade-offs Fetcher thread resource usage includes memory for fetch buffers, network connections, and thread overhead requiring careful capacity planning and resource allocation balancing replication performance with broker resource utilization. Thread pool sizing affects replication parallelism and resource usage with suboptimal sizing potentially causing replication bottlenecks or excessive resource consumption requiring performance monitoring and tuning. Network failure handling can cause temporary replication delays during leader failover or network partition scenarios affecting ISR membership and cluster availability characteristics.

Corner Cases Leader broker failures during fetch operations can cause temporary replication delays until leader discovery and connection reestablishment complete, potentially affecting ISR membership and

partition availability. Network congestion or broker overload can cause fetch timeouts and retry storms across multiple fetcher threads potentially overwhelming recovering brokers and extending replication delays. Thread pool exhaustion during high partition counts or network issues can cause replication delays requiring thread pool sizing optimization or partition distribution rebalancing.

Limits / Boundaries Fetcher thread count typically ranges from 1-16 threads per broker with optimal configuration depending on partition count, network characteristics, and broker hardware capacity affecting replication parallelism and resource utilization. Fetch request size limits depend on broker memory configuration and network capacity with typical ranges from KB to MB balancing latency with throughput optimization and memory efficiency requirements. Maximum concurrent fetch connections are limited by network and broker capacity with practical limits around hundreds to thousands of connections depending on cluster size and configuration.

Default Values Default fetcher thread count is 1 (`num.replica.fetchers=1`), fetch size defaults to 1MB (`replica.fetch.max.bytes=1048576`), and fetch timeout is 500ms (`replica.fetch.wait.max.ms=500`). Connection idle timeout defaults to 30 seconds with automatic connection management and leader discovery refresh intervals.

Best Practices Configure fetcher thread count based on partition count and replication requirements, typically 1 thread per 100-200 partitions depending on throughput characteristics and network performance. Monitor fetcher thread performance including replication lag, fetch success rates, and resource utilization to optimize thread pool sizing and identify replication bottlenecks. Implement network capacity planning and monitoring for replication traffic ensuring adequate bandwidth for peak replication loads while maintaining performance for client operations and cluster coordination.

Leader & Follower Roles

Definition Leader brokers serve as the authoritative source for partition data handling all client read and write operations, coordinating replica synchronization, and managing partition-level coordination including offset assignment and ISR membership decisions. Follower brokers maintain synchronized copies of partition data through continuous replication from leaders, participating in ISR membership, and serving as standby replicas ready for leader promotion during failure scenarios with automatic coordination and state transition procedures.

Key Highlights Leader election uses controller coordination and ISR membership prioritization ensuring optimal leader selection based on data completeness, broker health, and cluster balance with automatic failover during leader broker failures. Follower synchronization implements continuous fetch protocols with adaptive performance optimization, offset tracking, and ISR participation enabling minimal replication lag and high availability characteristics. Role transitions during broker failures, planned maintenance, or cluster rebalancing use coordination protocols ensuring data consistency, client redirection, and minimal service disruption during leadership changes.

Responsibility / Role Leader brokers coordinate all partition operations including client request handling, replica synchronization coordination, ISR membership management, and high-water mark progression ensuring consistency and availability for partition operations. Follower brokers maintain data synchronization through continuous replication, ISR participation, and standby readiness for leader promotion with minimal transition overhead and data consistency guarantees. Both roles coordinate with cluster controllers for metadata consistency, partition assignment changes, and failure coordination ensuring optimal cluster operation and fault tolerance characteristics.

Underlying Data Structures / Mechanism Leader state management includes high-water mark coordination, ISR membership tracking, and client request processing with transaction coordination and exactly-once semantics for complex producer operations. Follower state synchronization uses fetch protocols with offset tracking, lag monitoring, and ISR eligibility determination coordinating with leader brokers for optimal replication performance and consistency maintenance. Leadership transition protocols coordinate state transfer, client redirection, and metadata updates ensuring consistent partition availability and data integrity during role changes and failure scenarios.

Advantages Clear leadership model simplifies consistency coordination eliminating complex multi-master scenarios while providing strong consistency guarantees and predictable performance characteristics for partition operations. Automatic leader election and failover enable high availability with minimal manual intervention during broker failures, maintenance scenarios, and infrastructure issues affecting cluster availability. Load distribution through partition leadership balancing across brokers enables optimal resource utilization and performance scaling across cluster members and varying workload patterns.

Disadvantages / Trade-offs Single-leader model creates potential bottlenecks for high-throughput partitions requiring careful partition count planning and leader distribution to prevent hotspots and resource concentration affecting cluster performance. Leader election overhead during failures can cause temporary partition unavailability affecting client operations until new leadership establishes and coordination completes requiring careful timeout tuning and failure detection optimization. Leadership imbalance across brokers can cause uneven resource utilization requiring monitoring and rebalancing procedures to maintain optimal cluster performance and resource distribution.

Corner Cases Simultaneous leader and controller failures can cause extended coordination delays affecting partition availability and cluster coordination requiring careful disaster recovery planning and manual intervention procedures. Network partitions affecting leader-follower communication can cause ISR membership instability and potential data consistency issues requiring careful coordination timeout tuning and partition design considerations. Leadership thrashing during broker performance issues can cause repeated leader elections and coordination overhead affecting cluster stability and performance requiring performance monitoring and capacity planning.

Limits / Boundaries Maximum partitions per leader broker depends on broker capacity and performance requirements with practical limits ranging from hundreds to thousands of leader partitions depending on throughput characteristics and resource allocation. Leader election time typically ranges from seconds to minutes depending on cluster size, coordination complexity, and failure detection timing affecting availability during transition scenarios. Follower lag tolerance affects ISR membership and availability characteristics with typical thresholds ranging from seconds to minutes balancing consistency with availability requirements.

Default Values Leader election timeout defaults to cluster controller session timeout (typically 18 seconds), preferred leader election is enabled by default with automatic rebalancing, and leader imbalance threshold defaults to 10% (`leader.imbalance.per.broker.percentage=10`). Unclean leader election is disabled by default (`unclean.leader.election.enable=false`) preventing data loss during leadership transitions.

Best Practices Monitor leader distribution across brokers and implement automatic rebalancing to prevent leader concentration and resource hotspots affecting cluster performance and availability characteristics. Configure leader election parameters based on network characteristics and availability requirements balancing failure detection speed with stability during temporary issues. Implement capacity planning for leader

partitions considering throughput requirements, resource utilization, and growth projections ensuring optimal performance across cluster scaling and workload evolution.

6.3 Metadata & Consensus

ZooKeeper-based vs KRaft (post-2023+)

Definition ZooKeeper-based coordination uses external ZooKeeper ensemble for cluster metadata management, broker discovery, controller election, and configuration storage requiring separate infrastructure management and operational procedures. KRaft (Kafka Raft) implements self-contained consensus protocol within Kafka brokers eliminating external dependencies through internal metadata management, integrated controller quorum, and simplified operational model with improved scalability characteristics and reduced complexity.

Key Highlights ZooKeeper coordination requires separate 3-5 node ensemble with independent scaling, monitoring, and operational procedures while KRaft integrates metadata management directly into designated Kafka controller brokers reducing infrastructure footprint. KRaft provides improved metadata scalability supporting larger partition counts and faster controller operations compared to ZooKeeper limitations, with controller quorum providing built-in consensus without external coordination systems. Migration from ZooKeeper to KRaft requires careful planning including metadata conversion, rolling upgrades, and verification procedures ensuring operational continuity during transition phases.

Responsibility / Role ZooKeeper coordination handles cluster membership registration, controller election coordination, topic metadata storage, and ACL management through hierarchical namespace with watch notifications and session management for distributed coordination. KRaft mode consolidates these responsibilities within Kafka infrastructure using designated controller brokers for metadata management, consensus coordination, and cluster state distribution eliminating external system dependencies. Both systems provide strong consistency guarantees for metadata operations while differing significantly in operational complexity, scalability characteristics, and infrastructure requirements.

Underlying Data Structures / Mechanism ZooKeeper implementation uses hierarchical znode structure with ephemeral nodes for broker registration, sequential nodes for controller election, and persistent nodes for configuration storage with watch mechanisms for change propagation. KRaft uses Raft consensus algorithm with replicated metadata log stored in internal `__cluster_metadata` topic providing linearizable consistency and automatic leader election through controller quorum coordination. Metadata operations in both systems ensure strong consistency but KRaft provides improved performance characteristics and eliminates ZooKeeper session management complexity.

Advantages ZooKeeper provides battle-tested reliability with extensive operational knowledge, mature tooling ecosystem, and well-understood failure scenarios enabling confident production deployments with comprehensive monitoring and management capabilities. KRaft eliminates external infrastructure dependencies reducing operational complexity, improves metadata operation performance, and provides better scaling characteristics for large clusters with hundreds of thousands of partitions. Both systems provide strong consistency guarantees while KRaft offers simplified deployment, faster controller operations, and reduced infrastructure footprint for modern Kafka deployments.

Disadvantages / Trade-offs ZooKeeper adds operational complexity including separate cluster management, monitoring requirements, and potential coordination bottlenecks affecting Kafka cluster performance and availability during ZooKeeper failures or performance issues. KRaft represents newer technology with evolving

tooling ecosystem, limited operational experience compared to ZooKeeper, and migration complexity for existing ZooKeeper-based deployments requiring careful transition planning. ZooKeeper session management can cause cluster unavailability during network issues while KRaft controller dependencies can affect cluster coordination if controller quorum fails.

Corner Cases ZooKeeper network partitions can cause Kafka cluster unavailability even when broker nodes remain healthy requiring careful network design and ZooKeeper placement considerations for cluster availability. KRaft controller quorum failures require majority availability for cluster coordination potentially causing operational issues during controller maintenance or failures requiring careful controller placement and capacity planning. Migration between coordination systems can encounter edge cases including metadata inconsistencies, timing dependencies, and rollback scenarios requiring comprehensive testing and contingency planning.

Limits / Boundaries ZooKeeper clusters typically support maximum 1 million znodes with recommended cluster sizes of 3-5 nodes, limiting metadata scalability for extremely large Kafka deployments with hundreds of thousands of partitions. KRaft supports significantly larger metadata volumes without znode limitations enabling clusters with millions of partitions while requiring careful controller sizing and resource allocation for metadata operations. Controller quorum size typically ranges from 3-7 nodes balancing availability with coordination overhead and resource utilization for metadata operations.

Default Values ZooKeeper session timeout defaults to 18 seconds with 6-second tick time requiring separate ZooKeeper configuration management, while KRaft mode uses controller quorum configuration with metadata topic replication factor typically set to 3-5 depending on availability requirements. Migration procedures require explicit configuration and operational planning with no automatic transition mechanisms provided between coordination systems.

Best Practices Plan ZooKeeper to KRaft migration during maintenance windows with comprehensive testing, staged rollout procedures, and rollback contingencies ensuring minimal operational disruption during coordination system transitions. Size KRaft controller resources appropriately for expected metadata volume and operation frequency, separating controller nodes from regular broker workloads for optimal performance and resource isolation. Monitor coordination system health including leader election frequency, metadata operation latency, and quorum health as critical cluster availability indicators regardless of coordination system choice.

Controller Quorum

Definition Controller quorum represents the distributed consensus system managing Kafka cluster metadata including partition assignments, topic configurations, broker membership, and administrative operations through replicated state machine coordination among designated controller brokers. The quorum uses Raft consensus protocol ensuring linearizable consistency for metadata operations while providing fault tolerance through majority-based decision making and automatic leader election procedures.

Key Highlights Controller quorum requires majority agreement for metadata changes ensuring consistent cluster state across distributed controller instances with automatic leader election providing seamless failover during controller failures. Metadata replication uses internal `__cluster_metadata` topic with configurable replication factor and retention policies optimized for metadata operation patterns and recovery requirements. Quorum membership uses dedicated controller brokers separate from regular broker workloads enabling resource optimization and failure isolation for metadata operations versus data operations.

Responsibility / Role Controller quorum coordinates all cluster metadata operations including topic creation, partition assignment, configuration updates, and ACL management ensuring consistent state propagation across all cluster members. Leader controller handles active metadata operations while follower controllers maintain synchronized state for automatic failover and load distribution of read-only metadata queries. Quorum coordination includes health monitoring, leader election, and metadata log compaction ensuring optimal performance and storage efficiency for ongoing metadata operations.

Underlying Data Structures / Mechanism Raft implementation uses replicated log structure with monotonically increasing sequence numbers ensuring ordered metadata operations and consistent state machine execution across controller quorum members. Metadata storage uses specialized topic format optimized for configuration data, administrative operations, and state machine snapshots with automatic compaction and retention management. Leader election uses Raft consensus with randomized timeouts preventing split votes while follower coordination includes log replication, heartbeat protocols, and automatic failover coordination.

Advantages Built-in consensus eliminates external coordination system dependencies while providing strong consistency guarantees and automatic failover capabilities for metadata operations without manual intervention. Dedicated controller resources enable optimization for metadata operation patterns while isolating metadata performance from data plane operations improving overall cluster stability and performance predictability. Scalable metadata management supports larger partition counts and faster administrative operations compared to external coordination systems with integrated monitoring and debugging capabilities.

Disadvantages / Trade-offs Controller quorum requires dedicated resources increasing cluster resource requirements and operational complexity for controller management, monitoring, and capacity planning compared to single-controller architectures. Quorum failures require majority availability potentially causing cluster administrative operations to become unavailable during controller maintenance or failures requiring careful capacity planning and operational procedures. Metadata operation performance depends on controller hardware and network characteristics potentially creating bottlenecks during high-frequency administrative operations or large-scale cluster changes.

Corner Cases Network partitions affecting controller quorum can prevent metadata operations even when data plane operations remain functional requiring careful network design and controller placement for operational continuity. Controller resource exhaustion during large-scale metadata operations can cause administrative operation delays requiring capacity monitoring and resource allocation optimization for controller workloads. Simultaneous controller failures beyond majority threshold require manual intervention and potential metadata recovery procedures affecting cluster availability and administrative capabilities.

Limits / Boundaries Controller quorum size typically ranges from 3-7 members with odd numbers recommended for optimal tie-breaking and resource efficiency balancing availability with coordination overhead. Metadata operation throughput depends on controller resources and network characteristics typically supporting thousands of operations per second for routine administrative tasks. Maximum cluster metadata size is primarily limited by controller memory and storage capacity with practical limits supporting millions of partitions and configuration objects.

Default Values Controller quorum size defaults to 3 members for new KRaft deployments, metadata topic replication factor matches quorum size, and election timeout defaults to 1 second

(`controller.quorum.election.timeout.ms=1000`). Metadata log segment size defaults to 128MB with accelerated compaction schedules optimized for configuration data characteristics.

Best Practices Deploy controller quorum on dedicated hardware separate from regular broker workloads ensuring resource isolation and optimal metadata operation performance without interference from data operations. Monitor controller quorum health including leader election frequency, metadata operation latency, and resource utilization as critical indicators of cluster administrative capability and stability. Size controller resources based on expected administrative workload and cluster growth projections ensuring adequate capacity for peak metadata operation periods and cluster scaling scenarios.

Partition Reassignment

Definition Partition reassignment enables dynamic redistribution of partition replicas across cluster brokers for load balancing, capacity expansion, broker decommissioning, and failure recovery through coordinated data migration and metadata updates. The process coordinates replica placement changes, data transfer procedures, and leadership transitions while maintaining partition availability and data consistency throughout reassignment operations.

Key Highlights Reassignment operations use multi-phase coordination including replica addition, data synchronization, leadership transition, and old replica removal ensuring minimal service disruption while redistributing partition data across target brokers. Throttling mechanisms control reassignment bandwidth preventing overwhelming cluster resources during large-scale migration operations while configurable parallelism enables optimization for various cluster sizes and reassignment scenarios. Progress monitoring and cancellation capabilities provide operational control enabling adjustment of reassignment parameters and recovery from problematic reassignment operations.

Responsibility / Role Partition reassignment coordinates with controller systems for metadata consistency, replica management for data transfer coordination, and ISR management for availability maintenance during replica placement changes. Data migration includes bandwidth management, progress tracking, and failure recovery ensuring consistent data transfer while maintaining cluster performance for ongoing operations. Administrative coordination provides progress visibility, parameter adjustment, and cancellation capabilities enabling operational control over complex reassignment operations affecting multiple partitions and brokers.

Underlying Data Structures / Mechanism Reassignment state management uses controller metadata tracking reassignment progress, target replica sets, and coordination status with persistent state enabling recovery from controller failures during ongoing reassignment operations. Data transfer coordination uses replica fetcher optimization with dedicated bandwidth allocation and progress checkpointing ensuring efficient migration while maintaining cluster performance characteristics. Replica lifecycle management coordinates addition, synchronization, promotion, and removal phases with ISR coordination ensuring availability throughout reassignment lifecycle.

Advantages Dynamic load balancing enables cluster optimization without downtime allowing response to changing workload patterns, broker additions, and capacity requirements through automated or administrative-driven partition redistribution. Operational flexibility supports various cluster maintenance scenarios including broker replacement, capacity expansion, and failure recovery with minimal service impact and coordinated data protection. Throttling and parallelism controls enable optimization for various cluster characteristics and operational requirements balancing reassignment speed with cluster performance impact.

Disadvantages / Trade-offs Reassignment operations consume significant network bandwidth and broker resources potentially affecting cluster performance during migration phases requiring careful scheduling and resource planning for large-scale operations. Complex coordination requirements can cause prolonged reassignment operations with potential for failure scenarios requiring recovery procedures and operational intervention during problematic reassignment situations. Administrative overhead includes planning, monitoring, and coordination complexity increasing operational burden for cluster maintenance and optimization procedures.

Corner Cases Broker failures during reassignment can cause partially completed migrations requiring recovery procedures and potential data consistency verification depending on reassignment phase and affected replicas. Network congestion or resource constraints can cause reassignment delays or failures requiring parameter adjustment, throttling configuration, or operational intervention to complete migration operations. Simultaneous reassignments across multiple partitions can overwhelm cluster coordination systems requiring careful scheduling and parallelism management for large-scale reassignment operations.

Limits / Boundaries Maximum concurrent reassignments depend on cluster capacity and network bandwidth with typical limits ranging from tens to hundreds of simultaneous reassignments depending on cluster size and configuration. Reassignment bandwidth throttling typically ranges from MB/s to GB/s per broker balancing migration speed with cluster performance impact requiring monitoring and adjustment based on operational characteristics. Reassignment operation duration scales with data volume and network capacity potentially requiring hours or days for large partition migrations across distributed infrastructure.

Default Values Reassignment throttling defaults to unlimited bandwidth requiring explicit configuration for production operations, reassignment batching defaults to 5 partitions simultaneously, and progress reporting intervals default to 30 seconds. No default reassignment policies are provided requiring administrative configuration and planning for cluster optimization and maintenance operations.

Best Practices Plan reassignment operations during low-traffic periods with appropriate bandwidth throttling preventing performance impact on production workloads while enabling reasonable migration completion times. Monitor reassignment progress including data transfer rates, completion status, and cluster performance impact enabling parameter adjustment and operational intervention when necessary for optimal results. Implement staged reassignment procedures for large-scale operations including verification phases, rollback capabilities, and comprehensive monitoring ensuring successful completion of complex cluster optimization and maintenance operations.