

Link <https://www.confluent.io/blog/turning-the-database-inside-out-with-apache-samza>
<https://www.confluent.io/blog/how-choose-number-topics-partitions-kafka-cluster>
<https://www.confluent.io/blog/apache-kafka-0-8-2-1-release> <https://www.confluent.io/blog/a-comprehensive-rest-proxy-for-kafka> <https://www.confluent.io/blog/bottled-water-real-time-integration-of-postgresql-and-kafka> <https://www.confluent.io/blog/compatibility-testing-for-apache-kafka>
<https://www.confluent.io/blog/using-logs-to-build-a-solid-data-infrastructure-or-why-dual-writes-are-a-bad-idea> <https://www.confluent.io/blog/hands-free-kafka-replication-a-lesson-in-operational-simplicity>
<https://www.confluent.io/blog/making-apache-kafka-elastic-with-apache-mesos>
<https://www.confluent.io/blog/compression-in-apache-kafka-is-now-34-percent-faster>
<https://www.confluent.io/blog/apache-kafka-samza-and-the-unix-philosophy-of-distributed-data>
<https://www.confluent.io/blog/log-compaction-highlights-stream-processing-kafka-2015-08>
<https://www.confluent.io/blog/distributed-consensus-reloaded-apache-zookeeper-and-replication-in-kafka>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-september-2015> <https://www.confluent.io/blog/apache-kafka-hits-1-1-trillion-messages-per-day-joins-the-4-comma-club> <https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-october-2015> <https://www.confluent.io/blog/apache-kafka-purgatory-hierarchical-timing-wheels>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-november-2015> <https://www.confluent.io/blog/apache-kafka-0-9-is-released>
<https://www.confluent.io/blog/how-to-build-a-scalable-etl-pipeline-with-kafka-connect>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-january-2016> <https://www.confluent.io/blog/290-reasons-to-upgrade-to-apache-kafka-0-9-0-0>
<https://www.confluent.io/blog/tutorial-getting-started-with-the-new-apache-kafka-0-9-consumer-client>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-february-2016> <https://www.confluent.io/blog/announcing-kafka-connect-building-large-scale-low-latency-data-pipelines> <https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-march-2016> <https://www.confluent.io/blog/log-compaction-highlights-in-the-kafka-and-stream-processing-community-april-2016> <https://www.confluent.io/blog/deploying-apache-kafka-on-aws-elastic-block-store-ebs> <https://www.confluent.io/blog/log-compaction-kafka-summit-edition-may-2016>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-june-2016> <https://www.confluent.io/blog/kafka-connect-sink-for-postgresql-from-justone-database> <https://www.confluent.io/blog/building-a-streaming-analytics-stack-with-apache-kafka-and-druid>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-july-2016> <https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-august-2016> <https://www.confluent.io/blog/kafka-connect-cassandra-sink-the-perfect-match> <https://www.confluent.io/blog/log-compaction-highlights-apache-kafka-stream-processing-community-september-2016> <https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-october-2016> <https://www.confluent.io/blog/streaming-data-oracle-using-oracle-goldengate-kafka-connect> <https://www.confluent.io/blog/announcing-apache-kafka-0-10-1-0> <https://www.confluent.io/blog/bloglog-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-november-2016> <https://www.confluent.io/blog/2017-kafka-summits>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-december-2016> <https://www.confluent.io/blog/data-dichotomy-rethinking-the-way-we-treat-data-and-services> <https://www.confluent.io/blog/log-compaction-highlights-apache-kafka-stream-processing-community-january-2017> <https://www.confluent.io/blog/apache-kafka-getting-started>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-february-2017> <https://www.confluent.io/blog/first-annual-state-apache-kafka-client-use-survey>

<https://www.confluent.io/blog/log-compaction-highlights-apache-kafka-stream-processing-community-march-2017> <https://www.confluent.io/blog/creating-data-pipeline-kafka-connect-api-architecture-operations>
<https://www.confluent.io/blog/stories-front-lessons-learned-supporting-apache-kafka>
<https://www.confluent.io/blog/2017-apache-kafka-survey-streaming-data-on-the-rise>
<https://www.confluent.io/blog/optimizing-apache-kafka-deployment> <https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-may-2017>
<https://www.confluent.io/blog/introduction-to-apache-kafka-for-python-programmers>
<https://www.confluent.io/blog/log-compaction-highlights-in-the-apache-kafka-and-stream-processing-community-june-2017> <https://www.confluent.io/blog/building-real-time-streaming-etl-pipeline-20-minutes>
<https://www.confluent.io/blog/exactly-once-semantics-are-possible-heres-how-apache-kafka-does-it>
<https://www.confluent.io/blog/upgrading-apache-kafka-clients-just-got-easier>
<https://www.confluent.io/blog/apache-kafka-for-service-architectures> <https://www.confluent.io/blog/chain-services-exactly-guarantees> <https://www.confluent.io/blog/simplest-useful-kafka-connect-data-pipeline-world-thereabouts-part-1> <https://www.confluent.io/blog/the-simplest-useful-kafka-connect-data-pipeline-in-the-world-or-thereabouts-part-2> <https://www.confluent.io/blog/simplest-useful-kafka-connect-data-pipeline-world-thereabouts-part-3> <https://www.confluent.io/blog/apache-kafka-tested>
<https://www.confluent.io/blog/okay-store-data-apache-kafka> <https://www.confluent.io/blog/build-deploy-scalable-machine-learning-production-apache-kafka> <https://www.confluent.io/blog/apache-kafka-goes-1-0>
<https://www.confluent.io/blog/toward-functional-programming-analogy-microservices>
<https://www.confluent.io/blog/handling-gdpr-log-forget> <https://www.confluent.io/blog/put-several-event-types-kafka-topic> <https://www.confluent.io/blog/event-sourcing-using-apache-kafka>
<https://www.confluent.io/blog/no-more-silos-how-to-integrate-your-databases-with-apache-kafka-and-cdc>
<https://www.confluent.io/blog/real-time-syslog-processing-with-apache-kafka-and-ksql-part-2-event-driven-alerting-with-slack> <https://www.confluent.io/blog/analytics-bare-metal-xenon-kafka-connect>
<https://www.confluent.io/blog/real-time-syslog-processing-apache-kafka-ksql-enriching-events-with-external-data> <https://www.confluent.io/blog/apache-kafka-vs-enterprise-service-bus-esb-friends-enemies-or-frenemies> <https://www.confluent.io/blog/apache-kafka-spring-boot-application>
<https://www.confluent.io/blog/apache-kafka-supports-200k-partitions-per-cluster>
<https://www.confluent.io/blog/kafka-connect-deep-dive-converters-serialization-explained>
<https://www.confluent.io/blog/using-apache-kafka-drive-cutting-edge-machine-learning>
<https://www.confluent.io/blog/easy-ways-generate-test-data-kafka> <https://www.confluent.io/blog/machine-learning-with-python-jupyter-ksql-tensorflow> <https://www.confluent.io/blog/kafka-connect-deep-dive-jdbc-source-connector> <https://www.confluent.io/blog/spring-for-apache-kafka-deep-dive-part-1-error-handling-message-conversion-transaction-support> <https://www.confluent.io/blog/spring-for-apache-kafka-deep-dive-part-2-apache-kafka-spring-cloud-stream> <https://www.confluent.io/blog/kafka-connect-deep-dive-error-handling-dead-letter-queues> <https://www.confluent.io/blog/consuming-messages-out-of-apache-kafka-in-a-browser> <https://www.confluent.io/blog/apache-kafka-to-amazon-s3-exactly-once>
<https://www.confluent.io/blog/12-programming-languages-walk-into-a-kafka-cluster>
<https://www.confluent.io/blog/apache-kafka-data-access-semantics-consumers-and-membership>
<https://www.confluent.io/blog/schemas-contracts-compatibility> <https://www.confluent.io/blog/spring-for-apache-kafka-deep-dive-part-3-apache-kafka-and-spring-cloud-data-flow>
<https://www.confluent.io/blog/spring-for-apache-kafka-deep-dive-part-4-continuous-delivery-of-event-streaming-pipelines> <https://www.confluent.io/blog/streaming-data-from-the-universe-with-apache-kafka>
<https://www.confluent.io/blog/building-a-scalable-search-architecture> <https://www.confluent.io/blog/whats-new-in-apache-kafka-2-3> <https://www.confluent.io/blog/microservices-apache-kafka-domain-driven-design>
<https://www.confluent.io/blog/designing-the-net-api-for-apache-kafka> [2 / 5](https://www.confluent.io/blog/kafka-</p></div><div data-bbox=)

listeners-explained <https://www.confluent.io/blog/getting-started-mongodb-connector-for-apache-kafka-and-mongodb> <https://www.confluent.io/blog/fault-tolerance-distributed-systems-tracing-with-apache-kafka-jaeger> <https://www.confluent.io/blog/announcing-apache-kafka-tutorials> <https://www.confluent.io/blog/kafka-connect-improvements-in-apache-kafka-2-3> <https://www.confluent.io/blog/transactional-systems-with-apache-kafka> <https://www.confluent.io/blog/incremental-cooperative-rebalancing-in-kafka> <https://www.confluent.io/blog/analytics-with-apache-kafka-and-rockset> <https://www.confluent.io/blog/apache-kafka-spring-boot-tutorial-for-pivotal-application-service> <https://www.confluent.io/blog/iot-with-kafka-connect-mqtt-and-rest-proxy> <https://www.confluent.io/blog/build-streaming-etl-solutions-with-kafka-and-rail-data> <https://www.confluent.io/blog/create-dynamic-kafka-connect-source-connectors> <https://www.confluent.io/blog/getting-started-with-rust-and-kafka> <https://www.confluent.io/blog/machine-learning-real-time-analytics-models-in-kafka-applications> <https://www.confluent.io/blog/kafka-connect-single-message-transformation-tutorial-with-examples> <https://www.confluent.io/blog/kafka-python-asyncio-integration> <https://www.confluent.io/blog/apache-kafka-2-4-latest-version-updates> <https://www.confluent.io/blog/apache-kafka-producer-improvements-sticky-partitioner> <https://www.confluent.io/blog/configure-kafka-to-minimize-latency> <https://www.confluent.io/blog/confluent-developer-offers-kafka-tutorials-resources-guides> <https://www.confluent.io/blog/kafka-elasticsearch-connector-tutorial> <https://www.confluent.io/blog/how-to-share-kafka-connectors-on-confluent-hub> <https://www.confluent.io/blog/apache-kafka-2-5-latest-version-updates> <https://www.confluent.io/blog/webify-event-streams-using-kafka-connect-http-sink> <https://www.confluent.io/blog/removing-zookeeper-dependency-in-kafka> <https://www.confluent.io/blog/stream-analyze-visualize-data-with-kafka-ksqldb-and-friends> <https://www.confluent.io/blog/kafka-client-cannot-connect-to-broker-on-aws-on-docker-etc> <https://www.confluent.io/blog/guide-to-kafka-pricing-and-diy-open-source-costs> <https://www.confluent.io/blog/spring-kafka-can-your-kafka-consumers-handle-a-poison-pill> <https://www.confluent.io/blog/kafka-data-privacy-security-and-compliance> <https://www.confluent.io/blog/simplified-robust-exactly-one-semantics-in-kafka-2-5> <https://www.confluent.io/blog/cdc-and-streaming-analytics-using-debezium-kafka> <https://www.confluent.io/blog/apache-kafka-2-6-updates> <https://www.confluent.io/blog/kafka-consumer-multi-threaded-messaging> <https://www.confluent.io/blog/kafka-fastest-messaging-system> <https://www.confluent.io/blog/prioritize-messages-in-kafka> <https://www.confluent.io/blog/devops-for-apache-kafka-with-kubernetes-and-gitops> <https://www.confluent.io/blog/5-things-every-kafka-developer-should-know> <https://www.confluent.io/blog/how-to-prepare-for-kip-500-kafka-zookeeper-removal-guide> <https://www.confluent.io/blog/streaming-data-pipelines-and-analytics-with-qlik-and-kafka> <https://www.confluent.io/blog/advanced-testing-techniques-for-spring-kafka> <https://www.confluent.io/blog/kafka-scala-tutorial-for-beginners> <https://www.confluent.io/blog/transactional-machine-learning-with-maads-viper-and-apache-kafka> <https://www.confluent.io/blog/spring-microservices-into-production-with-kubernetes-gitops> <https://www.confluent.io/blog/apache-kafka-2-7-features-updates-improvements> <https://www.confluent.io/blog/apache-kafka-intro-how-kafka-works> <https://www.confluent.io/blog/best-kafka-tools-that-boost-developer-productivity> <https://www.confluent.io/blog/getting-started-with-kafkajs> <https://www.confluent.io/blog/write-a-kafka-connect-connector-with-configuration-handling> <https://www.confluent.io/blog/42-ways-zookeeper-removal-improves-kafka> <https://www.confluent.io/blog/kafka-sap-integration-with-kafka-connect-odp-source-connector> <https://www.confluent.io/blog/containerized-testing-with-kerberos-and-ssh> <https://www.confluent.io/blog/push-or-pull-data-into-kafka-connect> [3 / 5](https://www.confluent.io/blog/real-time-</p></div><div data-bbox=)

analytics-with-kafka-and-pinot <https://www.confluent.io/blog/kafka-without-zookeeper-a-sneak-peek>
<https://www.confluent.io/blog/modern-cpp-kafka-api-for-safe-easy-messaging>
<https://www.confluent.io/blog/kafka-2-8-0-features-and-improvements-with-early-access-to-kip-500>
<https://www.confluent.io/blog/how-to-survive-a-kafka-outage> <https://www.confluent.io/blog/error-handling-patterns-in-kafka> <https://www.confluent.io/blog/data-analytics-pipeline-with-kafka-and-rstudio>
<https://www.confluent.io/blog/confluent-developer-launches-free-apache-kafka-courses-and-tutorials-online>
<https://www.confluent.io/blog/build-cross-platform-kafka-applications-using-c-and-dotnet-5>
<https://www.confluent.io/blog/kafka-connect-tutorial> <https://www.confluent.io/blog/intro-to-graphql-an-api-for-kafka-data> <https://www.confluent.io/blog/apache-kafka-3-0-major-improvements-and-new-features>
<https://www.confluent.io/blog/how-baader-built-a-predictive-analytics-machine-learning-system-with-kafka-and-rstudio> <https://www.confluent.io/blog/using-spring-with-apache-kafka-full-introduction>
<https://www.confluent.io/blog/best-tutorials-for-getting-started-with-apache-kafka>
<https://www.confluent.io/blog/getting-started-with-apache-kafka-in-python>
<https://www.confluent.io/blog/top-10-kafka-use-cases-lessons-guides-from-200-experts>
<https://www.confluent.io/blog/5-common-pitfalls-when-using-apache-kafka>
<https://www.confluent.io/blog/streaming-data-pipeline-with-apache-kafka-and-ksqldb>
<https://www.confluent.io/blog/apache-kafka-3-1-version-features-and-updates>
<https://www.confluent.io/blog/how-to-make-apache-kafka-even-faster> <https://www.confluent.io/blog/why-replace-zookeeper-with-kafka-raft-the-log-of-all-logs> <https://www.confluent.io/blog/apache-kafka-architecture-and-internals-by-jun-rao> <https://www.confluent.io/blog/apache-kafka-3-2-0-new-features-and-updates> <https://www.confluent.io/blog/apache-kafka-ci-cd-with-github> <https://www.confluent.io/blog/using-data-pipelines-for-real-time-alerting-with-ksqldb> <https://www.confluent.io/blog/real-time-detection-monitoring-with-apache-kafka> <https://www.confluent.io/blog/kafka-security-how-to-secure-real-time-data-streams> <https://www.confluent.io/blog/what-is-kraft-and-how-do-you-use-it>
<https://www.confluent.io/blog/real-time-gaming-infrastructure-kafka-ksqldb-websockets>
<https://www.confluent.io/blog/sync-databases-and-remove-silos-with-kafka-cdc>
<https://www.confluent.io/blog/event-driven-microservices-with-python-and-kafka>
<https://www.confluent.io/blog/apache-kafka-3-3-0-new-features-and-updates>
<https://www.confluent.io/blog/designing-events-and-event-streams-introduction-and-best-practices>
<https://www.confluent.io/blog/schema-registry-for-beginners> <https://www.confluent.io/blog/build-real-time-iot-application-with-apache-kafka-and-ksqldb> <https://www.confluent.io/blog/debug-apache-kafka-reduced-message-throughput> <https://www.confluent.io/blog/how-to-secure-kafka-clusters>
<https://www.confluent.io/blog/debug-apache-kafka-pt-2> <https://www.confluent.io/blog/debug-apache-kafka-pt-3> <https://www.confluent.io/blog/how-to-turn-apache-kafka-project-into-a-talk>
<https://www.confluent.io/blog/debug-apache-kafka-pt-4> <https://www.confluent.io/blog/configuring-apache-kafka-consumer-group-ids> <https://www.confluent.io/blog/apache-kafka-3-4-0-new-features-and-updates>
<https://www.confluent.io/blog/windowing-in-kafka-streams> <https://www.confluent.io/blog/co-partitioning-in-kafka-streams> <https://www.confluent.io/blog/dynamic-vs-static-kafka-consumer-rebalancing>
<https://www.confluent.io/blog/is-kafka-a-database-with-ksqldb> <https://www.confluent.io/blog/multi-geo-replication-in-apache-kafka> <https://www.confluent.io/blog/kafka-python-developer-guide>
<https://www.confluent.io/blog/how-to-fix-unknown-magic-byte-errors-in-apache-kafka>
<https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-1-infrastructure>
<https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-2-development-and-operations> <https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-3-lower-cloud-costs> <https://www.confluent.io/blog/understanding-and-optimizing-your-kafka-costs-part-4-savings-challenge> <https://www.confluent.io/blog/introducing-apache-kafka-3-5> [4 / 5](https://www.confluent.io/blog/what-</p></div><div data-bbox=)

is-an-event-in-the-apache-kafka-ecosystem <https://www.confluent.io/blog/kafka-clusters-and-cluster-management> <https://www.confluent.io/blog/standardized-hashing-across-java-and-non-java-producers> <https://www.confluent.io/blog/what-is-an-apache-kafka-cluster> <https://www.confluent.io/blog/how-to-increase-throughput-on-kafka-connect-source-connectors> <https://www.confluent.io/blog/design-and-deployment-considerations-for-deploying-apache-kafka-on-aws> <https://www.confluent.io/blog/set-up-and-run-kafka-on-windows-linux-wsl-2> <https://www.confluent.io/blog/apache-kafka-message-compression> <https://www.confluent.io/blog/kafka-consumer-group-partition-strategy> <https://www.confluent.io/blog/introducing-apache-kafka-3-6> <https://www.confluent.io/blog/best-kafka-tutorials-examples-and-learning-resources> <https://www.confluent.io/blog/how-to-send-messages-with-librdkafka> <https://www.confluent.io/blog/enterprise-kafka-cluster-strategies-and-best-practices> <https://www.confluent.io/blog/introducing-apache-kafka-3-7> <https://www.confluent.io/blog/process-github-data-with-kafka-streams> <https://www.confluent.io/blog/how-to-use-kafka-docker-composer> <https://www.confluent.io/blog/how-to-write-KIPs-for-apache-kafka> <https://www.confluent.io/blog/sizing-scaling-kafka-clusters-with-confluent-cloud> <https://www.confluent.io/blog/dual-write-problem> <https://www.confluent.io/blog/running-kafka-at-the-edge-requires-confluent> <https://www.confluent.io/blog/building-full-stack-app-with-kafka-and-nodejs> <https://www.confluent.io/blog/introducing-apache-kafka-3-8> <https://www.confluent.io/blog/kafka-producer-internals-preparing-event-data> <https://www.confluent.io/blog/kafka-producer-internals-handling-producer-request> <https://www.confluent.io/blog/preparing-the-consumer-fetch-kafka-producer-and-consumer-internals-part-3> <https://www.confluent.io/blog/introducing-apache-kafka-3-9> <https://www.confluent.io/blog/kafka-producer-and-consumer-internals-4-consumer-fetch-requests> <https://www.confluent.io/blog/introducing-confluent-kafka-javascript> <https://www.confluent.io/blog/queues-on-kafka> <https://www.confluent.io/blog/confluent-and-lambda> <https://www.confluent.io/blog/event-driven-multi-agent-systems> <https://www.confluent.io/blog/delete-kafka-topic> <https://www.confluent.io/blog/latest-apache-kafka-release> <https://www.confluent.io/blog/ai-agents-using-anthropic-mcp> <https://www.confluent.io/blog/guide-to-consumer-offsets> <https://www.confluent.io/blog/building-streaming-data-pipelines-part-1> <https://www.confluent.io/blog/the-future-of-ai-agents-is-event-driven> <https://www.confluent.io/blog/querying-kafka-natural-language> <https://www.confluent.io/blog/kip-848-consumer-rebalance-protocol> <https://www.confluent.io/blog/apache-kafka-tableflow-vs-zero-etl> <https://www.confluent.io/blog/introducing-apache-kafka-4-1>